Multi-View RGB-D Fusion for Measuring Sediment Rugosity

An In Depth Characterisation of the D456 Camera, and Multi-View Fusion Pipeline

Joab Gray Kloppers

Department of Electrical Engineering
University of Cape Town
Rondebosch, Cape Town
South Africa

Supervisor
RA **Verrinder**

Co-supervisor(s)
A/Prof D **Pillay**, G **Maniel**

October 2026

BSc (Eng) final year project report submitted in partial fulfilment of the requirements for the degree of Bachelor of Science in Electrical and Computer Engineering in the Department of Electrical Engineering at the University of Cape Town

Keywords: Active Atereo; RGB-D; Multi-view Fusion; Surface Rugosity; Sediment Microtopography; Bioturbation

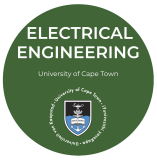

Declaration

I, **Joab Gray Kloppers**, hereby:

1. know that plagiarism is wrong. Plagiarism is to use another's work and pretend that it is one's own.
2. have used the **IEEE** convention for citation and referencing. Each contribution to, and quotation in, this report from the work(s) of other people has been attributed, and has been cited and referenced. Any section taken from an internet source has been referenced to that source.
3. declare that:
    a. this report is my own work, and apart from the normal guidance from my supervisor(s), I have received no assistance except as acknowledged.
    b. neither the substance nor any part of the above report has been submitted in the past, or is being, or is to be submitted for a degree at this University or at any other university.
    c. I have not paid a third party to complete my work on my behalf.
    d. my use of artificial intelligence software has been limited to:

    Anthropic: Claude for the use of generating code and assistance with report writing. Google: Gemini for the use of exploring and summarising literature.

4. have not allowed and will not allow anyone to copy my work with the intention of passing it off as their own work.
5. acknowledge that copying someone else's work, or part of it, is wrong, and declare that this is my own work.
6. grant the University of Cape Town free license to reproduce the above thesis in whole or in part, for the purpose of research.

Word count 21147 words (excluding preamble, appendices and bibliography)
Page count 55 pages (excluding preamble, appendices and bibliography)

Joab Gray Kloppers
Department of Electrical Engineering
University of Cape Town

Sat 24 October, 2026

Abstract

Multi-View RGB-D Fusion for Measuring Sediment Rugosity

Joab Gray **Kloppers**

Sat 24 October, 2026

Sediment microtopography, and rugosity in particular, has been proposed as a non-destructive proxy for the bioturbation activity of the sand prawn Callichirus kraussi on the intertidal flats of Langebaan Lagoon. An existing scanning system based on an Intel RealSense D456 active stereo camera has been implemented, however has never been characterised for these environmental conditions, and it analysed each of its repeated multi-view scans individually. This project characterised the D456 for close-range scanning of sediment-like surfaces and developed a processing pipeline that crops, registers and fuses multi-view scans before computing ISO 25178-2 areal roughness parameters.

Planar tests showed that noise increased with working distance, tilt and ambient light. Plane-fit RMS rose from 1.22 mm at 604 mm to 4.58 mm at 903 mm. Scans of 3D-printed reference objects showed that viewing geometry matters more than software settings. Increasing resolution to 1280×720 reduced RMSE by 13.4%, and the High Density preset gave the best F-score. Post-processing filters improved RMSE by at most 0.24 mm.

The pipeline uses AprilTags and IMU gravity readings for coarse alignment, refines the alignment with point-to-plane ICP, and reconciles the overlapping scans with a confidence-weighted trimmed mean. Fusion improved every metric at every height. On the large reference surface, it raised the F-score by 14.0 percentage points, reaching 90.9% at 641 mm. Twelve well-spread views matched the accuracy of a full 44-view sweep.

On flat laboratory sand, fused scans measured a rugosity of 1.0034, against 1.0193 for single scans. On shaped sand, single scans added 2.1 to 2.6 percentage points of spurious excess area, an amount comparable to the signal itself. Rugosity is therefore only comparable between surveys that use a fixed capture and processing protocol, and this report recommends such a protocol.

Acknowledgements

This final year project has been the capstone for the past four years of my life. It represents my entire journey going through the electrical engineering degree, every fascinating discovery, disappointing result and late night. Despite the ups and downs throughout the degree, I am so glad that i pursued this direction. These four years, and especially this project reminded me about my long standing fascination for understanding, and building systems that can help us better understand and influence the world around us. To this, I owe a great deal of appreciation for the people around me, who allowed me to partake in this journey.

Firstly, I would like to thank my supervisor Robyn Verrinder, who provided me with the opportunity to work on this project, and for providing me with guidance at every stage in the project. She provided me with the ability to explore the project on my own, whilst simultaneously making sure I never got lost. Whilst daunting at first, I really appreciate the lessons it taught me, and gives me great satisfaction with the final stages.

Secondly, I would like to extend a massive thank you to my family. Without their continuous support, I would never have been able to get to the point I am now. They have provided me with endless opportunities, and have always been there to help me when needed, and to guide me through everything.

I would also like to thank all my friends, who have taken this journey with me. Alex, Colby, Amir, Rafael, Justin, Ebrahiem and Glen. Thank you for keeping me company, and most importantly an excuse to procrastinate, during those long days where doing any work was extremely tedious.

Lastly, a thank you to all the masters and postgraduate students in the ARU, who welcomed us to their office space, and gave help with all equipment, especially when it came to us naive fourth years.

Contents

Declaration . . . . . **i**
Abstract . . . . . **ii**
Acknowledgements . . . . . **iii**
List of Figures . . . . . **viii**
List of Tables . . . . . **ix**
List of Symbols . . . . . **x**
List of Symbols . . . . . **x**
List of Acronyms and Glossary . . . . . **xiii**
Chapter 1 Introduction **1**
1.1 Background 1
1.1.1 Sand prawns and bioturbation 1
1.1.2 Measuring sediment microtopography 1
1.2 Problem statement 2
1.3 Project objectives 2
1.4 Scope, limitations and assumptions 2
1.5 Plan of development 3
Chapter 2 Literature review **4**
2.1 Sand prawn bioturbation and its ecological significance 4
2.2 Sediment microtopography as a proxy for benthic activity 5
2.3 Non-contact measurement of sediment surfaces 6
2.4 Characterisation of active stereo depth cameras 6
2.5 Registration and fusion of multi-view depth scans 7
2.6 Summary and research gap 8
Chapter 3 Theoretical Background **9**
3.1 Depth Capture 9
3.1.1 Camera Model 9
3.1.2 Stereo Triangulation 10
3.1.3 Passive Stereo, Structured Light and Active Stereo 10
3.2 Registration 10
3.2.1 Rigid Transformations and Closed-Form Alignment 11
3.2.2 Fiducial Markers 11
3.2.3 Feature-Based Coarse Alignment 11
3.2.4 Iterative Closest Point 12
3.3 Fusion 12
3.4 Evaluation 13
3.5 Roughness Calculation 13

Chapter 4 System Design **15**
4.1 Requirements Definition 15
4.2 System Overview 15
4.3 Hardware Rig 16
4.3.1 Frame 16
4.3.2 Camera, Tilt and Rotation Mechanism 16
4.3.3 AprilTag Boundary Markers 16
4.4 Software Architecture and Processing Pipeline 16
4.4.1 Capture Software 16
4.4.2 Cropping and Multi-Angle Fusion 17
4.4.3 Microtopography Processing 17
4.5 Evaluation Framework and Ground Truth 18
4.5.1 Reference Surfaces 18
4.5.2 Alignment and Metrics 18
4.6 Summary 18
Chapter 5 Intel RealSense Characterisation **19**
5.1 Plane Surface Characterisation 19
5.1.1 Capture procedure and metrics 19
5.1.2 Planar Characterisation Methodology 20
5.1.3 Planar testing Results 21
5.1.4 Discusion 21
5.2 3D Reference Characterisation 23
5.2.1 Reference objects and evaluation 24
5.2.2 Testing Procedure 24
5.2.3 3D Characterisation Results 25
5.2.4 Discussion 27
Chapter 6 Data Processing Pipeline **32**
6.1 Point Cloud Cropping 32
6.1.1 Finding the tags in each scan 33
6.1.2 Learning the tag layout of a session 33
6.1.3 Relating 3D points to image pixels 33
6.1.4 Building the crop volume 34
6.1.5 Review and manual fallback 35
6.1.6 Outputs 35
6.1.7 Limitations 35
6.2 Multi-Scan Fusion 35
6.2.1 Preprocessing 36
6.2.2 Initial poses from the rig and IMU 36
6.2.3 Coarse alignment 37
6.2.4 Fine registration 37
6.2.5 Reconciling the overlapping scans 38
6.2.6 Limitations 39
6.3 Surface Roughness Estimation 39
6.3.1 Levelling 40
6.3.2 Detrending 40
6.3.3 Gridding 40

6.3.4 Height parameters . . . 41
6.3.5 Rugosity and slope . . . 41
6.3.6 Spatial and spectral parameters . . . 41
6.3.7 Additional measures and uncertainty . . . 42
6.3.8 Validation . . . 42
6.3.9 Limitations . . . 43
Chapter 7 Experimental Validation . . . 44
7.1 Controlled Targets . . . 44
7.1.1 Process . . . 44
7.1.2 Results . . . 44
7.2 Controlled Laboratory Sand . . . 45
7.2.1 Results . . . 47
Chapter 8 Discussion . . . 48
8.1 Camera Configuration . . . 48
8.2 Working Distance, Viewing Angle and Lighting . . . 48
8.3 Fusion and Roughness Estimation . . . 49
8.4 Implications for the Overarching Project . . . 50
Chapter 9 Conclusions . . . 51
Chapter 10 Further work . . . 54
Chapter A GitHub repository . . . 55
A.1 Title for appendix A.1 . . . 55
Appendix B Theory . . . 56
Appendix C Reference Objects . . . 57
Appendix D Characterisation Results . . . 59
D.1 Planar Characterisation Results . . . 59
D.2 Camera Configuration Results . . . 61
D.3 Filter Configuration Results . . . 62
D.4 Fusion Combination Results . . . 63
Appendix E Characterisation Plots . . . 64
E.1 Filter Heatmaps . . . 64
E.2 Combination Heatmaps . . . 67
E.3 Parameter Heatmap . . . 68
E.4 Accuracy Vs Tilt . . . 69
E.5 Single Vs Fused Error Histograms . . . 70
Appendix F Validation Results . . . 72
F.1 Reference Shapes Rugosity . . . 72
Appendix B: Ethics approval . . . 75
Appendix C: Topic description . . . 76
Appendix D: ECSA Graduate Attributes . . . 77
Project mapping to ECSA Graduate Attributes . . . 77
Bibliography . . . 80

List of Figures

2.1 Surface Mounds and Burrow Openings Produced by *Callichirus kraussi* on an Intertidal Sandflat. Taken on 14/07/2026 During a Field Trip . . . 4
2.2 Measuring microtopography as an ecological proxy. . . . 5
2.3 The RGB-D framework of Azhar et al. [8]. . . . 6
2.4 Reflectance of sand decreasing non-linearly as moisture content increases [2]. . . . 7
2.5 A single noisy depth frame compared with the surface fused from many frames [33]. . . . 8

3.1 Rectified stereo geometry: camera centres O_L, O_R, baseline B, focal length f and the projections x_L, x_R of a point P at depth Z. . . . 10

4.1 System overview. [TODO: Replace with a fuller diagram showing the data formats (.bag, .ply) passed between stages.] . . . 15
4.2 System overview. [TODO: Replace with a fuller diagram showing the data formats (.bag, .ply) passed between stages.] . . . 17

5.1 Plots Showing Planar Performance Against Height. Per-trial Median with IQR Over n = 385–552 Frames. . . . 21
5.2 Plots Showing Planar Performance Against Scan Angle. Per-trial Median with IQR Over n = 527–574 Frames. . . . 22
5.3 Plots Showing Planar Performance Against Light Intensity, for Indoor Conditions with Foam Board. Per-trial Median with IQR Over n = 592–599 Frames. . . . 22
5.4 Plots Showing 3D Performance Against Camera Resolution. Median with IQR of n = 15 scans per point, Taken Over Three Depth Units, and Five Presets. . . . 26
5.5 Plots Showing 3D Performance Against Depth Units. Median with IQR of n = 15 scans per point, Taken Over Three Resolutions, and Five Presets. . . . 26
5.6 Plots Showing 3D Performance Against Camera Preset. Median with IQR of n = 15 scans per point, Taken Over Three Depth Units, and Three Resolutions. . . . 27
5.7 Plots Showing 3D Performance Comparison Between No Filter and Best Performing Filter, Best Fit (R^2 0.00 to 1.00). Median with IQR of n = 4–5 Scans Per Point. . . . 28
5.8 Plots Showing 3D Performance Between No Fusion and All View-Point Fusion. Median with IQR of n = 1 Point Clouds Per Point. . . . 29
5.9 Grey cells: no data for that combination. Median with IQR of n = 12 point clouds per cell. Fixed: Kind = fused; Fusion cell = 0.0005; Scan voxel = 0.0005; Status = ok. Median taken over: Object (2), Height (6). . . . 30
5.10 Visual Representation of Single Scan Performance Using Small Egg Box Surface . . . 30
5.11 Visual Representation of Fused Scan Performance Using Small Egg Box Surface . . . 30

7.1 Mean ± SD of n = 1 rows per bar. . . . 46
7.2 Mean ± SD of n = 1 rows per bar. Fixed: Grid nx = 45. . . . 46
7.3 Mean ± SD of n = 4 rows per bar. Averaged over: Detrend (4), Primary (2), Flags (3), Str bounded (2), Detrend n coef (4), Detrend order (4) and 1 more. . . . 47

B.1 D456 Semi-Random IR Dot Pattern Projected Onto a White Wall, With a Hand . . . 56

C.1 Large Egg Box Surface Reference Object. . . . 57
C.2 Small Egg Box Surface Reference Object, Printed with amplitude A = 15mm. . . . 57

C.3 Large Pyramid Surface Reference Object. 58

E.1 Median with IRQ % CI of the mean of n = 5–295 scans per cell. Fixed: Tilt = 0. Averaged over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial filter on (2), Spatial magnitude (6), Spatial α (8) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded. 64
E.2 Median with IRQ % CI of the mean of n = 5–815 scans per cell. Fixed: Tilt = 0. Averaged over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial filter on (2), Spatial magnitude (6), Spatial α (8) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded. 64
E.3 Median with IQR of n = 5–315 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial magnitude (6), Spatial δ (9), Spatial hole fill (7) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded. 65
E.4 Median with IQR of n = 5–385 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial magnitude (6), Spatial α (8), Spatial hole fill (7) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded. 65
E.5 Median with IQR of n = 5–950 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Spatial filter on (2), Spatial magnitude (6), Spatial α (8), Spatial δ (9), Spatial hole fill (7) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded. 66
E.6 Median with IQR of n = 55–915 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial filter on (2), Spatial magnitude (6), Spatial α (8) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded. 66
E.7 Grey cells: no data for that combination. Median with IQR of n = 6 point clouds per cell. Fixed: Kind = fused; Fusion cell = 0.0005; Scan voxel = 0.0005; Status = ok. Median taken over: Height (6). 67
E.8 Median with IQR of n = 1–3 scans per cell. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Object (2), Preset (5). 68
E.9 Dashed lines: quadratic least-squares fit per series and panel (R^2 0.27 to 0.92). Median with IQR of n = 5–6 point clouds per point. Fixed: Kind = single; View family = single; Number of rotations = 1; Number of tilts = 1; Number of scans = 1. Median taken over: Height (6), Status (3). 3 rows without a value excluded. 69
E.10 Dashed lines: quadratic least-squares fit per series and panel (R^2 0.09 to 0.86). Median with IQR of n = 5–6 point clouds per point. Fixed: Kind = single; View family = single; Number of rotations = 1; Number of tilts = 1; Number of scans = 1. Median taken over: Height (6), Status (3). 3 rows without a value excluded. 70
E.11 Histogram Showing Spread of Signed Error 70
E.12 Histogram Showing Spread of Signed Error 71

F.1 Big Wave Detrend Explained Single 72
F.2 Big Wave Detrend Explained Fused 73
F.3 Small Wave Detrend Explained Single 73
F.4 Small Wave Detrend Explained Fused 74

List of Tables

2.1 Non-contact methods for measuring sediment surfaces. 6
2.2 Published depth-camera characterisations relevant to this project. SL: structured light; AS: active stereo; PS: passive stereo; ToF: time of flight. 7
2.3 Research gaps and the objectives that address them. 8

4.1 Hardware summary. 16

5.1 Planar characterisation under different lighting conditions at a height of approximately 0.9 m. Lux values are indicative (smartphone sensor). 23
5.2 Filter parameter ranges tested (one factor at a time). 25
5.3 Post-processing filter settings of the best-performing configuration for all objects combined and for each object. 25

7.1 Each value is over n = 1 rows. 45
7.2 Each value is over n = 1 rows. Fixed: Grid nx = 45. 45
7.3 Each value is over n = 4 rows. Averaged over: Detrend (4), Primary (2), Flags (3), Str bounded (2), Detrend n coef (4), Detrend order (4) and 1 more. 47

D.1 Results for Planar Characterisation Against Distance. Per-Trial Values Over n = 385-552 Frames. 59
D.2 Results for Planar Characterisation Against Scan Angle. Per-trial values over n=527-574 frames. 60
D.3 Results for Planar Characterisation Against Light Intensity. Per-trial values over n = 592-599 Frames. 60
D.4 Each value is over n = 15 scans. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Depth unit (3), Preset (5), Preset actual (5). 61
D.5 Each value is over n = 15 scans. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Resolution (3), Preset (5), Preset actual (5). 61
D.6 Each value is over n = 8–10 scans. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Resolution (3), Depth unit (3), Preset (5). 61
D.7 Each value is over n = 4–5 scans. Fixed: Tilt = 0; Spatial filter on = Off; HDR merge on = Off; Status = ok. Median taken over: Rotation (4). 62
D.8 Each value is over n = 1 point clouds. Fixed: Status = ok. 63

List of Symbols

Symbol	Description	Unit
Camera model (Section 3.1.1)		
(u, v)	Pixel coordinates in the image	px
λ	Projective scale factor (equal to the depth Z)	m
$\mathbf{X}_w = (X_w, Y_w, Z_w)^\top$	Point in the world frame	m
$(X, Y, Z)^\top$	Point in the camera frame	m
Z	Depth of a point along the optical axis	m
$\mathbf{K}$	Camera intrinsic matrix	–
f_x,f_y	Focal lengths along the image axes	px
c_x, c_y	Principal point	px
$\mathbf{R}$	Rotation matrix ($\mathbf{R}^\top\mathbf{R} = \mathbf{I}$, $\det \mathbf{R} = +1$)	–
$\mathbf{t}$	Translation vector	m
$\mathbf{I}$	Identity matrix	–
(x_d, y_d)	Normalised image coordinates after lens distortion, $(X/Z,\ Y/Z)$	–
Stereo triangulation (Section 3.1.2)		
O_L, O_R	Camera centres of the left and right imagers	–
P	Scene point observed by both imagers	–
x_L, x_R	Horizontal image coordinates of P in the left and right imagers	px
f	Focal length of the rectified imagers	px
B	Stereo baseline (95 mm for the D456)	mm
d	Disparity, $x_L - x_R$	px
σ_d	Standard deviation of the disparity matching error	px
σ_Z	Standard deviation of the depth error	mm
$Z_{\min}$	Minimum measurable depth, $fB/126$	mm
Registration (Section 3.2)		
$\mathbf{p}_i, \mathbf{q}_i$	Corresponding source and target points	m
$\bar{\mathbf{p}}, \bar{\mathbf{q}}$	Centroids of the source and target points	m
N	Number of point pairs (Section 3.2.1) or number of RANSAC trials (Section 3.2.3)	–
$\mathbf{H}$	Cross-covariance matrix of centred point pairs (Eq. 3.5)	m^2
$\mathbf{U}, \mathbf{\Sigma}, \mathbf{V}$	Factors of the singular value decomposition of $\mathbf{H}$	–
$\mathbf{H}$	Homography relating tag corners to image positions (Section 3.2.2)	–
$\mathbf{r}_1, \mathbf{r}_2, \mathbf{r}_3$	Columns of the tag rotation matrix	–
w	Fraction of correct correspondences (inlier ratio)	–
p	Desired probability of drawing an all-inlier RANSAC sample	–
$\mathbf{n}_i$	Surface normal at the target point $\mathbf{q}_i$	–
$E(\mathbf{R}, \mathbf{t})$	Point-to-plane ICP cost function	m^2
Fusion (Section 3.3)		
$z(x, y)$	Height field of the fused surface	mm
w_k	Weight of measurement k, $\propto 1/\sigma_k^2$	–
σ_k^2	Variance of measurement k	mm^2

Continued on next page

Symbol	Description	Unit
h	Moving least squares (MLS) fitting radius	mm
Evaluation (Sections 3.4 and 5.2.1)		
$\mathbf{s}_i$	Reconstructed point	mm
d_i	Signed distance from $\mathbf{s}_i$ to the closest point on the reference mesh	mm
N	Number of reconstructed points (Section 3.4)	–
μ	Bias (mean signed error)	mm
σ	Standard deviation of the signed error	mm
RMSE	Root-mean-square error	mm
MAE	Mean absolute error	mm
τ	Distance threshold for precision, completeness and F-score	mm
S	Set of scanned (reconstructed) points	–
R	Set of points sampled from the reference model	–
$P(\tau)$	Precision at threshold τ	%
$C(\tau), R(\tau)$	Completeness at threshold τ	%
$F(\tau)$	F-score at threshold τ	%
$\mathscr{S}$	Reconstruction point set (Eq. 3.8)	–
$\mathscr{M}$	Reference point set (Eq. 3.8)	–
$d(\mathbf{a}, \mathscr{B})$	Distance from point $\mathbf{a}$ to the nearest point of set $\mathscr{B}$	mm
d_{CD}	Chamfer distance	mm
d_{H}	Hausdorff distance	mm
Point cloud cropping (Section 6.1)		
ρ	Root-mean-square distance of reprojected points from the nearest integer pixel	px
$\mathbf{u}_i$	Reprojected pixel position of point i	px
$\mathbf{r}$	Direction of a pixel ray	–
$\mathbf{n}$	Normal of a fitted plane	–
d	Offset of a fitted plane (Eq. 6.3)	m
$\mathbf{X}$	Intersection of a pixel ray with a plane	m
$\mathbf{p}_1, \mathbf{p}_2, \mathbf{p}_3, \mathbf{p}_4$	Detected and predicted tag corners (parallelogram completion)	px
Fusion pipeline (Section 6.2)		
w_i	Weight of measurement i (Eq. 6.5)	–
$c_s(i)$	Registration confidence of the scan containing point i	–
$\hat{\mathbf{n}}_i$	Unit surface normal at point i	–
$\hat{\mathbf{v}}_i$	Unit vector from point i to the camera	–
r_i	Range of point i from the camera	m
r_0	Reference range for range weighting	m
$\mathbf{q}_j$	Map point paired with scan point $\mathbf{p}_i$	m
$\mathbf{n}_j$	Map normal at $\mathbf{q}_j$	–
Roughness estimation (Sections 3.5 and 6.3)		
$p(x, y)$	Polynomial trend surface removed by detrending	mm
a_{ij}	Coefficients of the detrending polynomial	mm
n	Degree of the detrending polynomial	–
z_k	Gridded residual height of cell k relative to the mean	mm
N	Number of valid grid cells (Section 6.3.4)	–
g	Grid spacing	mm
N_q	Number of valid grid squares	–
A_{true}	True surface area	mm^2
A_{proj}	Projected surface area	mm^2
$A_{\mathrm{triangles}}$	Area of an individual surface triangle	mm^2

Continued on next page

Symbol	Description	Unit
σ_n	Standard deviation of uncorrelated height noise	mm
S_a	Arithmetic mean height	mm
S_q	Root-mean-square height	mm
S_{sk}	Skewness of the height distribution	–
S_{ku}	Kurtosis of the height distribution	–
S_p	Maximum peak height	mm
S_v	Maximum valley depth	mm
S_z	Maximum height, $S_p + S_v$	mm
S_{dq}	Root-mean-square gradient (slope)	–
S_{dr}	Developed interfacial area ratio	‰
S_{al}	Autocorrelation length	mm
S_{tr}	Texture aspect ratio	–
S_{td}	Texture direction	°
q	Spatial wavenumber	rad/m
$C(q)$	Power spectral density as a function of wavenumber	mm^4
H	Hurst exponent	–
D_f	Fractal dimension, $3 - H$	–

Chapter 1

Introduction

1.1 Background

1.1.1 Sand prawns and bioturbation

Bioturbation, the reworking of sediment particles and the ventilation of burrows by animals [3], is one of the main ways in which benthic fauna shape soft-sediment ecosystems. Burrowing thalassinidean shrimps, the group to which sand prawns belong, are among the most influential bioturbators. Pillay and Branch [4] describe how their burrowing brings buried sediment to the surface, builds mounds that change the structure and increase the area of the sediment surface, and alters the physical and chemical properties of the sediment through burrow construction and irrigation.

On the South African coast, the sand prawn *Callichirus kraussi* is a key ecosystem engineer. Venter et al. [5] showed experimentally that sand prawns filter phytoplankton from the overlying water by pumping water in both directions through their burrows. Long-term estuarine data further showed reductions in phytoplankton biomass of roughly 70% in sand prawn-dominated areas, which suggests that sand prawns help to buffer coastal ecosystems against eutrophication. Their methodology, however, relied on manual visual counts of burrow openings, which are laborious and describe only one aspect of the altered surface. This project forms part of a wider study of sand prawn bioturbation on the intertidal flats of the Langebaan Lagoon. [TODO: Name the overarching project and research group, and cite or acknowledge it.]

1.1.2 Measuring sediment microtopography

Bioturbators leave mounds, burrow openings and depressions at the sediment surface. The microtopography of the surface, its relief at millimetre to centimetre scales, therefore carries a signature of the activity beneath it. Schenone et al. showed in laboratory incubations that sediment microtopography predicted localised benthic ecosystem functioning [1]. In a later field study using structure-from-motion (SfM) photogrammetry, they found that it explained a substantial proportion of the variance in biodiversity and ecosystem functioning across subtidal habitats [6]. Microtopography is summarised by areal roughness parameters. Of these, rugosity is the ratio of the true surface area to its projected area, and is closely related to the developed interfacial area ratio defined in ISO 25178-2 [7].

Measuring microtopography in the field is not straightforward. Traditional sampling methods are invasive and damage the fragile structures being measured. Non-contact alternatives such as LiDAR and SfM photogrammetry can be expensive and cumbersome to process [8]. Azhar et al. [8] addressed this with a low-cost framework. A consumer RGB-D camera, the Asus Xtion Pro Live, captured depth images of intertidal sediment, from which roughness parameters were computed after detrending. Low-cost RGB-D sensors have similarly been used to quantify habitat complexity [9]. The Xtion relies on structured light, which is degraded by strong ambient light [10]. It therefore had to be operated inside an opaque enclosure, which limited the area of each scan.

1.2 Problem statement

The overarching project adapted the framework of Azhar et al. [8] to an Intel RealSense D456, mounted on a portable aluminium rig above the sediment. The D456 is an active stereo camera [11]. It computes depth on the camera by matching texture between two infrared imagers, and a projected infrared dot pattern adds texture where the surface has little of its own [12, 13]. Unlike a structured-light camera, it can therefore still recover depth from naturally textured surfaces when ambient light washes out the projected pattern. As inherited, however, the system had three shortcomings.

First, its configuration had not been systematically evaluated. This covers the working distance, viewing angle, lighting, stream resolution, depth unit, visual preset and post-processing filters. Published evaluations of depth cameras [10, 14, 15] consider other camera models and mostly indoor scenes of planar or manufactured objects. None characterises the D456 on sediment-like relief at the sub-metre working distances of the rig.

Second, each surface was scanned at least four times from different mount rotations, but each scan was analysed individually. A single scan is noisy and has viewpoint-dependent errors. It also contains holes wherever relief hides the surface from the camera. The repeated scans were not combined to reduce these effects.

This project therefore characterises the D456 for close-range scanning of sediment-like surfaces. It also develops a processing pipeline that fuses multi-view scans into a single reconstruction, from which roughness parameters are computed and validated.

1.3 Project objectives

The objectives of this project were to:

- Characterise the D456 as functions of working distance, camera tilt and ambient lighting, using planar targets.
- Characterise the D456 stream settings (resolution, depth unit and visual preset) and post-processing filter settings, by evaluating scans of 3D-printed reference objects of known geometry across the different parameters.
- Develop a pipeline that automatically crops each scan to the analysis region marked by AprilTags, registers scans taken at different tilts and rotations, and fuses them into a single surface. Characterise how reconstruction accuracy depends on the combination of views.
- Develop a roughness estimation pipeline that levels, detrends and grids the fused surface and computes standardised areal roughness parameters, including rugosity. Validate it against the reference objects and against laboratory sand.

1.4 Scope, limitations and assumptions

Scope. The project was carried out over 12 weeks, with final submission on 27 October 2026. It was limited to the existing hardware, and no other cameras were evaluated. The hardware comprised:

- A single Intel RealSense D456 on a 1 m aluminium-extrusion frame with a height-adjustable cross member.
- A manually rotated camera mount, turned in 45° increments, on a servo-driven tilt stage.
- Four adjustable AprilTag boundary markers.

- An overhead shading tarp.

Testing used planar targets, three 3D-printed reference objects, and sand collected from the Langebaan Lagoon. The sand was scanned in the laboratory in dry, wet and waterlogged states. Field deployment on the intertidal flats was not possible within the time available.

Limitations. The characterisation and the chosen configuration target smooth, gently undulating relief of a few centimetres, representative of sand prawn mounds. They may not transfer to objects with sharp edges. The fused surface is represented as a single-valued height field, so overhangs such as burrow walls cannot be represented. It is also smoothed with a 10 mm moving least squares radius by default, so relief at shorter wavelengths is attenuated. Each characterisation condition was recorded once, so the repeatability of the set-up was not quantified. One of the proposed processing steps was automated feature detection, however due to the limited time, this was removed from the scope and is discussed as further work in chapter 10.

Assumptions. The following assumptions were made:

- The factory calibration of the D456 (camera intrinsics and depth-to-colour extrinsics) is accurate, and the camera was not recalibrated.
- The surface does not change during a capture session, and the AprilTags remain fixed relative to it, so every scan in a session observes the same region.
- The sediment within the analysis region is approximately planar at the scale of the region, with relief superimposed on it.
- The camera is stationary while each scan is captured, so its accelerometer measures gravity alone.
- The 3D-printed reference objects match their CAD models closely enough to serve as ground truth. These reference objects were printed using a Bambu lab X1 carbon.

1.5 Plan of development

Chapter 2 reviews the literature on sand prawn bioturbation, microtopography as a proxy for benthic activity, and depth sensing for surface measurement. Chapter 3 presents the theory of stereo depth capture, registration, fusion, evaluation and roughness estimation. Chapter 4 defines the system requirements and describes the rig and software architecture. Chapter 5 characterises the D456 camera across different conditions. Chapter 6 describes the cropping, fusion and roughness estimation pipeline. Chapter 7 validates the complete pipeline on reference objects and laboratory sand. Chapter 8 provides a summary of the results found in the project, chapter 9 draws conclusions against the objectives, and chapter 10 outlines further work.

Chapter 2

Literature review

This chapter reviews the literature behind the project and identifies the gap it addresses. It covers why sand prawn bioturbation matters (section 2.1), whether surface microtopography is a valid proxy for it (section 2.2), how sediment surfaces can be measured (section 2.3), how depth cameras have been characterised (section 2.4) and how multiple depth scans have been fused (section 2.5). section 2.6 summarises the resulting research gaps.

2.1 Sand prawn bioturbation and its ecological significance

Bioturbation is the reworking of sediment particles and the ventilation of burrows by animals [3]. Burrowing thalassinidean shrimps are among the most intense bioturbators, and Pillay and Branch [4] describe them as ecosystem engineers. Excavating burrows and ejecting sediment alters grain size, organic content, porewater chemistry and surface topography. At the surface this appears as mounds, depressions and burrow openings (Figure 2.1). Reported mounds are about 3 to 5.4 cm high and 12 to 15 cm wide [4], so a measurement system must resolve centimetre-scale heights over decimetre-scale wavelengths.

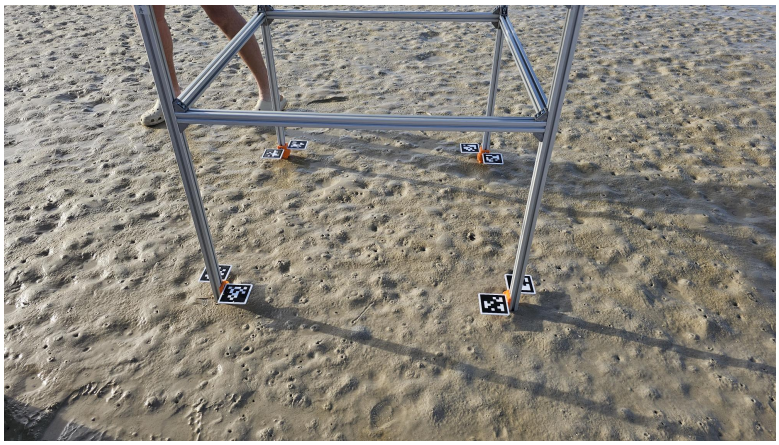

Figure 2.1: Surface Mounds and Burrow Openings Produced by *Callichirus kraussi* on an Intertidal Sandflat. Taken on 14/07/2026 During a Field Trip

On the South African coast the dominant species is the sand prawn *Callichirus kraussi* (*Callianassa kraussi* in older literature), which occupies most of the intertidal area of Langebaan Lagoon [16]. Its effects are substantial:

- sediment turnover of about $12\,\mathrm{kg\,m^{-2}}$ of wet sediment per day, burying benthic microalgae and altering bacterial and meiofaunal communities [17];
- increased abundance of surface grazers and of suspension- and deposit-feeding polychaetes and bivalves when prawns were excluded [18];
- reductions in phytoplankton biomass of about 50 % in experiments and 70 % in prawn-dominated estuarine areas, through pumping water across the burrow walls, which helps buffer coastal ecosystems against eutrophication [5].

Measuring this activity usually relies on sediment cores, which are destructive and slow, or on counts of burrow openings, which are laborious and capture only one surface feature. Surface structure also matters in its own right, as homogenising it reduces soft-sediment biodiversity [19]. Surface relief is therefore both a product of bioturbation and a habitat variable, which raises the question of whether it can be measured in place of the animals.

2.2 Sediment microtopography as a proxy for benthic activity

The case for microtopography as a proxy rests mainly on three studies by Schenone and colleagues:

- in a multi-site field study, biogenic surface features predicted ecosystem functioning better than species density [20];
- in flow-through laboratory tanks (Figure 2.2a), rugosity from a single 2D profile imaged through the tank wall explained 35 to 61 % of the variability in oxygen and nutrient fluxes [1];
- in the field, microtopography reconstructed with structure-from-motion (SfM) explained more than half of the variance in environmental variables, benthic fluxes, biodiversity and functional traits [6].

These studies were subtidal, in New Zealand, with different species, and the relationships are correlative. The link has therefore not been shown for *C. kraussi* on exposed intertidal flats. Testing it at Langebaan is the purpose of the wider study, and this project must provide measurements accurate enough for that test.

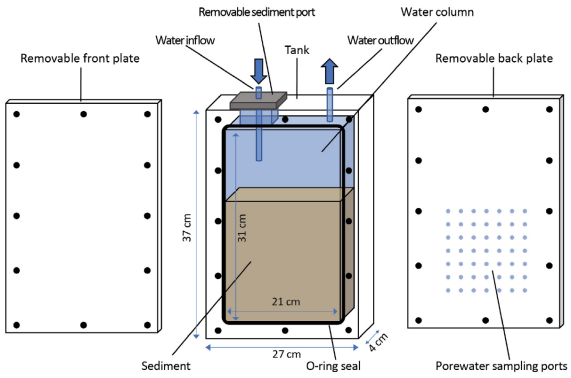

(a) Flow-through incubation tank [1].

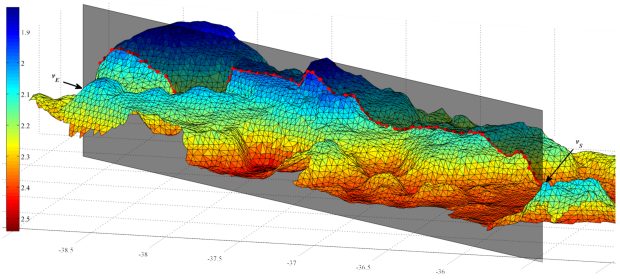

(b) Rugosity as true surface area over area projected onto a plane of best fit [21].

Figure 2.2: Measuring microtopography as an ecological proxy.

How roughness is quantified matters as much as whether it is measured. Traditional chain-and-tape rugosity compares the length of a chain laid along the bed with its straight-line length. It samples a single profile and overestimates rugosity when the chain wanders sideways [21]. Friedman et al. [21] instead computed rugosity from stereo meshes relative to a plane of best fit (Figure 2.2b), which decouples rugosity from slope, and SfM-derived rugosity has proved more powerful than chain-and-tape [22]. ISO 25178-2 standardises these quantities as areal parameters [7, 23], where the developed interfacial area ratio S_{dr} is rugosity expressed as a percentage excess area (chapter 3). Azhar et al. [8] used a comparable set of nine parameters from detrended depth images.

These parameters depend on how the surface is sampled. They vary with scan size and pixel resolution [24], and slope-based parameters such as S_{dq}, S_{dr} and rugosity are dominated by the shortest resolved wavelengths, where sensor noise and smoothing also act. Rugosity from a depth camera is therefore only interpretable when the noise floor and effective resolution of the measurement are known.

2.3 Non-contact measurement of sediment surfaces

A system for this application must resolve millimetre-scale relief over patches of a few tenths of a square metre, be portable and fast enough for a low-tide window, leave the surface undisturbed, and work in daylight on dry, wet and waterlogged sand. Coring and contact profiling damage the structures being measured [8], so only non-contact methods are considered (Table 2.1).

Table 2.1: Non-contact methods for measuring sediment surfaces.

Method	Strengths	Limitations
Terrestrial laser scanning	Accurate; return intensity calibrated to sand moisture [2]	Expensive; cumbersome processing [8]
SfM photogrammetry	Cheap hardware; very fine resolution [22]	Many images, scale reference and offline processing
Structured light	Immediate metric depth; ≤1 mm at 0.5 m [9, 25]	Pattern washed out by sunlight, needs shading [8]
Passive stereo	Tolerates sunlight	Needs a GPU host [15]; fails on textureless surfaces
Active stereo	Tolerates sunlight on textured surfaces; depth computed on board [12]	Field evidence only for other models [26]

Azhar et al. [8], the framework this project builds on, imaged 0.3 m^2 patches of intertidal sediment at seven sites with an Asus Xtion Pro Live inside a semi-opaque container (Figure 2.3). Of four detrending models tested on simulated surfaces, linear and quadratic performed best and agreed strongly on field data. The container was a necessity: a structured-light camera triangulates against its own projected pattern, which sunlight washes out, so the shading fixed both the size of each scan and its single viewpoint.

(a) RGB-D device and container in the field [8].

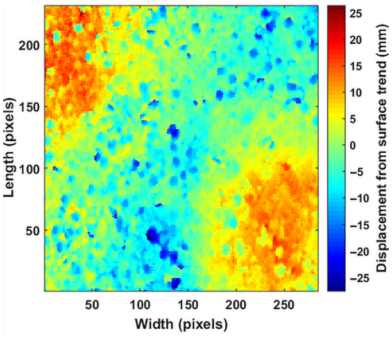

(b) Captured and detrended sediment surface [8].

Figure 2.3: The RGB-D framework of Azhar et al. [8].

Sunlight tolerance therefore decides the choice of sensor. An active stereo camera matches texture between two infrared imagers, and its projected dots only supplement the scene texture, so on a textured surface in strong light it degrades towards passive stereo rather than failing [12, 13]. Outdoors, Vit and Shani [26] found the active stereo D435 the most viable of four close-range sensors. The D456 used here is the IP65-rated variant of the D455 [11], which also suits wet intertidal work.

2.4 Characterisation of active stereo depth cameras

Intel's documents describe the D400 correlation algorithm and its expected performance [12], the choice of resolution, preset, depth unit and laser power [27], the post-processing filters [28], external projectors [13] and the plane-fit test methodology used in this project [29]. None evaluates these settings against ground-truth relief at the millimetre scale relevant to sediment.

Independent studies are summarised in Table 2.2. The closest precedent is Servi et al. [30], who applied ISO 10360-13 to the D415, D455 and L515 and found the D455 best in the standardised tests. The D456 shares the D455 optical module [11], but that study was confined to manufactured artefacts in the laboratory. Results for the D435, which Rustler et al. [15] found best within 1 m, do not transfer, because its 50 mm baseline differs from the 95 mm of the D455 module. Nguyen et al. [31] showed that noise models from distance and angle sweeps improve fusion, and the sweeps in this project follow the same logic.

Table 2.2: Published depth-camera characterisations relevant to this project. SL: structured light; AS: active stereo; PS: passive stereo; ToF: time of flight.

Study	Sensors (technology)	Targets	Environment
Halmetschlager-Funek et al. [10]	Ten sensors (SL, AS, ToF)	Planar targets, several materials	Indoor
Carfagni et al. [14]	D415 (AS)	VDI/VDE 2634 artefacts, 0.1–1 m	Laboratory
Servi et al. [30]	D415, D455 (AS); L515 (ToF lidar)	ISO 10360-13 artefacts, flat and 3D objects, 0.1–1.5 m	Laboratory
Rustler et al. [15]	D435, D455, OAK-D Pro (AS); ZED 2 (PS)	Plane, doll, household objects, up to 4 m	Indoor
Vit and Shani [26]	Kinect v2 (ToF), Astra S (SL), SR300 (SL), D435 (AS)	Maize plants, close range	Outdoor, varied light
Azhar et al. [8]	Xtion Pro Live (SL)	Plane; intertidal sediment	Outdoor, shaded
This work	D456 (AS)	Planes, 3D-printed relief, sand, 0.6–0.93 m	Indoor, outdoor

No study was found that characterises any active stereo camera on wet or waterlogged sediment. Nolet et al. [2] found that sand reflectance falls non-linearly with moisture from 350 to 2500 nm (Figure 2.4), so wet sand is darker in the near-infrared and offers less pattern return and contrast for matching. Standing water adds specular reflection and refraction [25], and some sensors fail on reflective materials [10]. The behaviour of the D456 on wet sand must therefore be established experimentally.

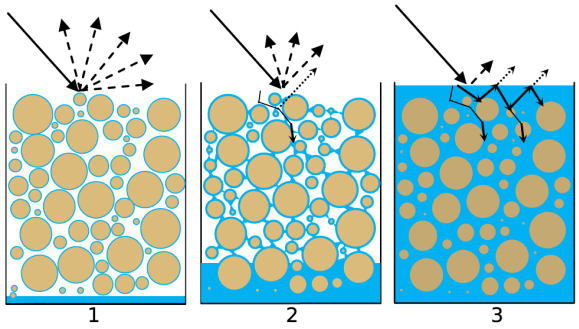

Figure 2.4: Reflectance of sand decreasing non-linearly as moisture content increases [2].

2.5 Registration and fusion of multi-view depth scans

Fusing aligned range images into a weighted signed distance function averages noise, fills gaps and weights measurements by reliability [32]. KinectFusion applied this in real time with a consumer depth camera and showed that the fused model contains detail absent from any single frame (Figure 2.5) [33]. These methods were developed for objects and rooms, and judged by geometric accuracy rather than by roughness parameters.

Two results bear on sediment. First, nearly flat surfaces are degenerate for registration, since ICP cannot constrain sliding within a plane or rotation about its normal [34]. A near-planar bed therefore needs external constraints such as AprilTags [35, 36] or inertial measurements. Second, benthic fusion is normally done at survey scale using vehicle navigation [21], not by fusing repeated close-range views of one patch. Azhar et al. [8] used a single viewpoint, and the inherited system analysed each repeated scan individually. No study was found

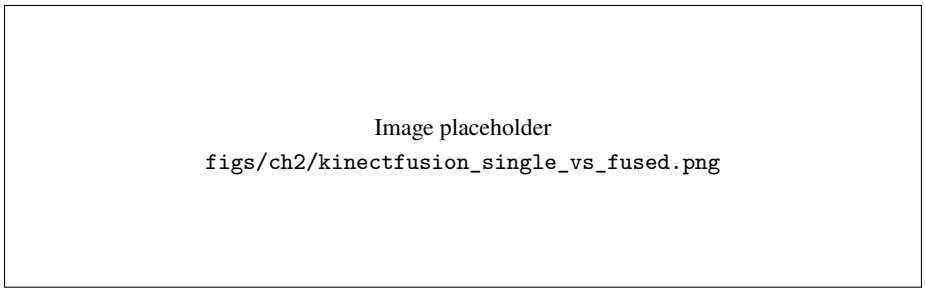

Figure 2.5: A single noisy depth frame compared with the surface fused from many frames [33].

that fuses multi-view scans of a sediment patch and quantifies how the choice of views affects reconstruction accuracy and roughness.

2.6 Summary and research gap

Callichirus kraussi is a dominant ecosystem engineer of the Langebaan intertidal, and sediment microtopography is a credible proxy for its activity. Roughness parameters, however, depend on the noise and resolution of the measurement. Active stereo is the best-supported low-cost sensor for outdoor use, but it has not been tested in this setting. Four gaps remain, each addressed by one objective in section 1.3 (Table 2.3).

Table 2.3: Research gaps and the objectives that address them.

Gap	Objective
The D455 module has been characterised only in the laboratory [30]. Its noise, bias and fill rate at the rig's distances, tilts and outdoor lighting are unknown.	1: Planar characterisation
Manufacturer guidance on resolution, depth unit, preset and filters is not validated against relief of the scale of sand prawn mounds.	2: Reference-object characterisation
Multi-view fusion has not been applied to close-range sediment scans, and near-flat beds defeat purely geometric registration.	3: Cropping and fusion pipeline
Rugosity from RGB-D data has not been validated against surfaces of known geometry.	4: Roughness pipeline and validation

Chapter 3

Theoretical Background

This chapter outlines the theory behind the project, seperated by each stage: depth measurement, registration, fusion, evaluation and roughness estimation. The main aim of this section is to provide a high level description of some of the fundamentals. Implementation details and parameter choices are given in chapter 4 and chapter 6.

3.1 Depth Capture

Seeing Double Vision Can Be Good

The Intel RealSense D456 is an IP65-rated version of the D455, built from the same D450 optical module and D4 vision processor [11]. It contains two global-shutter imagers separated by a baseline of 95 mm, an infrared (IR) laser projector, a global-shutter RGB camera and an inertial measurement unit (IMU). The Camera features the D4 processor, which computes depth on the camera, so the host receives depth images rather than stereo image pairs.

3.1.1 Camera Model

Each imager is modelled as a pinhole camera [37]. A point $\mathbf{X}_w$, in a world frame is mapped to the pixel (u, v) by

$$\lambda\,(u,\,v,\,1)^\top = \mathbf{K}\,[\mathbf{R}\mid\mathbf{t}]\,(\mathbf{X}_w^\top,\,1)^\top, \qquad \mathbf{K} = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}, \tag{3.1}$$

where $\mathbf{X}_w = (X_w, Y_w, Z_w)^\top$ is the point in the world frame, and $(X, Y, Z)^\top = \mathbf{R}\mathbf{X}_w + \mathbf{t}$ is the same point in the camera frame, whose depth Z equals λ. The intrinsic matrix $\mathbf{K}$ holds the focal lengths f_x,f_y in pixels and the principal point (c_x, c_y). The extrinsic parameters $\mathbf{R}$ and $\mathbf{t}$ give the pose of the camera relative to another frame, such as the world or a second sensor. Real lenses add radial and tangential distortion, and calibration estimates all of these parameters, typically from images of a known planar target [38]. The D456 is factory calibrated, and because its stereo images are rectified on the camera, the depth stream is described by a distortion-free pinhole model while the colour stream has its own distortion coefficients.

Inverting Equation 3.1 for a pixel with measured depth Z gives the corresponding point in the camera frame,

$$X = \frac{(u - c_x)\,Z}{f_x}, \qquad Y = \frac{(v - c_y)\,Z}{f_y}, \tag{3.2}$$

which is how a depth image becomes a point cloud. Depth is aligned to the colour image by back-projecting each pixel, transforming it with the depth-to-colour extrinsics and projecting it with the colour intrinsics.

3.1.2 Stereo Triangulation

After rectification the image rows of the two cameras are aligned, so a scene point appears on the same row in both images and the search for its match is one-dimensional [37]. In the geometry of Figure 3.1, a point at depth Z projects to horizontal image coordinates x_L and x_R, measured from each principal point, and similar triangles give

$$Z = \frac{fB}{d}, \qquad d = x_L - x_R, \tag{3.3}$$

where f is the focal length in pixels, B the baseline and d the disparity. The camera finds d for each pixel by comparing small image patches along the row, and interpolates the best match to a fraction of a pixel [12].

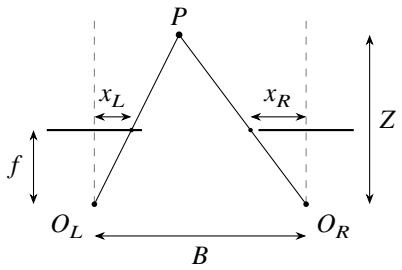

Figure 3.1: Rectified stereo geometry: camera centres O_L, O_R, baseline B, focal length f and the projections x_L, x_R of a point P at depth Z.

Differentiating Equation 3.3 gives the depth uncertainty caused by an uncertainty σ_d in disparity,

$$\sigma_Z = \frac{Z^2}{fB}\,\sigma_d. \tag{3.4}$$

For a fixed matching error, depth noise grows with the square of the range and falls with a longer baseline or focal length [12]. Since f in pixels is proportional to the horizontal resolution, a lower resolution also increases the noise. A constant σ_d is optimistic: it depends on the contrast of the matched texture, which falls as the projected pattern dims with distance or is washed out by ambient light.

The disparity search also has a finite range. The D4 processor searches 126 disparities [39], so the minimum measurable depth is $Z_{\min} = \frac{fB}{126}$. At 1280×720, $f \approx \frac{640}{\tan(87^\circ/2)} \approx 674\,\text{px}$, giving $Z_{\min} \approx 0.51\,\text{m}$, consistent with the specified 0.52 m [11]. Equation 3.4 and $Z_{\min}$ together set the working-distance trade-off characterised in chapter 5. the camera should be as close as $Z_{\min}$ and the required field of view allow.

3.1.3 Passive Stereo, Structured Light and Active Stereo

Passive stereo matches natural texture between two cameras, and fails where the surface is uniform because many patches along a row look alike. Structured light, used by the Asus Xtion in the system of [8], projects a known pattern and triangulates between the projector and a single camera, so the pattern itself must be known and calibrated. Active stereo, the principle of the D456, combines a stereo pair with a projector of a semi-random dot pattern, which can be seen in chapter B, [12, 13]. Matching is still performed between the two cameras, so the pattern does not need to be know, the main purpose is to add aditional texture onto the surface. On a featureless surface, such as the white board of chapter 5, depth therefore depends entirely on the dots and fails when sunlight, whose near-infrared component can outshine them, washes them out. On a naturally textured surface in sunlight the camera behaves as a passive stereo camera and still returns dense depth. Both behaviours were observed in chapter 5. The RGB camera has an IR-cut filter, so the dots do not appear in the colour image.

3.2 Registration

Each scan is expressed in the frame of the camera at the moment of capture, so the same point on the sediment has different coordinates in scans taken at different rotations and tilts. Registration finds the rigid transformation that maps one scan into the frame of another, and is needed both for fusion (section 6.2) and for comparison with a reference model (section 4.5).

3.2.1 Rigid Transformations and Closed-Form Alignment

A rigid transformation maps a point $\mathbf{p}$ to $\mathbf{R}\mathbf{p} + \mathbf{t}$, where $\mathbf{R}$ is a rotation matrix ($\mathbf{R}^\top\mathbf{R} = \mathbf{I}$, $\det\mathbf{R} = +1$) and $\mathbf{t}$ a translation, giving six degrees of freedom. If N corresponding point pairs $(\mathbf{p}_i, \mathbf{q}_i)$ are known, the transformation that minimises $\sum_i \|\mathbf{R}\mathbf{p}_i + \mathbf{t} - \mathbf{q}_i\|^2$ has a closed-form solution [40, 41]. With centroids $\bar{\mathbf{p}}$ and $\bar{\mathbf{q}}$, and the singular value decomposition $\mathbf{H} = \sum_i (\mathbf{p}_i - \bar{\mathbf{p}})(\mathbf{q}_i - \bar{\mathbf{q}})^\top = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$,

$$\mathbf{R} = \mathbf{V}\,\mathrm{diag}(1, 1, \det(\mathbf{V}\mathbf{U}^\top))\,\mathbf{U}^\top, \qquad \mathbf{t} = \bar{\mathbf{q}} - \mathbf{R}\bar{\mathbf{p}}. \tag{3.5}$$

The determinant term ensures a proper rotation rather than a reflection when the data are noisy [41]. At least three non-collinear pairs are needed because if the points lie close to a line, rotation about that line is poorly determined, which is why the tag-based alignment in section 6.2 takes the tilt from the IMU when the visible tags are nearly collinear. A stationary accelerometer measures only gravity, which fixes the two tilt angles but not the rotation about the vertical axis, which can be determined from tag identification or surface geometry.

3.2.2 Fiducial Markers

AprilTags are square black-and-white markers that encode an identifier in a binary grid. The codes are chosen to keep a large minimum Hamming distance between all identifiers in all four rotations, so a tag is rarely misidentified even when some bits are read incorrectly [35]. The detector fits quadrilaterals to the boundaries between light and dark regions, decodes the bits inside each candidate, and returns each tag's identifier and the sub-pixel positions of its four corners [36].

The corners of a planar tag of known size are related to their image positions by a homography $\mathbf{H} \simeq \mathbf{K}[\mathbf{r}_1\ \mathbf{r}_2\ \mathbf{t}]$, where $\mathbf{r}_1$ and $\mathbf{r}_2$ are the first two columns of the tag's rotation matrix [37, 38]. Given $\mathbf{K}$, the pose follows from $\mathbf{K}^{-1}\mathbf{H}$ with $\mathbf{r}_3 = \mathbf{r}_1 \times \mathbf{r}_2$. This monocular estimate rests on four closely spaced points and is sensitive to corner noise. For a small or nearly fronto-parallel tag, two different orientations can fit the corners almost equally well [42]. The pipeline instead uses the depth data, intersecting the corner pixel rays with a plane fitted to the measured depth of each tag (section 6.1). This gives 3D corners in every scan, so aligning two scans becomes a 3D-to-3D problem solved with Equation 3.5.

3.2.3 Feature-Based Coarse Alignment

When no prior pose is available, correspondences must be found from the shape of the surface. Fast Point Feature Histograms (FPFH) describe the geometry around each point by the angles between its surface normal and those of its neighbours, expressed in a local coordinate frame. Three angular features are each binned into 11 bins and concatenated into a 33-element histogram [43]. Because only relative angles are used, the descriptor does not change under rigid motion. Candidate correspondences are formed by pairing points with similar descriptors, and many of these pairs are wrong.

Random Sample Consensus (RANSAC) estimates the transformation in the presence of these outliers [44]. It repeatedly selects three correspondences at random, computes their transformation with Equation 3.5, and counts the inliers (correspondences brought within a distance threshold), keeping the transformation with the most inliers. If a fraction w of the correspondences is correct, drawing at least one all-inlier sample with probability p takes $N = \log(1-p)\,/\,\log(1-w^3)$ trials, or about 4600 for $w = 0.1$ and $p = 0.99$. On flat or smoothly undulating

sediment most points have nearly identical descriptors, so w is small and an incorrect transformation can gather as many inliers as the correct one. This is why feature matching is only a fallback in this pipeline (section 6.2).

3.2.4 Iterative Closest Point

Iterative Closest Point (ICP) refines an approximate alignment by alternately pairing each source point with its nearest target point and solving for the transformation that minimises the distance between the pairs. The point-to-point form [45] minimises Euclidean distance and is solved with Equation 3.5. The point-to-plane form [46] minimises the distance from each source point to the tangent plane at its paired target point,

$$E(\mathbf{R}, \mathbf{t}) = \sum_i \Big((\mathbf{R}\mathbf{p}_i + \mathbf{t} - \mathbf{q}_i) \cdot \mathbf{n}_i \Big)^2, \tag{3.6}$$

where $\mathbf{n}_i$ is the target normal. This has no closed-form solution, but for the small rotations of a refinement step it can be linearised and solved as a six-parameter linear least-squares problem [47].

Two scans never sample exactly the same points. Point-to-point ICP penalises the resulting tangential offsets, which slows convergence, whereas point-to-plane ICP lets points slide along the surface and on smooth surfaces usually converges in far fewer iterations and to a more accurate pose, at the cost of requiring normals [48]. Both forms are local, descending to the nearest minimum of the cost, so the initial pose must already be close, and both are limited by the geometry of the surface rather than by the error metric. On a plane, translation within the plane and rotation about its normal leave the point-to-plane cost unchanged and change the point-to-point cost only slightly, so these motions are effectively unconstrained [34]. A nearly flat sediment bed is close to this degenerate case, so its in-plane pose must be fixed by the tags or the coarse stages, with ICP refining the remaining degrees of freedom.

3.3 Fusion

Putting Bits Together

Volumetric methods store a weighted running average of the signed distance to the surface in a voxel grid and extract the surface where it is zero [32, 33]. For a nearly flat bed, a 2.5D height field $z(x, y)$ on a grid aligned with the mean plane is a simpler equivalent in which each cell collects the heights from all scans (section 6.2). One consideration with this medthod, is that it cannot represent overhangs and sharp features, suuch as burrow walls.

First all scans points are weighted. When measurements have independent errors with different variances σ_k^2, the combination with the lowest variance is the weighted mean with weights $w_k \propto 1/\sigma_k^2$. Depth error grows with range (Equation 3.4) and with obliquity, therefore a surface seen at a grazing angle is foreshortened and its depth varies across each matching window, and the axial noise of depth sensors has been observed to rise sharply as the angle between the surface normal and the viewing direction approaches 90° [31]. Measurements are therefore down-weighted with both, in the spirit of the angle-based weights of [32].

A misregistered scan or stray point produces errors that are neither small nor normally distributed. One such value can move the mean arbitrarily far, while the median tolerates up to half the data being bad but discards much of the information in good data. The trimmed mean, which discards a fixed fraction of the lowest and highest values (by cumulative weight, in a weighted version) before averaging, lies between the two [49].

Moving Least Squares (MLS) fits a low-order polynomial surface to the neighbours of each point within a radius h, with weights that decrease with distance, and projects the point onto it [50, 51]. This removes noise and residual layer offsets on scales below h, but it is a low-pass filter that also attenuates relief with wavelengths comparable to or shorter than h, so h sets a lower bound on the scale of roughness the fused surface can represent.

3.4 Evaluation

Are We Winning?

Each reconstruction is first registered to the reference model (section 4.5). For each reconstructed point $\mathbf{s}_i$ lying over the part, d_i is its signed distance to the closest point on the reference surface, positive on the side of the outward surface normal. Points on the table or other surroundings are excluded from the accuracy statistics. Distances are computed to the reference triangles rather than to its vertices, which avoids any dependence on how densely the reference is sampled.

Over N points the bias μ, standard deviation σ and root-mean-square error are

$$\mu = \frac{1}{N}\sum_i d_i, \qquad \sigma^2 = \frac{1}{N}\sum_i (d_i - \mu)^2, \qquad \mathrm{RMSE} = \sqrt{\frac{1}{N}\sum_i d_i^2} = \sqrt{\mu^2 + \sigma^2}. \tag{3.7}$$

The bias measures systematic offset, σ random spread, and the RMSE combines both. The RMSE is not the average distance, since squaring weights large errors more heavily. The mean absolute error $\frac{1}{N}\sum_i |d_i|$ is less affected by outliers, and percentiles of $|d_i|$ and a trimmed RMSE (worst 1 % and 5 % of points removed) are also reported. Because registration chooses the pose that best fits the reference, part of any systematic error, such as a small tilt, is absorbed into the pose, so the metrics measure shape error after the best rigid fit rather than absolute accuracy.

Accuracy metrics are computed only over points that exist, so a reconstruction that covers only the easy parts of an object can still score well. Following [52], the precision $P(\tau)$ is the fraction of reconstructed points within a distance τ of the reference. For the completeness, the visible reference surface is sampled uniformly by area. A sample is *covered* if a reconstructed point lies within a coverage radius $r_c = 10\,\mathrm{mm}$ of it, and the covered set is denoted $\mathcal{M}_c$. The completeness $R(\tau)$ is the fraction of $\mathcal{M}_c$ within τ of the reconstruction, so a reconstruction of half the object is judged on the half it saw, and the covered fraction $|\mathcal{M}_c|/|\mathcal{M}|$ is reported separately. The F-score, $F = 2PR/(P+R)$, is high only when both are high. Thresholds of $\tau = 0.5, 1, 2$ and $5\,\mathrm{mm}$ are used.

With $d(\mathbf{a}, \mathcal{B})$ the distance from a point $\mathbf{a}$ to the nearest point of a set $\mathcal{B}$, the Chamfer distance averages the mean nearest-neighbour distances in both directions between the reconstruction $\mathcal{S}$ and the covered reference samples $\mathcal{M}_c$, while the Hausdorff distance takes the worst case [53]:

$$d_{\mathrm{CD}} = \frac{1}{2}\left[\frac{1}{|\mathcal{S}|}\sum_{\mathbf{s}\in\mathcal{S}} d(\mathbf{s},\mathcal{M}) + \frac{1}{|\mathcal{M}_c|}\sum_{\mathbf{m}\in\mathcal{M}_c} d(\mathbf{m},\mathcal{S})\right], \qquad d_{\mathrm{H}} = \max\Big\{\max_{\mathbf{s}\in\mathcal{S}} d(\mathbf{s},\mathcal{M}),\ \max_{\mathbf{m}\in\mathcal{M}_c} d(\mathbf{m},\mathcal{S})\Big\}. \tag{3.8}$$

Here $d(\mathbf{s}, \mathcal{M})$ is the exact point-to-surface distance used above, and unsquared distances are used throughout. The second Chamfer term penalises missing regions within the covered area. Because $\mathcal{M}_c$ excludes reference surface farther than r_c from the reconstruction, d_{H} is a covered-region Hausdorff distance: it does not capture large unscanned regions, which are reported through the covered fraction instead. The Hausdorff distance is set by the single worst point and is very sensitive to outliers, so a 95th-percentile version, $\max\{P_{95}(d(\mathbf{s},\mathcal{M})),\ P_{95}(d(\mathbf{m},\mathcal{S}))\}$, is also reported.

3.5 Roughness Calculation

Rough Times Ahead

A measured surface combines form, waviness and roughness at progressively shorter wavelengths. ISO 25178-2 defines areal texture parameters on a scale-limited surface, obtained by removing the form and restricting the wavelength band with filters [7, 23]. Here the form is removed by subtracting a fitted low-order polynomial $p(x,y) = \sum_{i+j\le n} a_{ij}\, x_n^i\, y_n^j$, and the band is set implicitly, at the short end by the grid spacing and MLS radius, and at the long end by the polynomial degree and patch size (section 6.3). A higher degree removes more long-wavelength relief, so the degree decides what counts as roughness. [8] found a planar fit the most robust for sediment in the intertidal region.

The parameters, defined in Section section 6.3, fall into three groups. Height parameters (S_a, S_q, S_{sk}, S_{ku}) describe the distribution of residual heights but not their arrangement, so two surfaces can share the same S_q yet have very different feature sizes. Spatial parameters (S_{al}, S_{tr}) describe feature size and directionality. Hybrid parameters combine height and spacing through the slope: S_{dq} is the root-mean-square gradient and S_{dr} the developed interfacial area ratio, the percentage by which the true surface area exceeds the projected area. Rugosity, as used in benthic ecology, is the same quantity expressed as the ratio $A_{\text{true}}/A_{\text{proj}} = 1 + S_{dr}/100$. For gentle slopes the area element $\sqrt{1 + |\nabla z|^2} \approx 1 + \frac{1}{2}|\nabla z|^2$, so $S_{dr} \approx 50\, S_{dq}^2$ % and rugosity is essentially a measure of the mean squared slope.

Scale dependence. Slopes are estimated from height differences between neighbouring grid cells a distance g apart. Natural surfaces have relief over a wide range of scales, often approximately self-affine with a power-law spectrum, so S_{dq}, S_{dr} and rugosity depend on g [24]. Noise makes this worse since uncorrelated height noise σ_n adds a variance of $2\sigma_n^2/g^2$ to a forward-difference slope along each axis. For example, $\sigma_n = 0.5$ mm on a 2 mm grid adds a root-mean-square slope of about 0.35 per axis, which can exceed the true slope of a gently rippled bed. Slope-based parameters are therefore comparable only between surfaces processed with the same grid spacing and smoothing, and the noise contribution must be removed before rugosity is attributed to the sediment. Height parameters are much less sensitive to g.

Chapter 4

System Design

This chapter defines the general requirements of the system as a whole and describes the hardware rig used to capture scans, and the software pipeline built to capture, and process the data to provide useful metrics about the sediment surface. The theory in chapter 3 is applied here, and the sensor characterisation in chapter 5 uses the system as specified below.

4.1 Requirements Definition

The aim of the overarching project is to quantify sandprawn bioturbation on intertidal sediment, within the Langebaan Lagoon, by measuring surface microtopography and extracting roughness metrics and visible biogenic features (mounds and burrow openings). The initial stages of the project were adapted from the framework of Azhar et al. [8], primarily modified to use the Intel RealSense D456.

This project aims to provide an investigation into the performance of the RealSense D456 camera currently used. The central aim is to characterise the depth camera in order to determine the optimal hardware and software configuration, and then to use scans taken under those configurations to build a high-quality reconstruction of the sediment on which the analysis is performed. The key output is an analysis of depth noise and reconstruction accuracy under different conditions.

4.2 System Overview

The system has three main operational stages, shown in Figure 4.1, with a fourth evaluation stage used for characterisation.

1. *Capture*: The capture software rs_marker scans the surface at multiple rotations and camera tilts.
2. *Fusion*: multi-angle scans are cropped to the tag-defined region, aligned and merged into a single reconstruction.
3. *Microtopography processing*: the fused surface is detrended, and roughness metrics are extracted

For the purpose of characterisation, a fourth evaluation stage is used to gauge the reconstruction and metrics performance of the system against known target relief's.

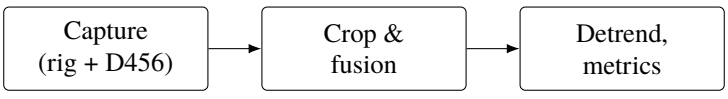

Figure 4.1: System overview. [TODO: Replace with a fuller diagram showing the data formats (.bag, .ply) passed between stages.]

Table 4.1: Hardware summary.

Component	Specification
Depth camera	Intel RealSense D456
Frame	Aluminium extrusion, $\approx$1 m^3, sliding H-frame
Tilt actuation	Servo, $\pm$60°, Arduino-controlled
Rotation	Manual, 45° increments, 360°
Boundary markers	Four adjustable AprilTags
Ambient light control	Overhead tarp

4.3 Hardware Rig

A summary of the hardware used for the capture rig is shown in Table 4.1.

4.3.1 Frame

The existing rig was built from 3030 aluminium extrusion in a 1 m × 1 m × 1 m frame. The top of the frame has a plastic tarp attatched, this allows for the shading of the scan area in the event of direct sunlight degrading the active stereo functionality of the D456. This frame was modified for the purpose of this project, in order to feature an H-shaped cross member that can slide vertically along the four legs. This enables the ability to adjust the camera height, and thus working distance, for purpose of characterisation and easy configuration. It was also modified to feature 3D printed April tag mounts that allow the position to be adjusted

4.3.2 Camera, Tilt and Rotation Mechanism

The Intel RealSense D456 is mounted on the H-frame through a 3D-printed bracket bolted to the extrusion. A servo motor in the bracket tilts the camera by up to $\pm$60°. The camera clips onto a rotating platform on the tilt stage and can be turned manually through a full 360° in 45° increments. Together, tilt and rotation allow each surface point to be observed from several viewing directions, which reduces occlusion within burrows and mounds. This multi-view scanning additionally averages out view-dependent depth error. The servo is driven by an Arduino microcontroller, connected to and commanded from the operator's laptop, which also runs the capture software.

4.3.3 AprilTag Boundary Markers

Four AprilTags are mounted on brackets on the frame legs, and each bracket can slide inward or outward. The tags' positions define a square region on the sediment surface to be processed and analysed. This gives a repeatable, physically marked analysis region, and its size can be changed to suit the desired scan region or height configuration of the camera. Because the tags are fixed to the frame, their detected positions also provide a geometric cue for aligning scans taken from different tilt and rotation settings.

4.4 Software Architecture and Processing Pipeline

4.4.1 Capture Software

Capture is implemented in C++ utilising the RealSense SDK. For each rig pose, it applies a configurable chain of depth post-processing filters and records the stream to a `.db3` file, together with IMU data. It detects the four

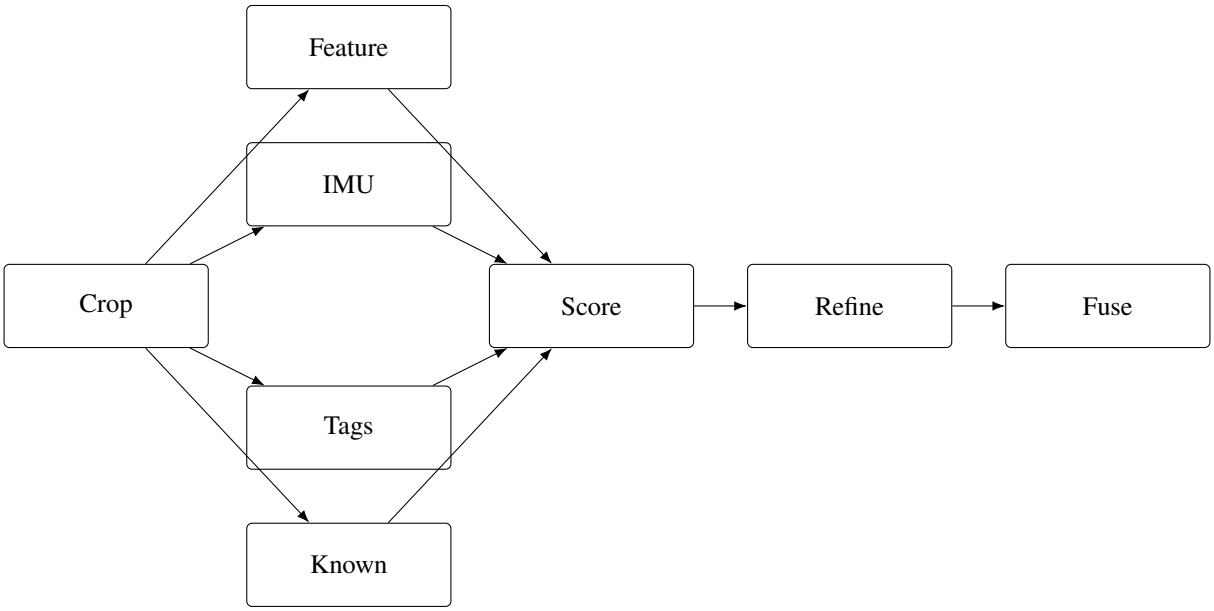

Figure 4.2: System overview. [TODO: Replace with a fuller diagram showing the data formats (.bag, .ply) passed between stages.]

AprilTags in the colour image, and the best frame of each pose is exported as a point cloud in PLY format. Tag detection is stored along side captures for use of later processing through cropping and alignment.

4.4.2 Cropping and Multi-Angle Fusion

Each scan is first cropped to the square defined by the four tags. The cropped scans, taken at different tilt and rotation settings, are then fused (**??**). Fusion needs a good initial pose before refinement, so the pipeline generates several *coarse* alignment hypotheses by independent methods:

1. from the detected tag locations;
2. from the known rotation and servo tilt angle set on the rig;
3. from the camera IMU orientation;
4. from RANSAC matching of FPFH features [**?** **?**].

Each hypothesis is scored with a common alignment-quality measure, and the highest-scoring one is selected. It is then refined with ICP [**?**] and the merged surface passes to the microtopography stage. Using four independent seeds makes the fusion robust to the failure mode of any single method (for example, feature matching on a nearly planar surface, or IMU drift).

4.4.3 Microtopography Processing

The fused surface is detrended with a robust polynomial fit, which removes the large-scale slope and leaves the microtopography. Roughness is then quantified with reference-free, ISO 25178-2-style areal height and area parameters, together with spectral and multi-scale descriptors, computed on the gridded residual height map. Mounds and burrow openings are segmented from the microtopography using thresholding and morphological operations.

4.5 Evaluation Framework and Ground Truth

For characterisation and validation, a separate Python tool, `evaluate_scans`, measures the geometric error of any scan, whether a single capture or the fused output, against a known reference. It reuses the structure of the fusion pipeline's alignment stage, but aligns each scan to the reference. The tool then evaluates scan accuracy and coverage against the reference object.

4.5.1 Reference Surfaces

The reference objects used for validation include a 3D printed egg carton surface $z = sin(x) + sin(y)$, at large and small scale, and a square based pyramid. These reference objects provide a balance of condition matching of the actual sediment surface, using the soft sinusoidal surface, and general performance to less natural sharp shapes. For the purpose of validation, the 3D model files (.stl) are used for evaluation as the 3D mesh surface allows for point to surface distance calculation across the entire 3D object.

4.5.2 Alignment and Metrics

Alignment uses several seeded feature-matching starts, systematic perturbations of the best pose (quarter turns, a flip, and slides of a quarter and half of the part), and point-to-plane ICP, followed by a final refinement against the exact surface. Candidate poses are ranked by a truncated least-squares cost computed in both directions (scan to reference and reference to scan), which prevents a pose that matches fewer points from winning by default. This matters on periodic targets such as the egg carton, where several look-alike poses exist. The tool flags a result as ambiguous when the two best poses are too close to call, and reports degrees of freedom the surface cannot constrain.

For each aligned scan the tool reports accuracy (RMSE, MAE, bias, standard deviation, trimmed statistics, percentiles), the vertical error for height-field references, and coverage of the reference. It also reports completeness, precision and F-score at several distance thresholds, Chamfer and Hausdorff distances, outlier fraction and point density. Every run appends a row to a shared table with the exact parameters, reference hash and random seed, so experiments are reproducible and comparable across configurations.

4.6 Summary

The rig provides repeatable, multi-angle capture of a user-defined surface region. The software fuses those scans, extracts microtopography, and is checked against analytic ground truth. chapter 5 uses this system to characterise the D456's depth noise.

Chapter 5

Intel RealSense Characterisation

The aim of the overarching project is to produce high-accuracy reconstructions of intertidal sediment from which surface roughness metrics can be extracted. The accuracy of those metrics is limited by the depth data they are computed from, so the Intel RealSense D456 was characterised before the capture and processing pipeline was finalised. The characterisation answers four practical questions about how the camera and rig should be operated in situ: at what working distance, at what viewing angle, under what ambient lighting, and with which depth-stream and filter settings the camera produces the most accurate surface.

The characterisation was performed in two stages. In the first stage the camera was tested against flat planar targets. A plane has an exact, easily measured ground truth, which allows the effect of working distance, tilt angle and ambient light to be isolated from the effect of surface shape. These tests give a baseline for sensor noise and provide the operating envelope used in the second stage. In the second stage the camera was tested against 3D-printed reference objects with known geometry (see Appendix C.) These tests identify the camera configuration, post-processing filter settings and combination of viewing directions that give the best reconstruction, using surface relief that is representative of the sediment of interest. Each characterisation below is presented with its method, results and a short discussion.

5.1 Plane Surface Characterisation

As Plain As That

5.1.1 Capture procedure and metrics

All planar characterisation used custom C++ software built on the Intel RealSense SDK, following the plane-fit metrics of Intel's Depth Quality Tool. Details of each trial, such as distance, tilt angle and a lux reading, are entered into the tool and the scan is started. A region of interest (ROI) is then applied whose pixel size is adjusted with distance so that it always covers a fixed 50×50 cm area of the target. A plane is fitted to the ROI points of every frame and compared with the ground-truth plane. Several metrics are reported by the tool, however there are four main metrics that will be discussed here to represent the general performance, each averaged over the valid frames of a trial.

- **Median depth error** (mm): the median signed difference between the measured and ground-truth depth, which measures systematic bias.
- **Depth RMSE** (mm): the root mean square representing depth noise from camera to plane.
- **Plane-fit RMS** (mm): the root-mean-square distance of the ROI points from the best-fit plane. This measures spatial noise and is independent of any constant offset in the depth.
- **Valid Pixels** (%): the proportion of ROI pixels that returned a valid depth value.

The distinction between precision and bias matters for this project. A smooth, systematic bias is removed when the fused surface is detrended to extract microtopography, whereas random noise adds directly to the measured roughness. Plane-fit RMS and fill rate are therefore treated as the primary metric, with depth error and depth RMSE treated as secondary indicators. Each condition was recorded as a single trial. The frames within a trial capture temporal variation but not the repeatability of re-setting the experiment, which is discussed in **??**.

5.1.2 Planar Characterisation Methodology

Working distance

The camera was mounted on a tripod and pointed at a white, featureless wall, as recommended by Intel for depth-quality testing. This forces the camera to depend on its own projected infrared pattern. The tripod was used instead of the experimental rig because the rig's mounting brackets could not hold the camera sufficiently parallel to the floor: the camera angle varied with every scan and consistently exceeded 1.5°, which would have contaminated a distance sweep. The camera was positioned at eight distances between 604 mm and 903 mm, and the angle between the camera and the wall was minimised by hand at each position. The fitted-plane angle confirmed that the residual pitch and roll stayed below 0.5°in all trials. The ground-truth distance was measured between the wall and the cameras glass front using a tape measure.

Camera tilt

The camera was mounted on the experimental rig so that its angle could be set with the tilt servo. A large flat white foam board was placed on the floor and re-centred in the field of view (FOV) after each change of angle. Two rotations were tested independently. The pitch axis is parallel to the line through the two IR imagers, and the roll axis is perpendicular to both the pitch axis and the optical axis. The angle was changed in steps of approximately 2°, covering pitch angles from $-27°$ to $-2°$ and roll angles from $+3°$ to $+25°$ for 26 trials in total. The camera mount height was kept fixed. For each trial the vertical height of the left IR imager, which is the centre of the resulting depth image, above the board was measured, and the axial ground-truth distance was calculated from that height and the tilt angle. This axial distance varied between 657 mm and 743 mm over the sweep. The angle of the fitted plane agreed with the set angle to within 1° in every trial (mean 0.4°), which verifies the angle settings.

Ambient lighting

The same foam board was scanned at a camera height of approximately 900 mm under a range of lighting conditions, with the camera exposure left on automatic. Indoor conditions ranged from a dark room in the evening (blinds closed, lights off) to an afternoon room with the blinds open and the lights on. Outdoor trials were performed around midday in shade, in shade within direct sunlight, and in direct sunlight. The outdoor trials were taken twice per condition, one with the foam board, and one without the foam board. This was done as the effects of the high ambient light resulted in exposure settings that were unable to detect the cameras built in projected pattern, and thus the featureless white surface lost detection. The brick flooring was able to be detected in this condition, thus the results for the outdoor trials provide relational performance compare to the rest of the trials. Illuminance was read with a smartphone running the Physics Toolbox Suite application, which uses the front camera as a lux meter. Such readings are uncalibrated and are used here to indicate the order of magnitude and ranking of conditions, not precise values. The lux was recorded as 0 in the darkest condition and 24 400 in full sun.

5.1.3 Planar testing Results

Working Distance

Results for planar distance characterisation are shown in Table D.1. These are then plotted in Figure 5.1.

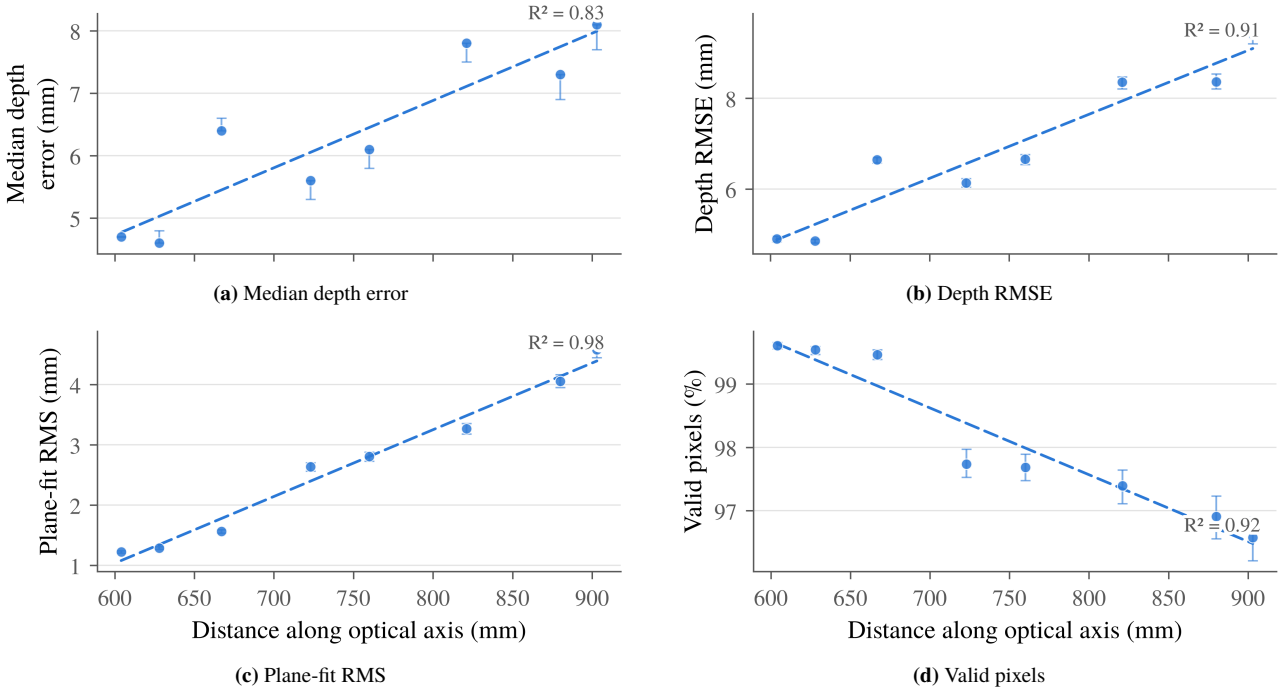

(a) Median depth error

(b) Depth RMSE

(c) Plane-fit RMS

(d) Valid pixels

Figure 5.1: Plots Showing Planar Performance Against Height. Per-trial Median with IQR Over n = 385–552 Frames.

Camera Tilt

Results for planar angle characterisation are shown in Table D.2. These are then plotted in Figure 5.2.

Ambient Lighting

Results for planar angle characterisation are shown in Table D.3. These are then plotted in Figure 5.3. Lighting conditions with corresponding illumination values are shown in Table 5.1

5.1.4 Discusion

Working Distance

Stereo theory predicts that depth noise grows with the square of the range, yet the data in Figure 5.1 follow closer to linear trend. Two factors likely contribute, firstly the ambient light was not constant as the trials were taken between 15:40 and 17:35, and the auto-exposure value fell from about 5300 for the three closest distances to between 1700 and 2050 for the rest. Second, the 604 to 903 mm range is narrow, so only a short, locally near-linear section of the quadratic curve is sampled. The practical trend is nonetheless clear. Plane-fit RMS (Figure 5.1c) increases with working distance while the proportion of valid pixels (Figure 5.1d) decreases, so the camera should be operated as close to the surface as the required field of view allows. Expressed as a percentage of distance, the median depth error rises only slightly, from 0.77% to 0.90%, with a weak fit (R^2

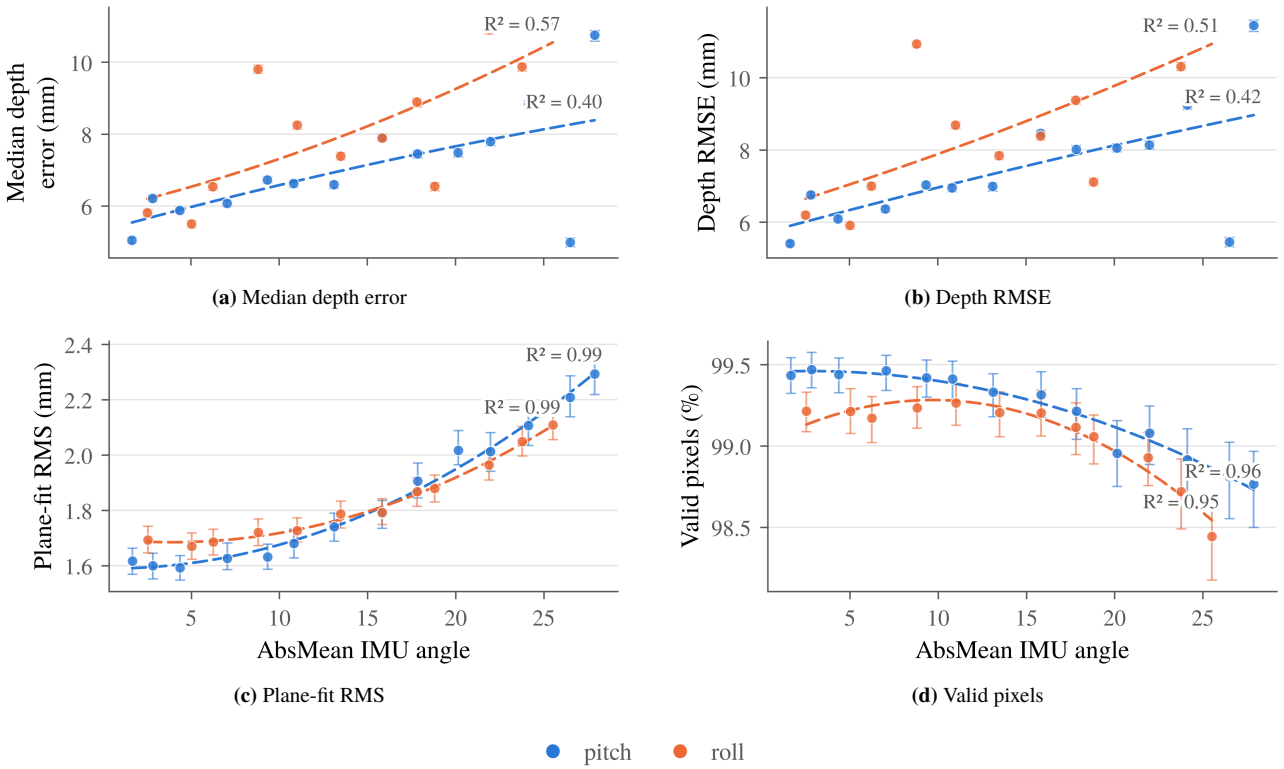

(a) Median depth error

(b) Depth RMSE

(c) Plane-fit RMS

(d) Valid pixels

Figure 5.2: Plots Showing Planar Performance Against Scan Angle. Per-trial Median with IQR Over n = 527–574 Frames.

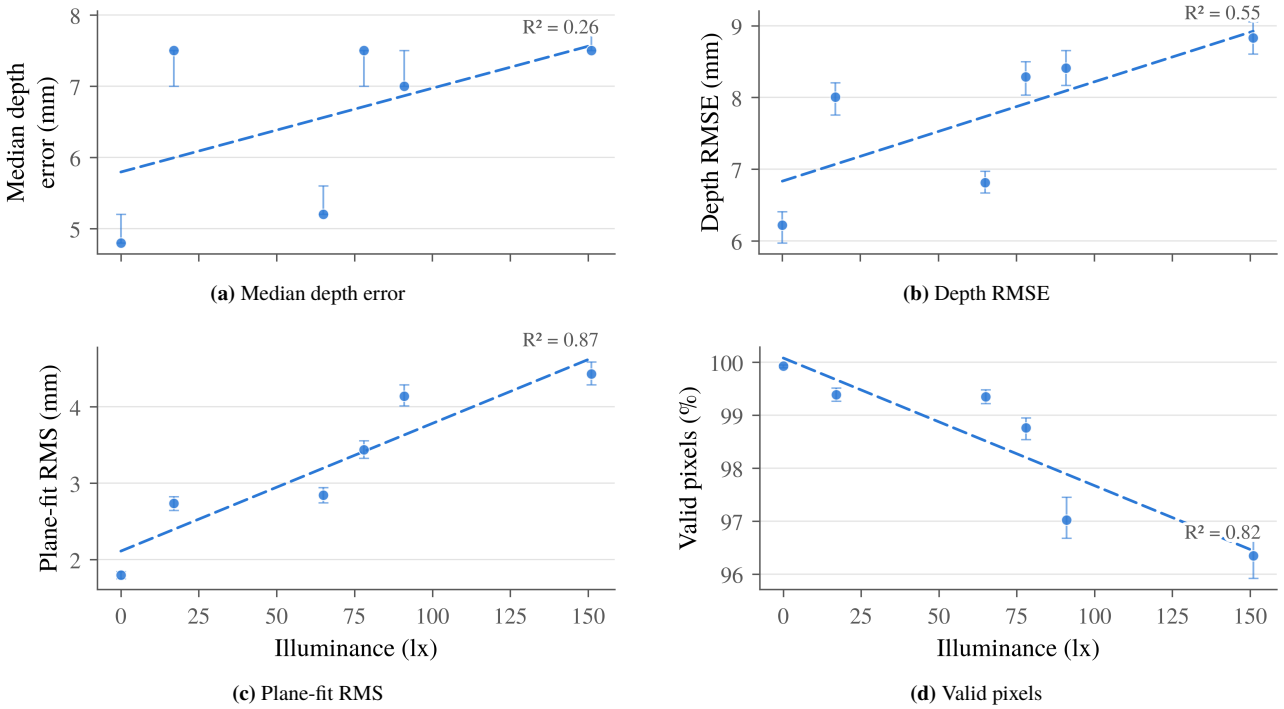

(a) Median depth error

(b) Depth RMSE

(c) Plane-fit RMS

(d) Valid pixels

Figure 5.3: Plots Showing Planar Performance Against Light Intensity, for Indoor Conditions with Foam Board. Per-trial Median with IQR Over n = 592–599 Frames.

Table 5.1: Planar characterisation under different lighting conditions at a height of approximately 0.9 m. Lux values are indicative (smartphone sensor).

Surface	Condition	Lux
Indoor foam board	Evening, blinds closed, Florescent lights off	0
	Afternoon, blinds closed, Florescent lights off	17
	Evening, blinds closed, Florescent lights on	65
	Afternoon, blinds closed, Florescent lights on	78
	Afternoon, blinds open, Florescent lights off	91
	Afternoon, blinds open, Florescent lights on	151
Outdoor foam board	Shade	388
	Direct sun, under shade	2 650
Outdoor, no board	Shade	213
	Direct sun	9 800
	Direct sun, under shade	10 000

= 0.18). This approximately constant relative error is consistent with a fixed bias or scale error rather than a distance-dependent effect, and is removed by detrending.

Camera Tilt

Depth noise is expected to rise with obliquity, as the surface is foreshortened and its depth varies across each matching window. A similar relationship is visible in the plane-fit RMS and valid pixels (Figure 5.2c and Figure 5.2d), which both degrade as the tilt increases, with the steepest change beyond about 15°. Quadratic trends were fitted to avoid overfitting with higher-order polynomials, and they describe both metrics well (R^2 ≥ 0.95). The median depth error and depth RMSE (Figure 5.2a and Figure 5.2b) show only weak trends, and neighbouring angles differ by up to 6 mm, which points to the repeatability of re-positioning the board and measuring its height as the dominant uncertainty in these metrics. Because the camera pivots at a fixed mount height, tilting it also increases the axial distance to the board, so these results combine the effects of angle and range and do not isolate angle alone. This is accepted here, since the mount height is also fixed when scanning sediment, so the results reflect the camera's performance in its intended use.

Ambient Lighting

Two conclusions follow. First, ambient light degrades active stereo depth even well within the sensor's range, because the projected infrared pattern must compete with ambient light for contrast. Daylight appears more damaging per lux than artificial light (Figure 5.3c and Figure 5.3d): 91 lx of daylight gave 4.14 mm RMS and 97.0% valid pixels, against 2.84 mm and 99.3% at 65 lx of fluorescent light. This is consistent with the greater near-infrared content of sunlight, so illuminance in lux is only a rough predictor of performance. Second, the surface matters as much as the illumination. The board failed completely at 2 650 lx in sun-exposed conditions, whereas the textured ground returned complete, low-noise data at 9 800 lx, so the failure cannot be attributed to light level alone. A bright, featureless white surface is a worst case for stereo matching, and the planar results should therefore be read as pessimistic bounds on the noise expected on naturally textured sediment. Wet or glossy sediment, which is more specular, was not tested outdoors and may behave worse. These results support keeping the shading tarp on the rig to limit the ambient infrared reaching the scan area.

5.2 3D Reference Characterisation

A Whole New Dimension

The planar tests establish the sensor's noise behaviour but not how accurately it reconstructs three-dimensional

relief, which is affected by resolution, filtering, distance and the fusion of several views. These effects were measured against reference objects with known geometry using a modified version of the main capture program, which records scans under controlled software settings and hardware set-ups.

5.2.1 Reference objects and evaluation

Three 3D-printed reference objects were used (chapter C). Each has a 100×100 mm footprint. Two are sinusoidal "egg carton" surfaces of the form $z = \sin(x) + \sin(y)$ at a large scale (feature spacing 50 mm, amplitude 20 mm) and a small scale (25 mm, amplitude 15 mm), and the third is a square-based pyramid of 50 mm height. The smooth sinusoidal surfaces are analogous to the smooth mounds produced by sand prawns, whose heights of 3 to 5 cm were reported in the literature review, and the small carton tests the lateral resolution of the camera. The pyramid has sharp edges and planar facets, which show whether the filters preserve geometric edges or smooth them.

Scans were evaluated with the `evaluate_scans` tool described in section 4.5. The tool rigidly aligns each scan to the reference .stl model, using feature-matching starts followed by point-to-plane ICP, and then computes point-to-surface distances. Scale is not estimated, so any scale error in the camera data appears as geometric error. The metrics used are Accuracy RMSE, signed bias, completeness at 2mm and computed F-score at 2mm. F-score is a combined metric which aims to represent both precision and completeness.

$$P(\tau) = \frac{1}{|S|} \sum_{s \in S} \mathbb{1}\Big[\min_{r \in R} \|s - r\| < \tau \Big] \tag{5.1}$$

$$C(\tau) = \frac{1}{|R|} \sum_{r \in R} \mathbb{1}\Big[\min_{s \in S} \|r - s\| < \tau \Big] \tag{5.2}$$

$$F(\tau) = \frac{2\,P(\tau)\,C(\tau)}{P(\tau) + C(\tau)} \tag{5.3}$$

Let S denote the set of scanned points and R the set of points sampled from the reference model. For a distance threshold τ, precision Equation 5.1) is the fraction of scanned points lying within τ of the reference surface, and completeness (Equation 5.2) is the fraction of reference points lying within τ of the scan. The F-score (Equation 5.3) is their harmonic mean, and therefore penalises both spurious points and missing coverage. It is calculated at $\tau = 0.5$, 1, 2 and 5 mm, however for the purposes of this report, all results consider the threshold of $\tau = 2$ mm as this gives a balance of high accuracy representation, whilst still including relatively large bias/point spread.

5.2.2 Testing Procedure

Camera configuration

The first sweep fixed the height and tilt of the camera and scanned the large egg carton and the pyramid with different stream configurations. The resolution ranged from 1280×720 to 640×360 pixels, the depth units from 1 mm to 0.1 mm, and all of the built-in visual presets were tested except the Hand preset, which is intended for gesture tracking. The depth unit sets the quantisation step of the 16-bit depth values, and the finest setting (0.1 mm) still provides a maximum range of 6.5 m, well beyond the working distance.

Post-processing filters

With the configuration fixed, five stationary scans of each object were recorded at each of the tested heights, without tilt or rotation. The data were recorded to files so that the RealSense post-processing filters could

be applied offline in many combinations to the same scans. The filters investigated were decimation, spatial, temporal and hole-filling, and a sweep was made over their smoothing alpha, smoothing delta and magnitude parameters shown in Table 5.2. Decimation reduces the resolution, the spatial filter performs edge-preserving smoothing, the temporal filter averages depth across frames, and hole filling interpolates invalid pixels. Each combination was evaluated against the reference, and the process was repeated at each height to determine whether the best filter settings depend on the working distance.

Table 5.2: Filter parameter ranges tested (one factor at a time).

Filter	Parameter	Values	Default
Decimation	magnitude	2–6, 8	–
Spatial	magnitude	1–5	2
	α	0.25–1.0	0.5
	δ	1–50	20
	holes fill	0–5	0
Temporal	α	0.05–1.0	0.4
	δ	1–100	20
	persistence	0–8	3
Hole filling	mode	0, 1, 2	–

Tilt and rotation

Finally, a full scan of each object was recorded at rotations of $-45°$, $0°$, $45°$ and $90°$ and at tilts from $-25°$ to $+25°$ in steps of $5°$. This gives 44 views of each object. These views were aligned and fused by the pipeline in chapter 6. Subsets of the views were fused and each fusion was evaluated against the reference, which shows how the reconstruction accuracy depends on the number of views and on the combination of tilt and rotation.

5.2.3 3D Characterisation Results

Camera Configuration

Results for 3D parameter characterisation are shown in Table D.4, Table D.5 and Table D.6. These are then plotted in Figure 5.4, Figure 5.5 and Figure 5.6 respectfully.

A full heatmap describing the performance of all parameter combinations can be found in section E.3.

Filters

A full comparison of all filter parameters is infeasible for the scope of this report. Thus only a comparison between a baseline evaluation, with no filters applied, and the best performing filters is shown in Figure 5.7, with the data shown in Table D.7. The filter settings can be found in Table 5.3

Table 5.3: Post-processing filter settings of the best-performing configuration for all objects combined and for each object.

	Decimation	Temporal Filter			Hole Filling	
Object	Magnitude	α	δ	Persistence	Mode	Spatial Filter
All objects	Disabled	0.1	20	Disabled	Fill From Left	Disabled
Big Egg Box	2	0.1	20	Disabled	Fill From left	Disabled
Small Egg Box	Disabled	0.4	20	Valid in 2/8	Disabled	Disabled

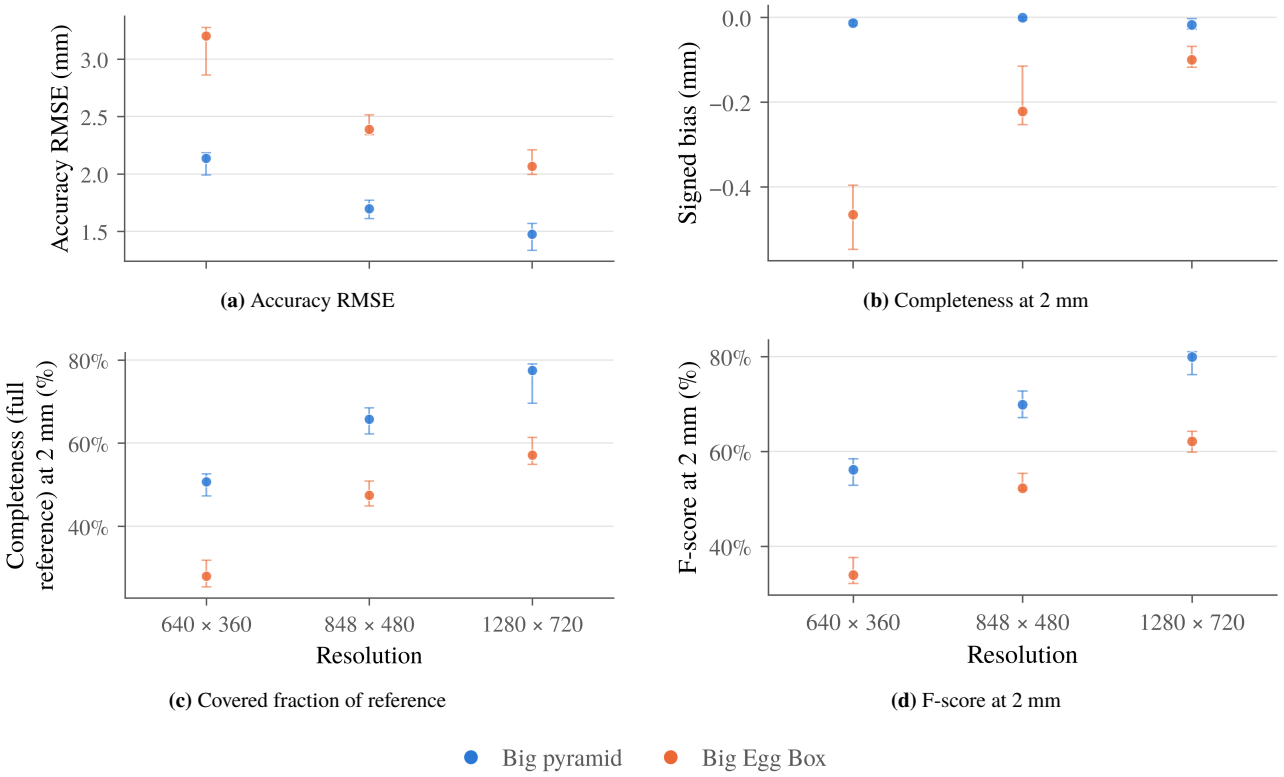

(a) Accuracy RMSE

(b) Completeness at 2 mm

(c) Covered fraction of reference

(d) F-score at 2 mm

Figure 5.4: Plots Showing 3D Performance Against Camera Resolution. Median with IQR of n = 15 scans per point, Taken Over Three Depth Units, and Five Presets.

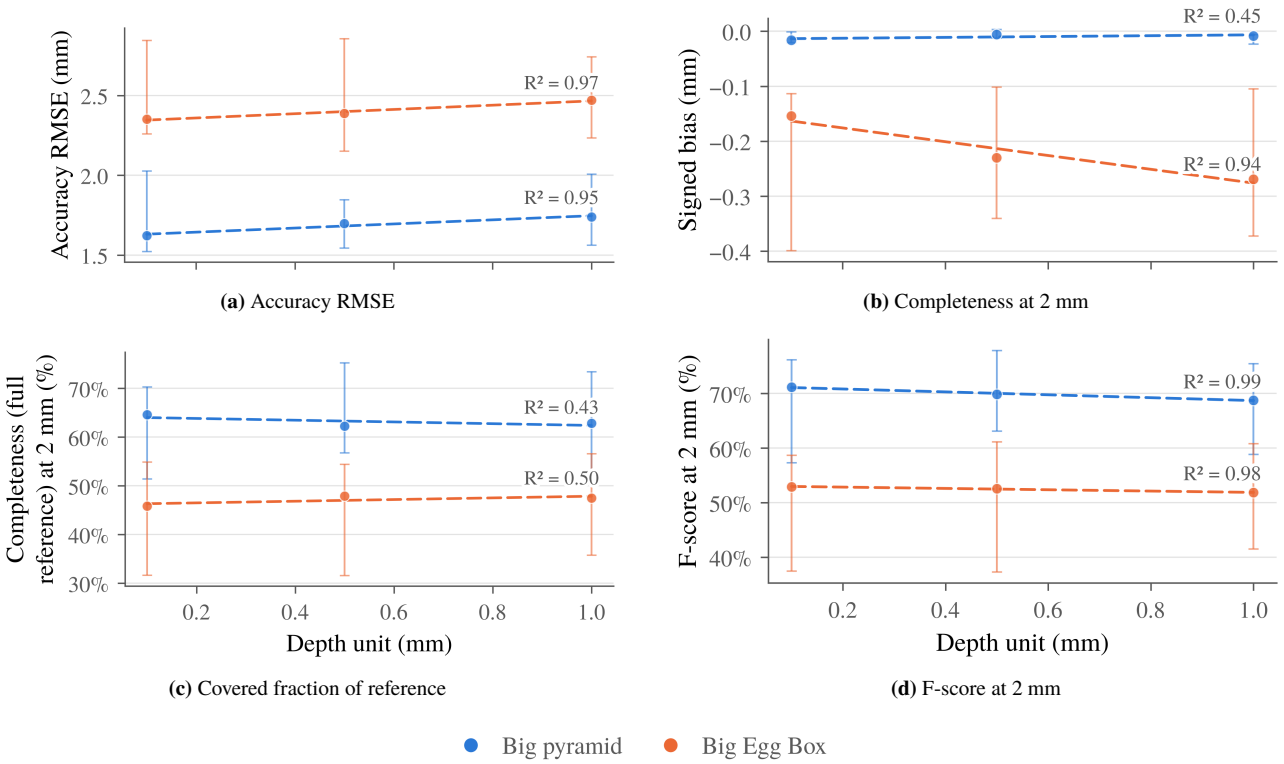

(a) Accuracy RMSE

(b) Completeness at 2 mm

(c) Covered fraction of reference

(d) F-score at 2 mm

Figure 5.5: Plots Showing 3D Performance Against Depth Units. Median with IQR of n = 15 scans per point, Taken Over Three Resolutions, and Five Presets.

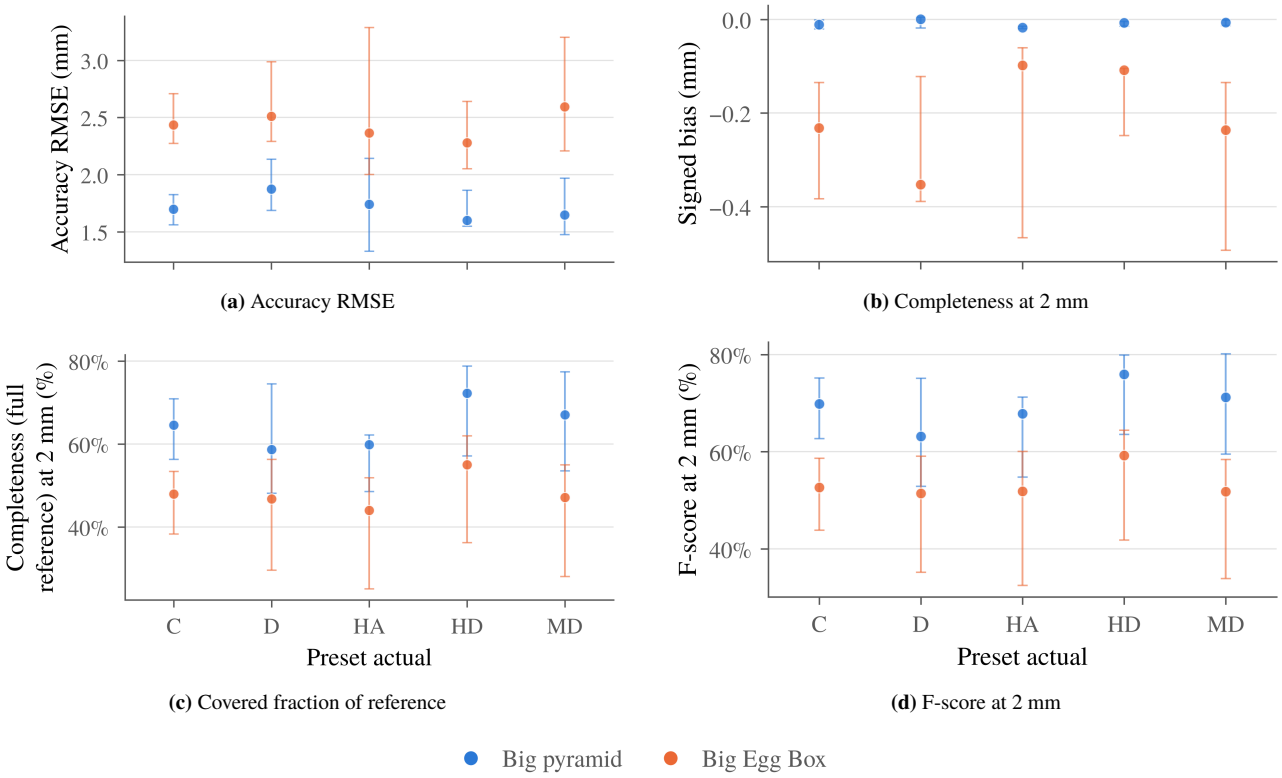

(a) Accuracy RMSE

(b) Completeness at 2 mm

(c) Covered fraction of reference

(d) F-score at 2 mm

Figure 5.6: Plots Showing 3D Performance Against Camera Preset. Median with IQR of n = 15 scans per point, Taken Over Three Depth Units, and Three Resolutions.

Tilt and Rotation

Again, a full comparison of all metrics agsainst all variables is infeasible for the scope of this project. The results shown in Table D.8 are plotted in Figure 5.8, and describe the comparison between a fusion of every tilt and rotation, compared to a single scan. This combination was chosen as it represents the best overall performance of the system as a whole.

A More in depth comparison of the different combinations can be seen in section E.2.

5.2.4 Discussion

Camera Configuration

Of the stream settings, resolution had the clearest effect (Figure 5.4). For the large egg box, moving from the default 848 × 480 to the maximum 1280 × 720 reduced the RMSE by 0.32 mm (13.4%) and increased completeness by 9.6 percentage points, as expected since the focal length in pixels, and therefore depth precision, scales with resolution. The pyramid scored better than the large egg box at every resolution. This was not expected, as sharp edges are usually smoothed more than gentle curves, but the egg box has a higher density of features within its footprint, and so more of its surface lies near curvature that the camera cannot fully resolve.

The depth unit had only a small effect (Figure 5.5), with performance improving slightly at finer units. The difference was smaller than expected, even on the pyramid, whose inclined facets were expected to benefit most from finer quantisation. This indicates that sensor noise, rather than quantisation, limits accuracy. The 0.1 mm unit gave the best results and costs nothing within the working range, although coarser units perform almost as well.

The High Density preset performed best on both objects (Figure 5.6). This is consistent with its aim of maxim-

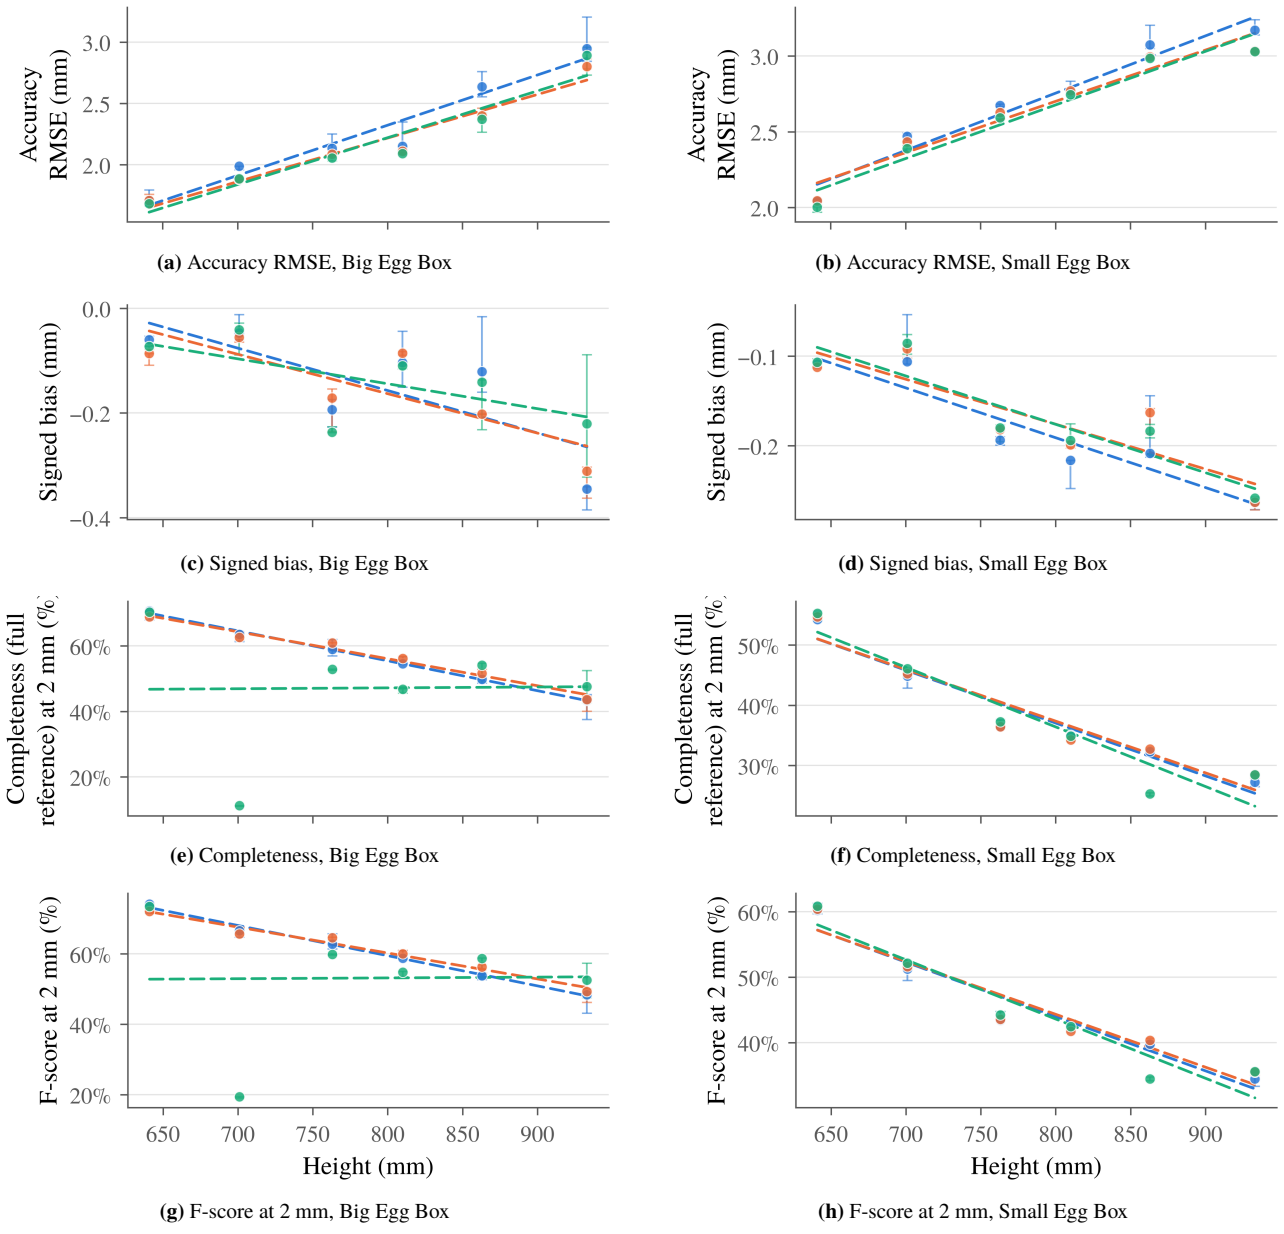

(a) Accuracy RMSE, Big Egg Box

(b) Accuracy RMSE, Small Egg Box

(c) Signed bias, Big Egg Box

(d) Signed bias, Small Egg Box

(e) Completeness, Big Egg Box

(f) Completeness, Small Egg Box

(g) F-score at 2 mm, Big Egg Box

(h) F-score at 2 mm, Small Egg Box

Figure 5.7: Plots Showing 3D Performance Comparison Between No Filter and Best Performing Filter, Best Fit (R^2 0.00 to 1.00). Median with IQR of n = 4–5 Scans Per Point.

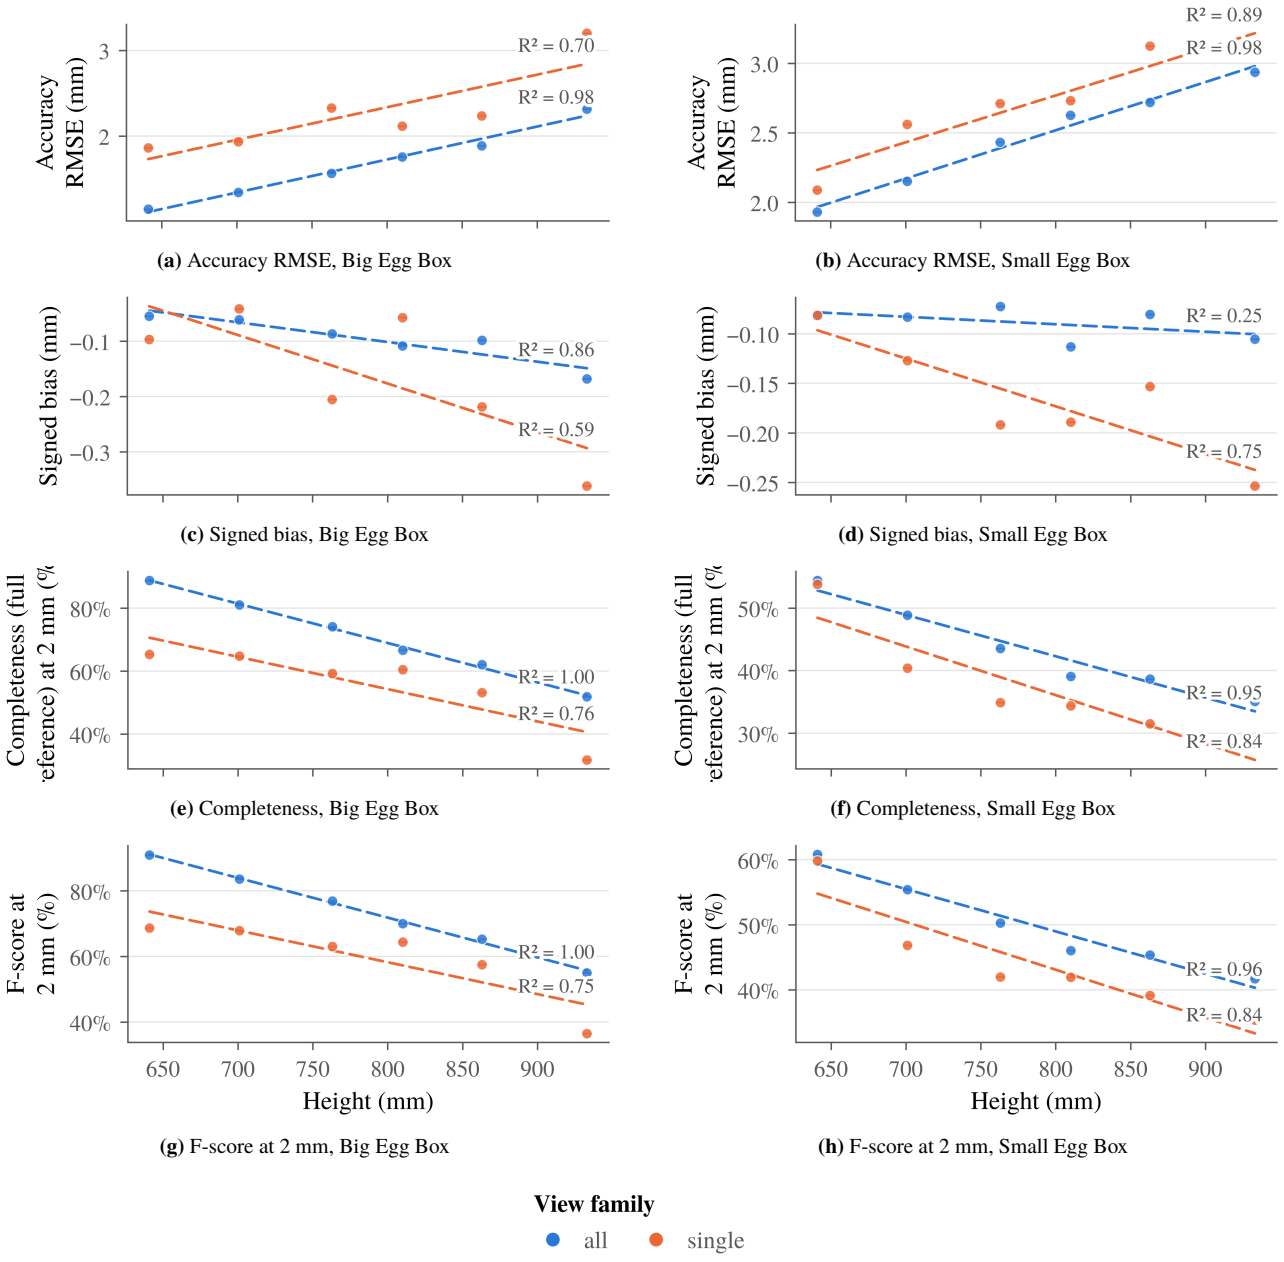

(a) Accuracy RMSE, Big Egg Box

(b) Accuracy RMSE, Small Egg Box

(c) Signed bias, Big Egg Box

(d) Signed bias, Small Egg Box

(e) Completeness, Big Egg Box

(f) Completeness, Small Egg Box

(g) F-score at 2 mm, Big Egg Box

(h) F-score at 2 mm, Small Egg Box

Figure 5.8: Plots Showing 3D Performance Between No Fusion and All View-Point Fusion. Median with IQR of n = 1 Point Clouds Per Point.

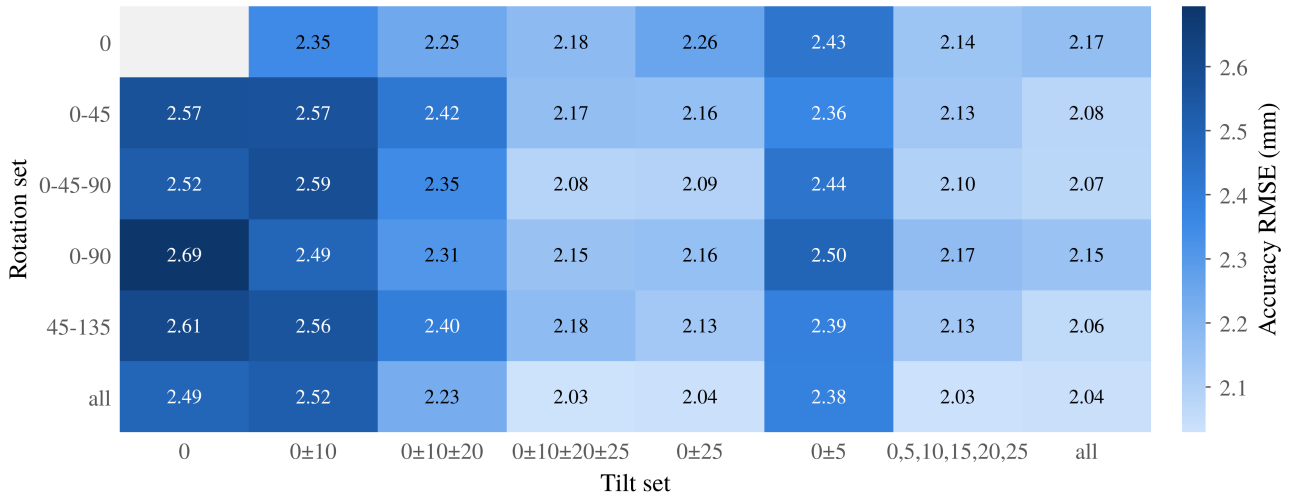

Figure 5.9: Grey cells: no data for that combination. Median with IQR of n = 12 point clouds per cell. Fixed: Kind = fused; Fusion cell = 0.0005; Scan voxel = 0.0005; Status = ok. Median taken over: Object (2), Height (6).

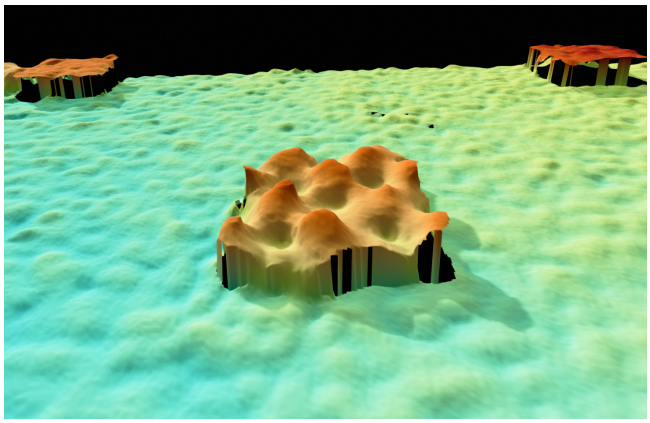

(a) 3D Mesh of Single Scan Viewed From Side

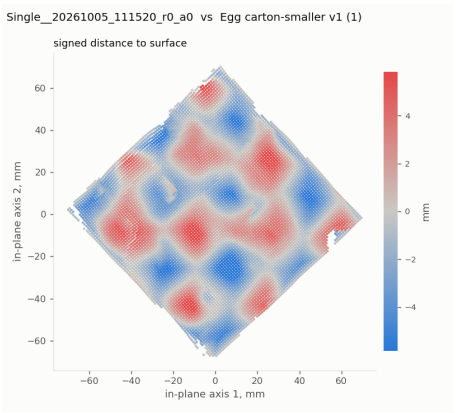

(b) Evaluation Showing Error Map of Single Scan Compared Against Reference Object

Figure 5.10: Visual Representation of Single Scan Performance Using Small Egg Box Surface

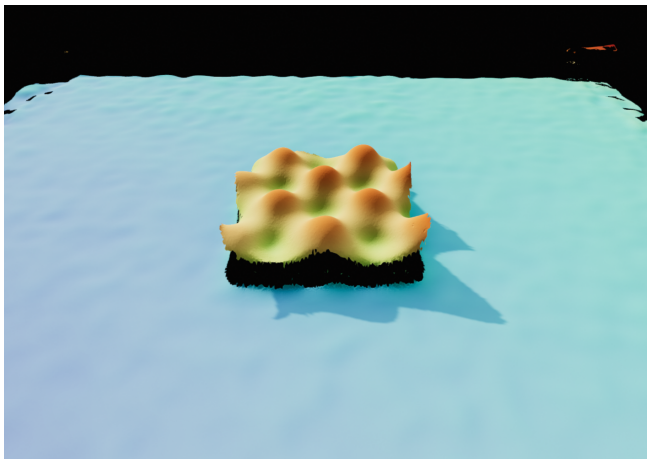

(a) 3D Mesh of Fused Scan Viewed From Side

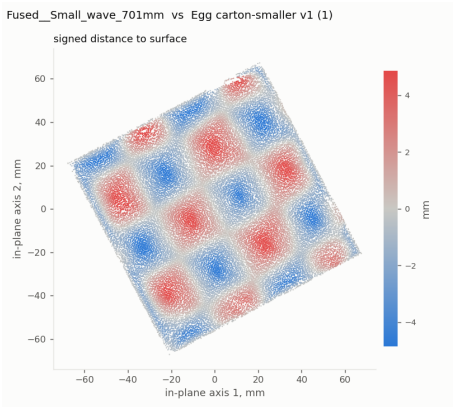

(b) Evaluation Showing Error Map of Fused Scan Compared Against Reference Object

Figure 5.11: Visual Representation of Fused Scan Performance Using Small Egg Box Surface

ising the number of valid depth pixels, which leaves fewer holes in the scan and improves completeness. The High Accuracy preset gave a low RMSE but poor completeness, so it may perform better once multi-view fusion fills its holes, which should be investigated further.

Filters

The post-processing filters gave only a small benefit (Figure 5.7), with the best settings listed in Table 5.3. The detailed heatmaps in section E.1 show a general trend that stronger temporal smoothing, little or no spatial filtering and minimal decimation give the best results. However, these effects are small compared with those of height: the variation in accuracy across filter settings at one height is smaller than the variation between heights. A larger benefit from aggressive temporal filtering was expected, since each scan is static and temporal noise should average out. A likely explanation is that the evaluation uses a single best frame on a small object, where the remaining error is dominated by spatially correlated error that averaging over time cannot remove.

Decimation caused only a small loss in accuracy, so it could be used to reduce processing time after capture. It is still not recommended, as the loss of detail is likely to be larger on large-area scans with finer relief.

Tilt and rotation

Multi-view fusion improved every metric at every height (Figure 5.8). Fusing all viewpoints raised the F-score by an average of 14.0 percentage points for the large egg box and 5.8 points for the small egg box, which shows that fusion improves accuracy without any change to the hardware. The error maps (Figure 5.10b and Figure 5.11b) show that fusion does not remove the largest errors, which remain at around ± 4 mm in both cases, but it smooths the noise considerably and gives a much more uniform error distribution. This matters more for this project than the absolute accuracy. Random noise adds directly to the measured slope, so a single scan will report a rougher surface than a fused one, even where their accuracy is similar. The same effect is visible on the flat floor around the object (Figure 5.10a and Figure 5.11a), which shows large undulations in the single scan but is consistently flat in the fused scan.

Chapter 6

Data Processing Pipeline

The capture stage (chapter 4) produces one session per experiment. A session contains multiple individual scans of the same patch of sediment, one per combination of mount rotation and tilt. Each scan stores the single best depth frame as a coloured point cloud (`.ply`), together with the AprilTag detections for the whole recording, the IMU sample taken with the best frame, the full recording and a text dump of the camera configuration.

A single scan of the surface contains several factors that result in lower quality reconstructions. It contains the rig and regions outside of the frames bound due to the wide field of view of the camera. It is noisy, particularly where the surface is seen at a grazing angle. It also has holes wherever relief hides part of the surface from that viewpoint. The processing pipeline turns a session of multiple of such scans into a single, clean reconstruction of the region of interest and then into numbers that describe its roughness. There are three main stages in the data processing pipeline.

1. **Cropping** (section 6.1) cuts every scan down to the square of sediment bounded by the four corner AprilTags.
2. **Fusion** (section 6.2) registers the cropped scans into one common coordinate frame and merges their overlapping regions into a single, lower-noise surface with fewer holes.
3. **Roughness estimation** (section 6.3) levels and detrends the fused surface and computes standardised surface-texture parameters, including rugosity.

Each stage is a separate program that reads the previous stage's output files, so any stage can be rerun or inspected on its own. The language of each program was chosen for its workload. Cropping is mostly bookkeeping and image geometry, so it is written in Python with Open3D and NumPy. Fusion runs nearest-neighbour searches, feature matching and iterative registration on millions of points, so it is written in C++ with the Point Cloud Library. The roughness metrics are array operations on a regular grid and are written in Python with NumPy and SciPy. A small Python script (`bulk_process.py`) chains cropping and fusion over many sessions without user interaction.

6.1 Point Cloud Cropping

In Too Deep

The raw point clouds have three problems that make them unsuitable for direct use.

1. **The field of view is larger than the region of interest.** The Intel RealSense D456 has a wide field of view, so even with the camera pointed straight down it records depth well outside the patch of sediment inside the experimental rig.

2. **The view changes with every scan.** Because the camera is carried on a rotating and tilting mount, tilted scans look past the edge of the rig and often capture the rig frame itself. These points do not belong to the surface and would corrupt both registration and the roughness statistics.

3. **Large scans are expensive and overlap poorly.** Fusion compares every scan with every other scan already placed. Points outside the shared region add computation without adding information.

The cropping step therefore reduces every scan to the same physical region: the square on the sediment bed whose corners are marked by four AprilTags. Because the tags are fixed in the world, this region is the same in every scan regardless of how the camera was rotated or tilted. This gives fusion a consistent, well-overlapping input. The step is implemented in `crop_pointclouds.py`, which processes an entire session at once. It crops automatically wherever it can and falls back to a manual crop when it cannot.

6.1.1 Finding the tags in each scan

AprilTag detection is performed during capture on the colour image, and every detection is written to a `.csv` file together with the frame number, the tag ID and the pixel coordinates of the tag's centre and four corners. Squares that look like tags but whose ID could not be decoded are also logged, with an ID of -1.

The point cloud of a scan comes from one particular frame of the recording. The cropper reads that frame number from the scan's IMU file, which was written at the same moment as the point cloud. It then collects the detections from that frame, however tag detections are accepted from ± 15 frames around the best frame, since each scan is taken from a static position. This allows for some redundancy against frames where detections momentarily failed. Decoded tags are always preferred, for region identification, however undecoded squares are used to fill corners where no identified tags are visible, and only when they lie close to where the missing tag is expected.

6.1.2 Learning the tag layout of a session

Not every tag will be detected, or visible in each scan wi thin a session. Glare, shadow or a steep tilt often hide tags. Therefore, in order to maintain the ability to automatically crop each scan to the correct region, the program first learns the layout of the tags from the scans in the session where all four were detected. From those frames it records the order in which the four tag IDs appeartgoing around the square, which says which tags share a side and which are diagonal, the ratio of the square's side length to a tag's side length and the aspect ratio of the square, i.e. how each side compares to its neighbours.

The median of each quantity over all four-tag frames is used, which makes the layout robust to a few poor detections. With this layout, missing corners can be predicted from the tags that were seen. With three tags, the fourth is placed to complete a parallelogram, $\mathbf{p}_4 = \mathbf{p}_1 + \mathbf{p}_3 - \mathbf{p}_2$. With two tags on the same side, the square is completed by stepping perpendicular to that side by the learned aspect ratio. With two diagonal tags, the side vector is solved from the diagonal and the aspect ratio before stepping out in the same way.

The crop region is then defined by the inner corner of each tag, i.e. the corner nearest the centre of the square region of sediment. This places the crop boundary just inside the tags, so the tags themselves are removed along with everything outside them.

6.1.3 Relating 3D points to image pixels

The tag positions are known in the 2D colour image, but the crop has to be applied to 3D points. The two are linked through the camera model. During capture the depth image is aligned to the colour image before the point cloud is exported, so every 3D point originates from one integer pixel of the colour image sensor. The cropper

reads the colour camera intrinsics (focal lengths f_x,f_y, principal point (c_x, c_y) and distortion coefficients) directly from the scan's recording file and projects every 3D point back into the RGB image plane. A 3D point (X, Y, Z) in the camera frame maps to

$$u = f_x\, x_d + c_x, \qquad v = f_y\, y_d + c_y, \tag{6.1}$$

where (x_d, y_d) are the normalised coordinates $(X/Z,\ Y/Z)$ after lens distortion has been applied. The distortion is computed exactly as the RealSense SDK computes it. Using the SDK's form reprojects the cloud onto the pixel grid to within about 10^{-5} pixels.

This property also provides a self-check. If the intrinsics really belong to this cloud, every reprojected point should land almost exactly on an integer pixel. The cropper measures the root-mean-square distance of the reprojected points from the nearest integer pixel centre,

$$\rho = \sqrt{\frac{1}{N}\sum_{i=1}^{N} \left\|\mathbf{u}_i - \mathrm{round}(\mathbf{u}_i)\right\|^2}, \tag{6.2}$$

and rejects the automatic crop if ρ exceeds 0.15 pixels. The same test is used to choose how the distortion coefficients should be applied, since the stored coefficients can describe either the forward or the inverse distortion. The interpretation with the smallest ρ is kept. This guards against silently cropping a scan with the wrong camera model, for example if a recording was captured with a different resolution.

6.1.4 Building the crop volume

The simplest crop would keep every point whose pixel falls inside the tag-bounded quadrilateral. This is equivalent to cutting the cloud with a pyramid whose apex is at the camera. The walls of that pyramid follow the camera's lines of sight, so a tilted scan would be cut along a different boundary on the sediment than an untilted scan, and the scans would no longer cover the same region.

Instead, the cropper builds a vertical prism through the tag corners, so that the crop boundary is the same physical line on the bed for every scan:

1. **Fit the bed surface.** A plane is fitted to the points in a narrow band (25 pixels) just inside the crop edge. The band is used rather the whole region so that the test object or sediment relief in the middle of the square does not influence the fit. The fit uses RANSAC: planes through three randomly chosen points are tried 500 times, and the plane with the most points within 5 mm of it is kept. RANSAC is used because it ignores outliers entirely instead of averaging them in, which matters near the edge of the crop where stray points from the rig may remain.

2. **Fit the tag plane.** A second RANSAC plane is fitted to the points lying on the tags themselves. The tags can sit slightly above the sediment, so the 3D corners of the crop region are placed where the pixel rays through the inner tag corners meet this plane. For a pixel ray with direction $\mathbf{r}$ and a plane with normal $\mathbf{n}$ and offset d, the intersection is
$$\mathbf{X} = \frac{-d}{\mathbf{n}\cdot\mathbf{r}}\,\mathbf{r}. \tag{6.3}$$

3. **Extrude along the surface normal.** All points are projected onto the bed plane, and a point is kept if its projection falls inside the four 3D corners. The result is a column whose walls are perpendicular to the bed and pass through the inner tag corners, independent of the camera's viewing angle.

If no bed plane can be found, the cropper falls back to the pyramid crop described above and notes this in its output. A crop that keeps fewer than 1000 points is treated as a failure.

6.1.5 Review and manual fallback

Unless run in batch mode, each automatic crop is shown in a preview window with the removed points dimmed and the crop boundary drawn as a wireframe. The user can accept it, keep the original cloud, or switch to a manual crop. The manual crop starts from points the user picks in the viewer and then lets the box be moved and resized with the keyboard. It is also used whenever the automatic crop fails, for example when fewer than two tags were detected. This keeps the user in the loop for the cases where automatic tag detection is least reliable.

6.1.6 Outputs

For each scan the cropper writes:

- **The cropped point cloud**, in the same coordinate frame as the input. It is written with 32-bit floating-point coordinates, because Open3D's default writer produces 64-bit coordinates that the PCL-based fusion program fails to read.
- **A tag sidecar file** giving the 3D positions of the four corners of every visible tag. Cropping removes the tags from the cloud, but fusion can use them for alignment (Section 6.2.3), so their positions are measured on the uncropped cloud before they are discarded. Each tag's pixel corners are taken as the median over every frame of the recording, which removes detector jitter. A plane is fitted by least squares to the depth points on the central 70% of the tag, this central region is selected to exclude the blending of the tag and bed surface. The 3D corners are the intersections of the corner pixel rays with that plane.
- **Copies of the IMU, detection and configuration files**, so that the cropped session has the same folder layout as a raw session and can be passed to fusion unchanged.

6.1.7 Limitations

The automatic crop assumes that the sediment near the edge of the square is roughly flat, so that the band plane represents the bed. If relief reaches the crop edge, the fitted plane, and therefore the prism, can tilt. The layout learned from four-tag frames is pooled across all poses, although perspective changes with tilt, so corners extrapolated from only two tags in a strongly tilted view may be slightly misplaced. The prism also has no height limit, so isolated noisy points directly above or below the square are kept.

6.2 Multi-Scan Fusion

Together is Almost Always Better

A single scan is noisy, has holes where relief hides the surface, and has errors that depend on the viewpoint. Combining registered scans addresses all three. Views from different directions see regions occluded in others, and the mean of n measurements with independent errors of standard deviation σ has a standard deviation of $\sigma / \sqrt{n}$. This however does not alleviate errors shared by all scans, such as a calibration or scale bias, all of which cannot be reduced through averaging.

Fusion is performed by a C++ program, `fuse_scans`. It runs as a sequence of stages: load and preprocess, estimate initial poses from the rig, coarse alignment, fine registration and reconciliation. The program can stop after coarse and after fine alignment to show every scan in an interactive viewer, colour-coded and listed with its alignment quality. From there the user can exclude a bad scan, nudge one into place by hand, or rerun a stage, before the session is fused. In batch processing these reviews are skipped.

The common frame is that of the first untilted scan, called the *reference*. The remaining scans are processed in order of increasing tilt, so that each new scan is as similar as possible to what has already been placed.

6.2.1 Preprocessing

Each cropped cloud is first thinned with a voxel grid of user defined size, default of 3 mm: space is divided into 3 mm cubes and all points in a cube are replaced by their centroid. This makes the density uniform, which most of the later algorithms assume, and it greatly reduces their cost. It also sets the finest detail the fused result can contain, so the voxel size is chosen close to the noise floor of the sensor. A statistical outlier filter then removes isolated points, defined as points whose mean distance to their 30 nearest neighbours is more than one standard deviation above the average.

Finally, a surface normal is estimated at every point from the principal axes of its neighbourhood within 15 mm. All normals are oriented to point towards the camera. This matters later: normals from several scans are averaged, and a normal and its opposite would cancel. The normals are computed once here and carried through every later transformation, because re-estimating them inside registration accounted for most of the program's run time.

6.2.2 Initial poses from the rig and IMU

Registration algorithms only converge when they start close to the right answer. The rig provides a natural starting point, because each scan's folder name records the commanded rotation and tilt. These commanded angles proved to be inaccurate, however. Commanded tilts of 15° produced measured tilts between 13.6° and 17.4°, and individual steps varied from 0.6 to 1.24 times the commanded step. In addition, because the rotating mount sits on top of the tilt stage, the tilt axis appears in a different direction in the camera frame for every rotation position, so a single fixed tilt axis cannot describe the rig.

The tilt is therefore *measured* rather than assumed. When the camera is stationary its accelerometer reads only gravity. The rotation that carries the gravity vector of a scan onto the gravity vector of the flattest scan at the same rotation position is exactly the tilt the mount applied, because the hinge axis is horizontal. Gravity readings with a magnitude outside 9.0–10.6 m/s^2 are rejected, since they indicate the camera was moving.

Two further quantities cannot be read from the IMU and are found by search:

- **The IMU-to-cloud axis convention.** The accelerometer and the exported point cloud use different axis conventions. Four candidate conventions (the identity and a 180° flip about each axis) are tried. For each, a few low-tilt scans are placed using the IMU tilt and compared with the reference scan. The convention with the best agreement is kept. On real data the correct convention gave residuals of 7–18 mm, against 42–76 mm for the wrong one.
- **The tilt pivot.** The IMU gives the angle of tilt but not the point the camera tilts about. The pivot is found by a coarse-to-fine grid search over a 0.12 m range, again scoring how well the placed scans agree with the reference. The convention and the pivot are searched together because an error in either can mask an error in the other.

The rotation direcition poses a different problem. Since the rotation stage is performed manually, the capture code names each rotation as a positive 45°increment, regardless of direction or order. IMU data cannot resolve this issue since the rotation axis is parallel to the gravity axis, and therefore can not be estimated using accelerometer data. The program therefore tries every combination of clockwise and anticlockwise steps, places one representative scan per rotation position accordingly, and keeps the combination under which the scans agree best. A joint search is used rather than deciding each step on its own, because a wrong early step makes all later steps fit worse, which only a joint search can detect. When AprilTag alignment is enabled, the tags measure the rotation angle of each rotation position directly, and the search is needed only for positions where no tags were seen.

6.2.3 Coarse alignment

The coarse stage places each scan to within a few millimetres of its final position. Scans are added one at a time, and each is aligned against a map built from every scan placed so far, rather than against only the previous scan. This reduces the accumulation of error (drift) along the sequence. Up to three candidate poses, or hypotheses, are generated for each scan:

1. **The rig seed**: the pose predicted by the measured tilt and inferred rotation of Section 6.2.2
2. **The AprilTag pose**, if tag alignment is enabled. Each tag contributes four 3D corners from the sidecar file, which are matched against the same tag corners already placed in the common frame. The best rigid transformation between the matched corners is found in closed form with the method of Umeyama [**?**], which uses a singular value decomposition of the corners' cross-covariance matrix. No scale term is estimated, because both point sets are already in metres. When gravity is available, a constrained version is used instead: the tilt is taken from the IMU and only the turn about the vertical axis and the translation are solved from the tags. The tags a scan sees often lie nearly in a line, so the tilt about that line would otherwise rest on depth noise across a single 80 mm tag, whereas the IMU measures tilt well. On one test session this reduced the disagreement between tag poses and the final registered poses from 1–2.5° to 0.2–0.6°.
3. **A feature-based pose**, computed only if the rig seed overlaps the map poorly. Fast Point Feature Histograms (FPFH) [**?**] describe the local shape around every point. Points with similar descriptors in the scan and the map are paired, and a RANSAC variant [**?**] repeatedly picks three pairs, rejects them early if the triangles they form have different shapes, and keeps the transformation that brings the most points into agreement. This method needs no initial guess, but it is unreliable on smooth, nearly flat sediment with few distinctive shapes, which is why it is only a fallback.

To choose between hypotheses, each is scored by comparing the transformed scan with the map. A sample of scan points is matched to the nearest map point. Points within 10 mm are counted as overlapping, and the score is the root-mean-square distance of the overlapping points divided by the fraction of points that overlap, and the lowest resulting value is chosen. Dividing by the overlap fraction prevents a hypothesis from scoring well by fitting a small sliver of the scan very tightly.

For nearly flat scenes such as the sediment bed, this nearest-neighbour score has a blind spot. If a scan is slid or spun within the plane of the bed, every point still has a nearby neighbour on the flat map, so the score hardly changes. In that case the program instead compares heights. Both clouds are expressed in the frame of the bed plane, binned onto a 10 mm grid, and the score is computed from the differences in mean height between corresponding cells. A 5 mm in-plane slide raised this height-field score by 36%, but the nearest-neighbour score by only 16%. Whether a scene counts as flat is decided once from the reference scan.

The rig seed is accepted directly if at least 15% of the scan overlaps the map. Only if it fails is FPFH tried, and the better of the two is kept. A tag hypothesis competes with the result on the same score. If the tag pose and the geometric pose disagree by more than 10°, a warning is printed, because a disagreement that large means that one of them is wrong, not merely imprecise.

6.2.4 Fine registration

The fine stage refines each coarse pose with the Iterative Closest Point (ICP) algorithm [**?**]. ICP alternates between pairing every scan point with its nearest map point and solving for the rigid transformation that minimises the distance between the pairs, until the pose stops changing. The point-to-plane form of ICP [**?**] is used, which minimises the distance of each scan point from the tangent plane at its paired map point rather than from the map point itself:

$$E(\mathbf{R}, \mathbf{t}) = \sum_{(i,j)} \Big((\mathbf{R}\mathbf{p}_i + \mathbf{t} - \mathbf{q}_j) \cdot \mathbf{n}_j \Big)^2, \tag{6.4}$$

where $\mathbf{p}_i$ is a scan point, $\mathbf{q}_j$ its paired map point, $\mathbf{n}_j$ the map normal, and $\mathbf{R}$, $\mathbf{t}$ the rotation and translation being solved for. Because of the voxel grid, a scan point and its true counterpart on the map are generally not the same physical point, only points on the same surface. Point-to-plane ICP lets points slide along the surface instead of penalising this offset, so it converges faster and more accurately on continuous surfaces than point-to-point ICP. Each step is solved with the linear least-squares approximation of Low [**?**]. Pairs further apart than 20 mm are ignored, so a coarse pose must be within roughly this distance for ICP to recover it, and at most 60 iterations are run.

As in the coarse stage, each scan is registered against a map of all scans placed so far. The map is stored as a voxel grid in which each voxel keeps running sums of the positions and normals of every scan that falls into it. Every scan therefore carries equal weight in the map. An earlier version re-thinned the merged map after every scan, which gave the newest scan about half the weight wherever it overlapped, so the map followed the latest scan rather than the consensus.

After the stage, a scan is flagged for review if ICP failed to converge, hit its iteration limit, or scored more than 1.4 times worse than the median scan of the session. The threshold is relative because what counts as a good score depends on the voxel size, the sensor noise and the relief of the surface. A scan whose refinement fails keeps its coarse pose, so that it remains visible in the review instead of silently disappearing.

6.2.5 Reconciling the overlapping scans

Once registered, the scans could simply be combined into one cloud and thinned with a voxel grid. This was the original approach, and it gives a poor result. The remaining registration error (a few millimetres) is larger than the output voxel, so where scans overlap, their surfaces land in neighbouring voxels and survive as two or three slightly offset sheets. A statistical outlier filter cannot remove a duplicate sheet, because the duplicate is just as dense as the real surface. The fusion step therefore reconciles the scans.

Weighting each measurement. Not all measurements are equally reliable, so each point i is given a weight

$$w_i = c_{s(i)} \cdot \max(0.05,\ |\hat{\mathbf{n}}_i \cdot \hat{\mathbf{v}}_i|) \cdot \min\left(1,\ \frac{r_0^2}{r_i^2}\right). \tag{6.5}$$

The first factor, $c_{s(i)}$, is the registration confidence of the scan the point came from, based on how much of the scan overlapped the map and how closely. The second factor is the cosine of the angle between the surface normal $\hat{\mathbf{n}}_i$ and the direction to the camera $\hat{\mathbf{v}}_i$. This factor is included as stereo depth error rises sharply when the surface is seen at a grazing angle, so such points are down-weighted. The third factor reflects the range r_i from the camera. This factor is included as the depth error of a stereo camera generally grows with the square of the range, so points beyond the reference range r_0 are down-weighted accordingly.

Grouping measurements into cells. For a nearly flat scene the points are expressed in the frame of the bed plane and binned into the user defined sized cells in the plane, each holding the heights of all points that fall into it. This *height-field* mode is chosen automatically when the scene is flat, which applies to the sediment bed. Scenes with strong relief use a volumetric mode instead, in which points are binned into 3D voxels and reconciled along the local surface normal.

Choosing one value per cell. Within each cell, the measurements are combined with a *trimmed weighted mean*. The samples are sorted by height, and the lowest and highest 25% of the total weight are discarded before the weighted mean of the remainder is taken. A plain weighted mean would let a single misregistered scan drag the surface towards it, while a weighted median discards most of the information. The trimmed mean lies between the two. Trimming by weight rather than by count means that a single high-confidence measurement cannot be outvoted by several poor ones. On a synthetic test with one of four scans misregistered by 8 mm, the error of the fused surface was 2.00 mm with a plain weighted mean and 0.35 mm with the trimmed mean. Overall, reconciliation halved the error of the original merge-and-thin approach, from 0.96 mm to 0.45 mm RMS.

Removing off-surface outliers. Points that still lie off the surface are then removed. For each point a plane is fitted to its 31 nearest neighbours, and the point is discarded if its distance from that plane is more than three standard deviations above the average. This replaced the statistical outlier filter used earlier, which judges points by local density. That filter removed genuine but sparsely sampled parts of the surface and punched holes in the fused result, whereas the plane test only removes points that are actually off the surface.

Smoothing. Finally, the surface is smoothed with Moving Least Squares (MLS) [**? ?**]. Around each point, MLS fits a quadratic surface to the neighbours within 10 mm and moves the point onto it, which also yields a clean normal. Because the fitting radius is larger than the residual misregistration, this removes the last traces of offset sheets and sensor noise along the surface normal while keeping features larger than the radius. It is, however, a low-pass filter and so relief with a wavelength comparable to or smaller than the 10 mm radius is attenuated. This must be kept in mind when interpreting the small-scale roughness parameters of section 6.3.

The result is a single point cloud per session, in the frame of the reference scan, with one point per occupied cell and a normal at every point. The program also reports the thickness of the overlapping surface, i.e. the spread of heights within a cell, before and after reconciliation.

6.2.6 Limitations

Registration is sequential: each scan is aligned against the scans before it, and no global optimisation revisits earlier poses afterwards, so an early error can propagate. The rig seed is accepted whenever it overlaps the map by 15% or more. On a flat bed even a pose with a wrong rotation can overlap that much, in which case only the tag hypothesis can correct it, so tag alignment should be enabled for sediment scans. The pivot of the rotation stage is not estimated, so an offset between the turntable axis and the camera must be absorbed by the tags or by ICP. The height-field mode assumes the surface has a single height at every point, so overhangs and vertical faces, such as the inside walls of a burrow, are averaged away. Finally, whilst the smoothing stage helps to remove sensor noise and point offsets, it does effectively low-pass filter the surface, and thus can remove sharp edges in the scan.

6.3 Surface Roughness Estimation

It's Not Flat?

The roughness program, `roughness_scans.py`, takes a fused point cloud and describes the roughness of the surface it represents. It runs five steps: level the scan, crop it, remove large-scale trends, resample it onto a regular grid, and compute the metrics. Most of the metrics follow the areal surface-texture parameters of ISO 25178-2 [7], so that results can be compared with other surface-metrology studies. Internally all units are in metres with reported units being millimetres.

6.3.1 Levelling

The fused scan is in the frame of the reference camera, so the sediment bed is generally tilted relative to the coordinate axes. Roughness parameters measure heights relative to a mean surface, so the scan is first levelled. This involves rotating scan so that the mean plane of the bed is horizontal and height is measured along the z axis. The mean plane is found by principal component analysis. The plane passes through the centroid of the points, and its normal is the direction in which the points vary least. To prevent large features from tilting the plane, the fit is repeated with sigma clipping, where points further than three times the spread estimate are dropped and the plane is refitted, this is then repeated up to five times. The spread is estimated with the median absolute deviation (MAD), scaled by 1.4826 so that it equals the standard deviation for normally distributed data. The reason for using MAD instead of normal standard deviation, is that MAD is less susceptible to the changes that result from discarded samples.

The levelled frame has its origin at the centroid, $+z$ along the plane normal pointing towards the camera, and x along the camera's original x axis projected onto the plane. Keeping the camera's horizontal direction means that the direction of ripples or other features reads the same in the results as in the images. The scan can then optionally be cropped to a region of interest or trimmed by a fixed margin around its border, where fused data are usually the sparsest.

6.3.2 Detrending

Although the surface has been levelled, there are still some sediment features forms that are not related to the roughness of said sediment. For example, a concave section of sediment, or residual tilt from the levelling stage. Removing these forms from the surface results in the isolated sediment texture surface, called the residual. For detrending, these non-textural forms are modelled as a 3D polynomial surface:

$$p(x, y) = \sum_{i+j \leq n} a_{ij}\, x^i y^j, \tag{6.6}$$

of degree n from 1 (planar) to 4 (quartic), and the residual $z - p(x, y)$ is kept as the roughness surface. The coordinates are scaled to the range $[-1, 1]$ before fitting so that the fit remains numerically stable at higher degrees, and the same sigma clipping as in levelling stops ripple crests, troughs or burrow mounds from dragging the trend.

The choice of degree is a trade-off. A low degree may leave some non-texture form in the residual and over-estimate roughness. A high degree may capture more complex forms but can also remove features of interest. A polynomial of degree n can follow roughly $n/2$ wavelengths, so on a small patch it can remove relief that is part of the texture being measured. The program therefore computes the metrics for every degree requested, using the same levelling, crop and grid for each so that the results are directly comparable, and raises a warning if a higher-order fit removes much more height variance than the planar fit while the patch spans fewer than n dominant wavelengths.

6.3.3 Gridding

Most roughness parameters are defined on a regular grid of heights, so the residual surface is resampled onto a square grid. By default the cell size is twice the median spacing between neighbouring points, so that most cells contain several points. The height of a cell is the mean of its points. Averaging c points with independent noise of standard deviation σ reduces the noise to $\sigma / \sqrt{c}$, so the gridding also acts as a mild low-pass filter at the scale of one cell.

Some cells contain no points, therefore any small enclosed holes, no larger than two cells across, are filled by

linear interpolation from the surrounding cells. Larger holes and gaps at the border are left empty and excluded from every metric, because interpolating across them would result in inaccurate metrics..

6.3.4 Height parameters

With z_k being the gridded heights relative to their mean and N being the number of valid cells, the height parameters are:

$$S_a = \frac{1}{N}\sum_k |z_k|, \qquad S_q = \sqrt{\frac{1}{N}\sum_k z_k^2}, \tag{6.7}$$

$$S_{sk} = \frac{1}{S_q^3}\frac{1}{N}\sum_k z_k^3, \qquad S_{ku} = \frac{1}{S_q^4}\frac{1}{N}\sum_k z_k^4, \tag{6.8}$$

together with the highest peak S_p, the deepest valley S_v and the total height $S_z = S_p + S_v$. S_a and S_q are the arithmetic and root-mean-square heights and are the most common single measures of roughness. The skewness S_{sk} indicates whether the surface is dominated by peaks ($S_{sk} > 0$) or by pits ($S_{sk} < 0$). A bed riddled with burrow openings would be expected to have negative skewness. The kurtosis S_{ku} indicates how sharp the features are, a surface with normally distributed heights has $S_{ku} = 3$, while isolated sharp peaks or pits raise it.

6.3.5 Rugosity and slope

Rugosity is the ratio of the true area of a surface to the area of its projection onto the horizontal plane. It is the areal equivalent of the chain-and-tape method used in benthic ecology, in which a chain is laid along the contours of the bed and its length compared with the straight-line distance. A perfectly flat surface has a rugosity of exactly 1, and any relief raises it.

To compute the true area, each square of four neighbouring grid cells is split along a diagonal into two triangles, and the area of each triangle is computed from the cross product of its edges. Summing over all squares whose corners are valid gives

$$\text{rugosity} = \frac{\sum A_{\text{triangles}}}{N_q\, g^2}, \tag{6.9}$$

where N_q is the number of valid squares and g the grid spacing. The ISO 25178 developed interfacial area ratio is the same quantity expressed as a percentage excess, $S_{dr} = (\text{rugosity} - 1) \times 100\%$. The root-mean-square slope of the surface, S_{dq}, is computed from the same squares.

Rugosity, S_{dr} and S_{dq} all depend strongly on the grid spacing. Measurement noise adds small random slopes to every cell, and the finer the grid, the steeper those slopes become, so these parameters grow as the grid is refined even if the surface is unchanged. Results computed with different grid spacings are therefore not comparable, and all scans in this work are processed with the same spacing of 2mm.

6.3.6 Spatial and spectral parameters

The height parameters describe how tall features are, but not how large or how far apart they are. Two further groups of parameters capture this.

Autocorrelation. The autocorrelation function measures how similar the surface is to a copy of itself shifted horizontally. It equals 1 at zero shift and decays as the shift grows. The autocorrelation length S_{al} is the shortest

shift at which it falls below 0.2, and gives the typical horizontal size of the features. The texture aspect ratio S_{tr} is the ratio of the shortest to the longest such distance over all directions: it is close to 1 for a surface that looks the same in every direction, such as a randomly burrowed bed, and close to 0 for a strongly directional one, such as a rippled bed. The autocorrelation is computed with the fast Fourier transform, counting only pairs of valid cells.

Power spectrum. The power spectral density (PSD) breaks the surface down into sinusoidal components and shows how much height variance each wavelength contributes. A Hann window is applied before the transform to reduce artefacts from the edges of the patch. The strongest peak of the PSD gives the dominant wavelength and direction of any regular pattern, such as ripples, and the texture direction S_{td} is reported as the direction of the crests. Many natural surfaces are *self-affine*: their PSD follows a power law $C(q) \propto q^{-2(1+H)}$ in wavenumber q over a wide range of scales [**?**]. Fitting a straight line to the log-log PSD gives the Hurst exponent H, which describes how roughness changes with scale, and the fractal dimension $D_f = 3 - H$. The fit is restricted to wavelengths between four grid cells, below which gridding distorts the spectrum, and a quarter of the patch width, above which there are too few repetitions to estimate the power reliably.

6.3.7 Additional measures and uncertainty

Several further outputs support the interpretation of the main parameters:

- **Multi-scale roughness.** S_q and rugosity are computed in square windows of increasing size, each with its own plane removed, to show how roughness depends on the scale of observation.
- **Confidence intervals.** 95% confidence intervals for S_q, S_{dr} and S_{al} are estimated with a block bootstrap: the surface is divided into square windows, which are resampled with replacement 1000 times. Neighbouring windows are correlated, so these intervals are likely to be somewhat optimistic.
- **Terrain ruggedness.** The Terrain Ruggedness Index [**?**] and the Vector Ruggedness Measure [**?**], widely used for seabed and landscape habitat analysis, are included for comparison with the ecological literature.
- **Noise correction.** Given an estimate of the sensor noise,either as a per-point standard deviation or from a scan of a flat plate processed in the same way, the noise contribution is subtracted from S_q and S_{dq}, assuming the noise adds independently to the true surface. A warning is raised if noise accounts for more than half of the measured slope, since rugosity then mainly measures noise.

All random choices use a fixed seed, so repeated runs give identical results, and every run records the hash of the code and of the input scan.

6.3.8 Validation

The program was checked in two ways. A test suite generates synthetic surfaces with known answers, such as sinusoids, for which S_q, S_a, the wavelength and even the exact rugosity follow in closed form, and random fractal surfaces with a prescribed Hurst exponent. Each metric is confirmed to recover the expected value, including when trends, noise, holes and a tilted camera are added. Secondly, a scan of a test object can be compared directly with its analytic definition: the reference surface is sampled, passed through exactly the same levelling, detrending and gridding, and the two sets of metrics are compared. This measures the error of the complete capture, fusion and roughness chain, and is used in chapter 7.

6.3.9 Limitations

The parameters follow ISO 25178-2 definitions but are not fully ISO-conformant. The standard defines them on a surface that has been band-limited with Gaussian filters, whereas here the band is set implicitly by the grid cell, the MLS smoothing in fusion, the polynomial degree and the size of the patch. The values are therefore directly comparable between scans processed in this pipeline with the same settings, but should be compared with values from commercial surface-metrology software with care. The surface is also treated as a single-valued height map, so overhangs are lost.

Chapter 7

Experimental Validation

In order to validate the processing pipeline as a whole, two methods are used. Firstly the pipeline is applied to experimental data captured during the characterisation stage of the known relief objects. Whilst this does not acurately describe the sediment surface in situ, it allows for an analytical evaluation the pipeline, whilst providing some resemblance for the true conditions.

Secondly, scans of the surface of sand, taken from the langebaan lagoon, was performed. This was taken for dry, wet and submerged sand, in both flat and shaped forms. This experimentation provides little analytical results, but helps to get a qualitative understanding of the performance in conditions closer to environmental.

7.1 Controlled Targets

7.1.1 Process

The rugosity estimation pipeline is applied to scans taken of reference objects shown in chapter C. These scans are re-used from the characterisation process chapter 5. These scans had the filters identified in chapter 5 applied to them. Cropping and fusion was then applied. This reconstructed surface was then cropped down to just the reference object, and the evaluation was performed on it. The results of which are then compared to the results obtained via analytical evaluation of the reference object file .stl. This provides an analytical comparison between the results of scanning and processing, and the known relief.

7.1.2 Results

Each smooth shape (big egg box and small egg box) were processed using the rugosity_metrics processing pipeline compared against their respective reference .stl files. The evaluation was performed on both fused and single scan results, and each metric applied all four detrending methods (planar, quadratic, cubic and quartic). These rugosity metrics are shown in Table 7.1 and Table 7.2, which are in turn plotted in Figure 7.1 and Figure 7.2.

There are two main points to be discussed with these results. Firstly is the effects of the different detrending results. The results show that planar and quadratic detrending methods tend to produce results that better match the ground truth. Care must be taken here since this analysis applies the same detrend to the ground truth. Thus the planar detrending should be the best representation of the ground truth rugosity. The primary reason for planar and quadraticdetrending giving the best performance can be seen in Figure 7.1, where the lower wavelength shape closely alignes with the higher order detrending methods, and thus they tend to remove this information. These can be visually seen in chapter F where the cubic and quartic detrending methods tend to remove actual shape of the big egg box surface, rather than underlying trends.

Secondly, and most importantly, is that these results are innacuratly representing the truth. The plots show that

single scans generally produce rugosity metrics that are more closely aligned with the reernce objects. This is contrary to what is expected, as the fused results generally provide a closer representation to the ground truth. The fundamental reason for this discrepency is actually a misleading result. The reconstrucitons for these reference objects are not completely accurate, and all feature some scale error, and underrepresetnt the amplitude of the reference waveforms. This means that the reconstructions will always report a lower rugosity compare to what the true waveforms have. This is where the inflated single scan performance arives. Since the single scans have a higher depth noise, this noise arteficially increases the roughness of the surface. This can be seen in Figure 5.11a and Figure 5.10a. This arteficial increase in the roughness, results in a higher rugosity, which results in an arteficially better result.

Table 7.1: Each value is over n = 1 rows.

Detrend Method	Scan	Rugosity mean ± SD	Ref rugosity mean ± SD
cubic	Big Egg Box Fused	1.19	1.27
	Big Egg Box Single	1.26	1.26
planar	Big Egg Box Fused	1.25	1.34
	Big Egg Box Single	1.31	1.34
quadratic	Big Egg Box Fused	1.25	1.34
	Big Egg Box Single	1.31	1.34
quartic	Big Egg Box Fused	1.19	1.27
	Big Egg Box Single	1.26	1.26

Table 7.2: Each value is over n = 1 rows. Fixed: Grid nx = 45.

Detrend Method	Scan	Rugosity mean ± SD	Ref rugosity mean ± SD
cubic	Small Egg Box Fused	1.26	1.69
	Small Egg Box Single	1.31	1.69
planar	Small Egg Box Fused	1.26	1.69
	Small Egg Box Single	1.31	1.69
quadratic	Small Egg Box Fused	1.26	1.69
	Small Egg Box Single	1.31	1.69
quartic	Small Egg Box Fused	1.25	1.69
	Small Egg Box Single	1.31	1.69

7.2 Controlled Laboratory Sand

Brining the Beach to the Lab

Whilst this reference object comparison provides an analytical evaluation, it does not align with real world conditions. Due to time constraints, real world implementation was not able to be performed, however a more realistic scan was taken, using sand collected from the langebaan lagoon. Approximately 13 kg of sand was collected dry, and was placed in a a 10 l plastic container. This was then flattened to produce a surface that should give rugosity metrics of 1. next, the sand was indented with the known reference shapes, and scans were taken. Next the sand was wetted with 2 kg of water. This approximates the wet intertidal sediment. The sand was then flattened again for scans, and then indented with the known reference shapes.

Finally, for purely qualitative performance, slightly more water was added, enough to produce pools of water on the surface. The sand was then shaped by hand to have some approximate natural texture. This was then scanned. This mainly gives an estimation of the performance on more water logged regions of the sediment, and also gives an estimate on what rugosity measurements would look.

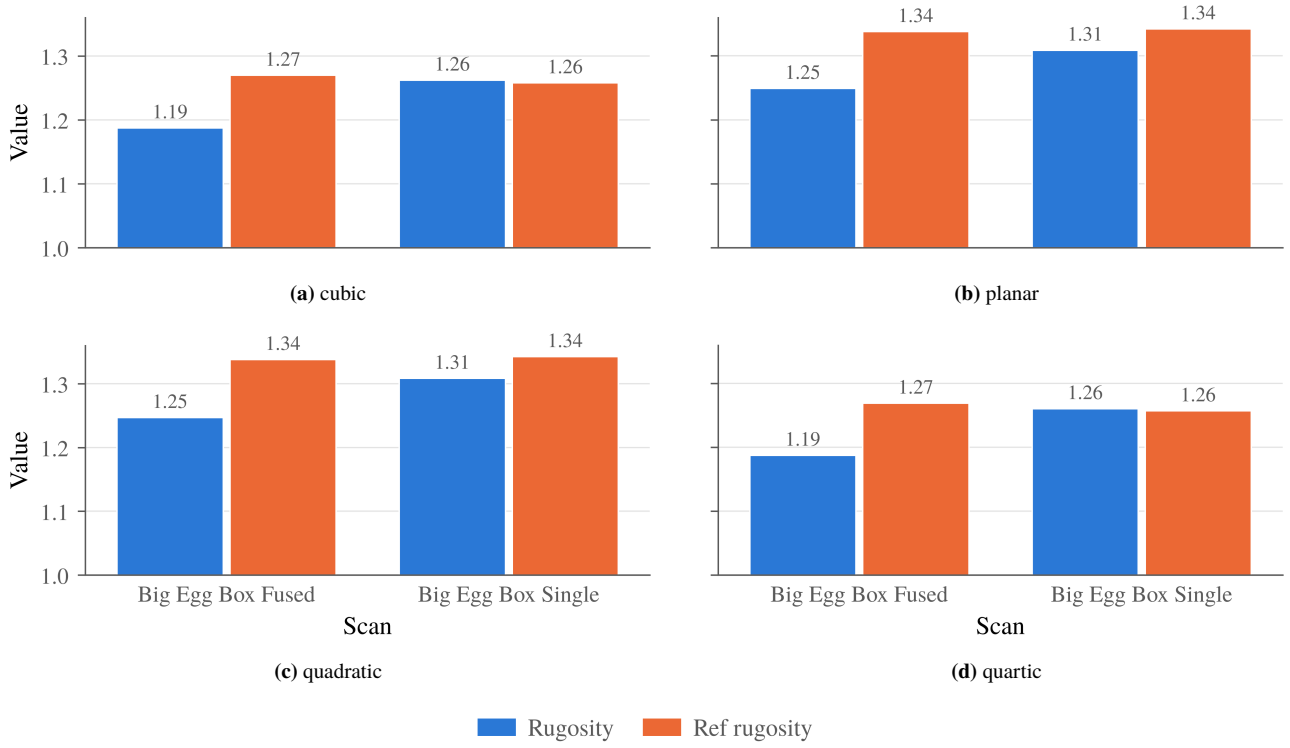

(a) cubic **(b)** planar **(c)** quadratic **(d)** quartic

Figure 7.1: Mean ± SD of n = 1 rows per bar.

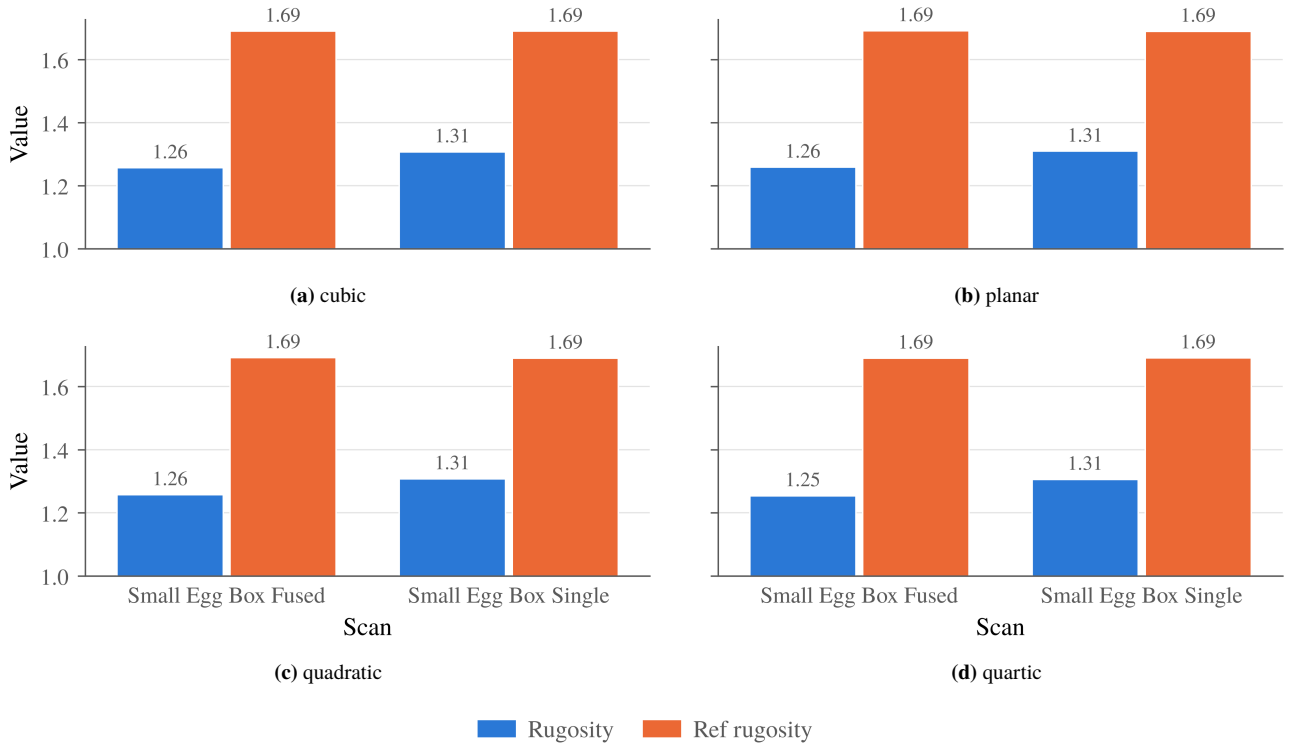

(a) cubic **(b)** planar **(c)** quadratic **(d)** quartic

Figure 7.2: Mean ± SD of n = 1 rows per bar. Fixed: Grid nx = 45.

7.2.1 Results

These results show a much better trend in the performance of the system. As can be seen in the dry flat analysis, the fused scan results in a rugosity of 1.00340, which is lower than that of the unfused scan which gives a rugosity of 1.01933. This is primarily due to the fact that single scans have a higher depth noise, and on flat planar surfaces, this raises the rugosity beyond the actual natural sediment. This is what is expected, and provdes a better representation than the referenc object validation.

Unfortunately these results dont have ground truth, since they were taken from sediment that was manually flattened and shaped. However te results do show that planar surfaces result in rugosirt that is far closer to 1 than shaped surfaces, which is expected since a flat surface would result in a rugosity of exactly 1.

An interesesting results can be seen with the soaked scan. This scan is the most representative of what sediment in situ might represetnt, and the scans show that system is able to accurately represent sediment shape, even when partially submerged by water. Providing useful insigt to show the robusteness to saturated sediment.

antoher intereting metric can be seen with the wet planar surface. In the midle of this surface is a large region of high noise. This is atributed to high specular reflections of the ir projector on the predominatly flat surface, with a water film. This shows that performance can degrade under such dconditions where the projector pattern can be reflected directly back into the ir sensors. Therefore this shows that care must be taken to prevent such specular reflections.

Table 7.3: Each value is over n = 4 rows. Averaged over: Detrend (4), Primary (2), Flags (3), Str bounded (2), Detrend n coef (4), Detrend order (4) and 1 more.

Label	Rugosity mean ± SD
Dry Flat	1.00340 ± 0.00021
Dry Flat Single	1.01933 ± 0.00018
Dry Shaped	1.05806 ± 0.00037
Dry Shaped Single	1.07912 ± 0.00036
Soaked	1.02659 ± 0.00104
Soaked Single	1.02922 ± 0.00077
Wet Flat	1.00931 ± 0.00013
Wet Flat Single	1.02429 ± 0.00014
Wet Shaped	1.05924 ± 0.00023
Wet Shaped Single	1.08549 ± 0.00011

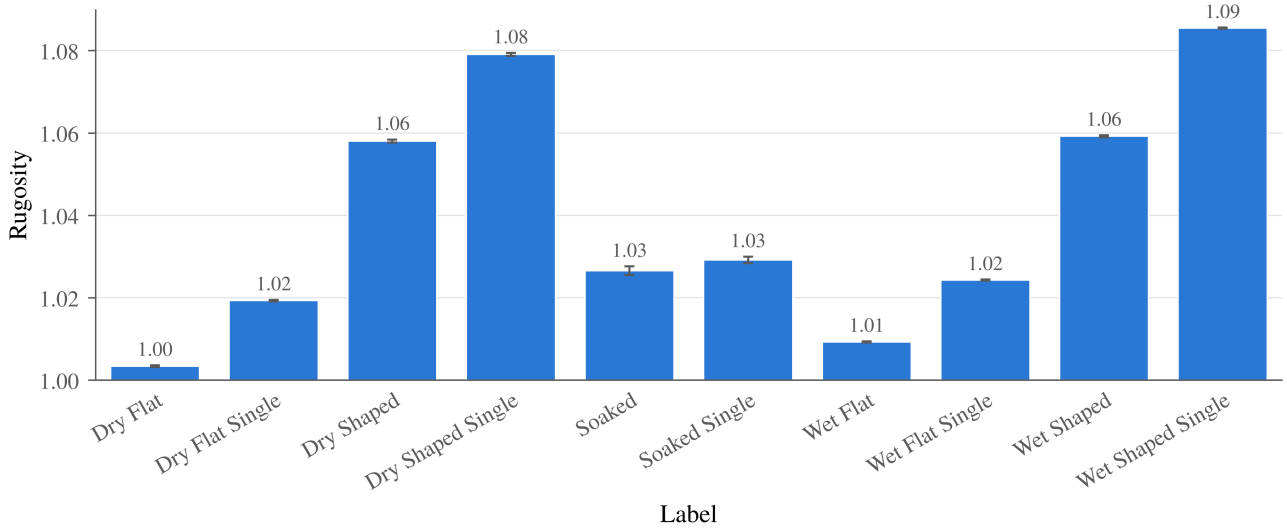

Figure 7.3: Mean ± SD of n = 4 rows per bar. Averaged over: Detrend (4), Primary (2), Flags (3), Str bounded (2), Detrend n coef (4), Detrend order (4) and 1 more.

Chapter 8

Discussion

This chapter summarises the key findings of chapter 5 and chapter 7, and relates them to the overarching project and the literature. Two findings recur throughout. First, reconstruction accuracy is set far more by geometry, meaning the working distance and the spread of viewpoints, than by the camera's software settings. Second, a single scan systematically overstates surface roughness, because its depth noise is measured as relief.

8.1 Camera Configuration

Of the stream settings, resolution had by far the largest effect (Figure 5.4). Moving from the default 848×480 to 1280×720 reduced the RMSE of the large egg box from 2.39 mm to 2.07 mm (13.4 %). Over the same change, completeness at 2 mm rose from 47.5 % to 57.1 %. This follows from Equation 3.4, since the focal length in pixels scales with resolution. The cost is that the minimum depth $Z_{\min} = fB/126$ also grows, to 0.52 m at 1280×720 [11].

The High Density preset performed best overall, with an egg box F-score of 59.2 % against 51.4 % for Default (Figure 5.6). High Accuracy gave among the lowest RMSE values (1.52 to 1.67 mm), but at much lower completeness. Since fusion can fill holes from other views, High Accuracy combined with fusion warrants testing, which is discussed in chapter 10.

The depth unit changed RMSE by at most 0.12 mm (Figure 5.5). This is expected, since even a 1 mm quantisation step adds only about $1/\sqrt{12} \approx 0.29$ mm of noise, small compared with the 1.2 to 4.6 mm of sensor noise measured on planes. The 0.1 mm unit costs nothing within the working range and can help with small scale relief and slight planar trends, which is why it is still recommended.

Post-processing filters improved RMSE by at most 0.24 mm (Figure 5.7). Raising the camera from 641 mm to 933 mm, by comparison, worsened RMSE by 1.13 to 1.24 mm. Strong spatial smoothing was determined harmful, because it flattens the relief being measured. In the characterisation, filters were applied offline to recorded scans. Recording unfiltered data in the field and filtering afterwards therefore keeps every option open and reduces the capture processing load on the host device.

8.2 Working Distance, Viewing Angle and Lighting

Working distance gave the most consistent trend in the report. Between 604 mm and 903 mm, the plane-fit RMS rose from 1.22 mm to 4.58 mm and the fraction of valid pixels fell from 99.6 % to 96.6 % (Figure 5.1). On the reference objects, the fused large egg box fell from an F-score of 90.9 % at 641 mm to 55.0 % at 933 mm (Figure 5.8). This agrees with the stereo noise model of Keselman et al. [12] and with other depth cameras [25, 31]. Over this narrow range, however, the trend is near linear rather than quadratic.

Height therefore trades directly against coverage. The footprint grows with the square of height, so moving

from 700 mm to 933 mm gives about 1.8 times the area. The D456 also cannot be lowered freely. Its 95 mm baseline is designed for longer range [11, 15], and tilting brings part of the view closer than the mount height. In initial experiments, depth anomalies appeared at 687 mm with tilts of around 17°, although fused results at 641 mm were still the best of any height. A height of about 700 mm is therefore recommended for accuracy. The maximum height should be used where area is needed, with at least one patch scanned at both heights so that height-dependent noise can be quantified.

Tilt degraded planar accuracy, with the plane-fit RMS rising from 1.62 mm to 2.29 mm at 28° pitch (Figure 5.2). On the 3D objects, by contrast, RMSE and completeness were generally best at the extreme tilts (Figure E.9 and Figure E.10). The two results are consistent. On a plane, tilt only increases obliquity, which raises noise [31]. On relief, tilt turns some slopes towards the camera at the expense of others. Each view therefore sees a different part of the surface well, which motivates fusion.

Ambient light raised the plane-fit RMS from 1.80 mm in darkness to 4.43 mm at 151 lx (Figure 5.3). Daylight was more damaging per lux than artificial light. The featureless board failed in sunlight, whereas textured ground returned complete data at 9 800 lx. This is the expected behaviour of active stereo, which degrades towards passive stereo on textured surfaces [12]. It agrees with outdoor tests of the D435 [26] and extends the indoor-only evaluation of [10]. The specular noise observed on wet flat sand (section 6.3) shows that water films remain a risk. The shading tarp should therefore be kept.

8.3 Fusion and Roughness Estimation

Fusing all views improved every metric at every height (Figure 5.8). For the large egg box, RMSE fell by 0.61 mm on average and the F-score rose by 14.0 percentage points. For the small egg box, RMSE fell by 0.25 mm and the F-score rose by 5.8 points. The fused surface is visibly smoother (Figure 5.11a and Figure 5.10a).

This extends the established benefits of fusion [32, 33] to close-range sediment-like relief. That was one of the gaps identified in section 2.6. On the small egg box, the gain is limited because its 25 mm features approach the camera's resolution and the 10 mm MLS radius. Fusion cannot recover relief that no single scan resolves.

The number of view points also matters less than their spread (Figure 5.9). Four rotations at tilts of 0° and $\pm$25° (12 scans) matched the full 44-view sweep, with an F-score of 0.73 in both cases. This makes multi-view capture feasible within a low-tide window.

On the reference objects, single scans gave rugosity closer to the reference than fused scans. With planar detrending, the large egg box gave 1.31 for the single scan, 1.25 for the fused scan and 1.34 for the reference (Table 7.1). This result is misleading as every reconstruction under-represents relief amplitude, as the negative signed bias shows. Rugosity scales with the squared slope ($S_{dr} \approx 50\, S_{dq}^2$ ‰), so this loss is amplified, most strongly on the steep small egg box (1.26 against 1.69). The single scan scores higher only because its noise adds slope variance that partly cancels the lost amplitude.

The laboratory sand confirms this. On flat dry sand, which should have a rugosity of 1, the fused scan gave 1.0034 and the single scan 1.0193 (Table 7.3). On shaped sand, single scans added 2.1 to 2.6 percentage points of excess area. This is roughly a third of the excess area that the fused scans measured on the same surfaces. The ± values in Table 7.3 reflect processing options rather than repeat captures, so repeatability remains unknown.

Planar and quadratic detrending preserved the reference relief, whereas higher orders removed part of it. This agrees with Azhar et al. [8]. As Jacobs et al. [24] caution, rugosity is therefore only comparable between scans with the same height, grid spacing, smoothing, detrending and views. It should also be referenced to a flat-surface noise floor.

8.4 Implications for the Overarching Project

Microtopography has explained a substantial share of benthic fluxes and biodiversity [1, 6]. It is only a useful proxy, however, if measured differences come from the sediment rather than the instrument. This project shows that single scans, or changes in height or processing between surveys, introduce instrument differences of the same order as the surface signal.

Fused scans captured with a fixed protocol make these effects small and consistent. That protocol is 1280×720 resolution, the High Density preset, a height of about 700 mm, four rotations at 0° and ±25°, and fixed processing settings. The large egg box, the closest analogue to the 3 to 5 cm high mounds of thalassinidean shrimps [4], was reconstructed best, with a fused F-score of 90.9 %. Finer features such as burrow openings remain near the system's limit. The main limitations are single trials per condition, uncalibrated light readings and the absence of field testing.

Chapter 9

Conclusions

This project set out to improve the system used to scan intertidal sediment for the wider study of sand prawn bioturbation on the Langebaan Lagoon. Bioturbators such as *Callichirus kraussi* reshape the sediment surface with mounds, depressions and burrow openings. Sediment microtopography, and rugosity in particular, has been shown to explain a substantial share of the variability in benthic fluxes and biodiversity [1, 6], which makes it a promising non-destructive proxy for prawn activity. That proxy is only useful if the measured roughness reflects the sediment rather than the instrument. The inherited system adapted the RGB-D framework of Azhar et al. [8] to an Intel RealSense D456 active stereo camera. However, its configuration had never been evaluated against known relief, and its repeated multi-view scans were each analysed on their own. This project addressed both problems. It characterised the D456 for close-range scanning of sediment-like surfaces, and it developed a cropping, fusion and roughness estimation pipeline that was validated on reference objects and on sand from the lagoon.

Characterisation of the D456

The first two objectives were met through planar tests and through scans of 3D-printed reference objects of known geometry. Across both, the results show that reconstruction accuracy is governed far more by geometry than by software settings.

Working distance had the most consistent effect. Between 604 and 903 mm, the planar fit RMS rose from 1.22 to 4.58 mm and the fraction of valid pixels fell from 99.6 to 96.6%. On the large egg box, the F-score at 2 mm of the fused reconstruction fell from 90.9% at 641 mm to 55.0% at 933 mm. Over this range the degradation was close to linear rather than the quadratic trend predicted by stereo theory, most likely because the range tested was narrow. Of the stream settings, resolution mattered most. Increasing it from the default 848×480 to 1280×720 reduced the RMSE of the large egg box by 13.4%, and raised completeness from 47.5 to 57.1%. The High Density preset gave the best overall F-score. The depth unit changed RMSE by at most 0.12 mm, and post-processing filters improved it by at most 0.24 mm. By comparison, raising the camera from 641 to 933 mm worsened RMSE by more than 1.1 mm. Filters therefore do not need to be applied during capture. Recording unfiltered data keeps every option open and lightens the load on the capture computer.

Tilt degraded planar scans, with the planar fit RMS rising from 1.62 to 2.29 mm at 28° of pitch. On the 3D reference objects, by contrast, single scans were often most accurate at the extreme tilts. Tilt increases obliquity on a plane, but on relief it turns some slopes towards the camera, so each tilted view sees part of the surface well. Ambient light also raised noise, from 1.80 mm in darkness to 4.43 mm at 151 lx indoors. Outdoors, the featureless white board failed completely in sunlight, while the textured ground still returned complete data at 9 800 lx. This shows that the D456 can operate in sunlight on textured surfaces, but it does not do so reliably. The shading tarp should therefore be kept, particularly since specular reflections from a water film produced a region of high noise on wet, flat sand.

Processing pipeline

The third objective was met with an automatic cropping tool and a multi-view fusion program. The cropping tool learns the layout of the four AprilTags from the scans of a session in which all four were seen. It then predicts missing corners from as few as two detected tags, and crops each scan with a vertical prism normal to the bed, so that every scan covers the same physical region regardless of tilt. It also records the 3D corners of each tag before removing them from the cloud. The fusion program combines these tag positions with IMU gravity readings for coarse alignment. On one test session, this reduced the disagreement with the final registered poses from 1–2.5° to 0.2–0.6°. The scans are then refined with point-to-plane ICP and reconciled with a confidence-weighted trimmed mean.

Fusing all views improved every metric at every height tested. For the large egg box, RMSE fell by an average of 0.61 mm and the F-score rose by 14.0 percentage points. For the small egg box, the gains were smaller (0.25 mm and 5.8 points), because its 25 mm features approach the resolution of the camera and the 10 mm smoothing radius of the fusion. Fused surfaces were also visibly smoother than single scans, with much less spurious undulation on the flat floor around the objects. The spread of viewpoints mattered more than their number. Four rotations at tilts of 0 and $\pm$25°, 12 scans in total, matched the F-score of 0.73 obtained from the full 44-view sweep. This makes multi-view capture practical within a low-tide window.

Roughness estimation and validation

The fourth objective was met with a roughness program that levels, detrends and grids the fused surface, and computes ISO 25178-2 areal parameters, including rugosity. On the reference objects, planar and quadratic detrending preserved the relief, while cubic and quartic fits removed part of the shape of the large egg box. At first glance, single scans gave rugosity closer to the reference than fused scans. For the large egg box with planar detrending, the single scan gave 1.31 and the fused scan 1.25, against a reference of 1.34. This result is misleading. Every reconstruction under-represents the amplitude of the relief, and because rugosity scales with the squared slope, this loss is amplified. The effect is strongest on the steeper small egg box, where the fused scan measured 1.26 against a reference of 1.69. The single scan scores higher only because its depth noise adds slope that partly offsets the lost amplitude.

The laboratory sand confirms this interpretation. Flat dry sand, which should have a rugosity of 1, measured 1.0034 when fused and 1.0193 as a single scan. On shaped sand, single scans added 2.1 to 2.6 percentage points of excess area, about a third of the excess area that the fused scans measured on the same surfaces. A single scan therefore overstates roughness by an amount comparable to the signal itself. The fused reconstruction also remained usable on soaked sand with standing water.

Implications and limitations

For the wider study, the main conclusion is that rugosity measured with this system is only comparable between surveys if the capture and processing protocol is fixed. The recommended protocol is 1280×720 resolution, the High Density preset, a working height of about 700 mm, four rotations at tilts of 0 and $\pm$25°, unfiltered capture with the shading tarp in place, and fixed fusion and roughness settings. The fused reconstruction of the large egg box, the closest analogue to the 3 to 5 cm mounds of thalassinidean shrimps, achieved an F-score of 90.9%. Finer features such as burrow openings remain near the limit of the system. The main limitations are that each condition was recorded only once, that lighting was measured with an uncalibrated smartphone sensor, that rugosity was not yet validated against a flat-surface noise floor, and that the system has not been tested in the field.

Overall, the project provides a characterised and documented operating configuration for the D456, together

with a cropping, fusion and roughness pipeline. The pipeline produces reconstructions that are more accurate and more complete than single scans, and roughness values that better reflect the sediment rather than the sensor.

Chapter 10

Further work

The 12-week timeframe limited how thoroughly parts of the system could be characterised and validated. This chapter identifies where the results are weakest and the work that would address them.

Illuminance was read with an uncalibrated smartphone sensor, and lux does not describe spectral content. Daylight degraded the depth data more per lux than fluorescent light, which is consistent with its larger near-infrared component, so the indoor results do not fully represent conditions in situ. The outdoor trials were also not directly comparable with the indoor ones, since the featureless foam board failed in direct sunlight and the brick floor had to be scanned instead. A repeat characterisation should use a calibrated light meter, ideally with a near-infrared or spectral measurement, and scan both a textureless board and a texture-rich surface under every condition, so that the effects of illumination and texture can be separated.

Ambient light also drifted during the long sweeps. Over the working-distance sweep (15:40 to 17:35) the auto-exposure value fell from about 5300 to between 1700 and 2050, so lighting is confounded with distance. The height and tilt sweeps should be repeated under controlled lighting, or in a randomised order.

Each condition was recorded once, so the repeatability of the set-up is unknown. The spread reported for the laboratory sand reflects processing options rather than repeat captures. Repeated captures of key conditions are needed, as is further testing on wet sand, where specular reflection from a water film produced a region of high noise. The system has also not been deployed on the Langebaan intertidal flats. Field trials within a low-tide window are needed to confirm the recommended protocol, and would provide the data on which to develop automated detection of mounds and burrow openings, which was removed from the scope of this project.

On the reference objects, single scans gave rugosity closer to the reference than fused scans, in conflict with every other metric. This is attributed to the reconstructions under-representing relief amplitude while single-scan noise adds slope, but the explanation has not been tested directly. A flat reference plate, processed with the same grid spacing and smoothing, would establish the noise floor and allow the pipeline's noise correction to be applied. Reference objects with lower-amplitude and shorter-wavelength relief would show where amplitude is lost, including the effect of the 10 mm MLS radius. Larger surfaces with verified roughness would give a more representative validation. The High Accuracy preset, which gave low RMSE but poor completeness in single scans, should also be tested in combination with fusion, which can fill its holes.

The D456 has a 95 mm baseline optimised for longer range [11]. Its minimum depth of about 0.52 m at 1280×720, together with the depth anomalies observed at 687 mm with tilts of around 17°, prevents it from operating at the shorter distances where noise is lowest. Shorter-baseline cameras are better suited to close range. The D415 has a minimum range of about 160 mm [14] and gave the best short-range reconstructions of the cameras compared by Servi et al. [30], while the D435 kept errors below 1 cm within 1 m [15]. These should be compared with the D456 on the same reference objects. Finally, an external projector can add texture beyond the camera's internal pattern [13], which may improve reconstruction of featureless or wet surfaces, but this could not be investigated within the timeframe.

Appendix A

GitHub repository

This appendix provides a link to the GitHub repository where the code and additional materials for this final year project report can be found.

A.1 Title for appendix A.1

- GitHub Repository: Repo name

Appendix B

Theory

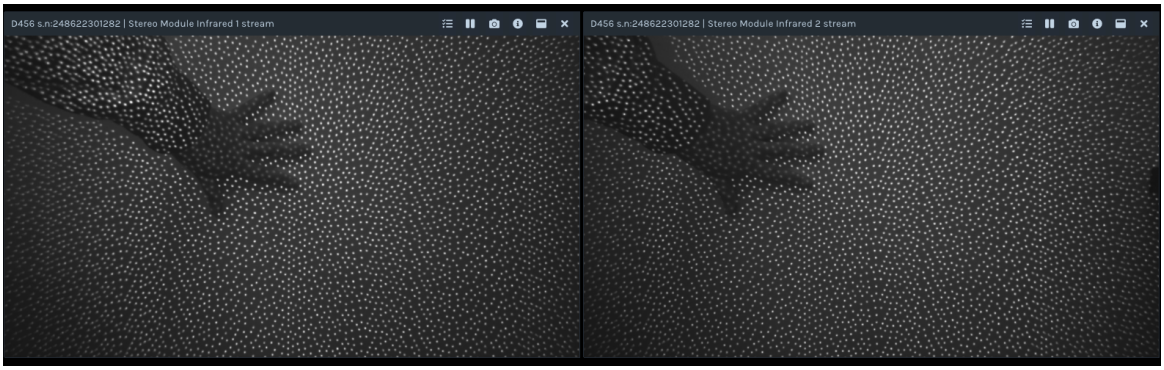

Figure B.1: D456 Semi-Random IR Dot Pattern Projected Onto a White Wall, With a Hand

Appendix C

Reference Objects

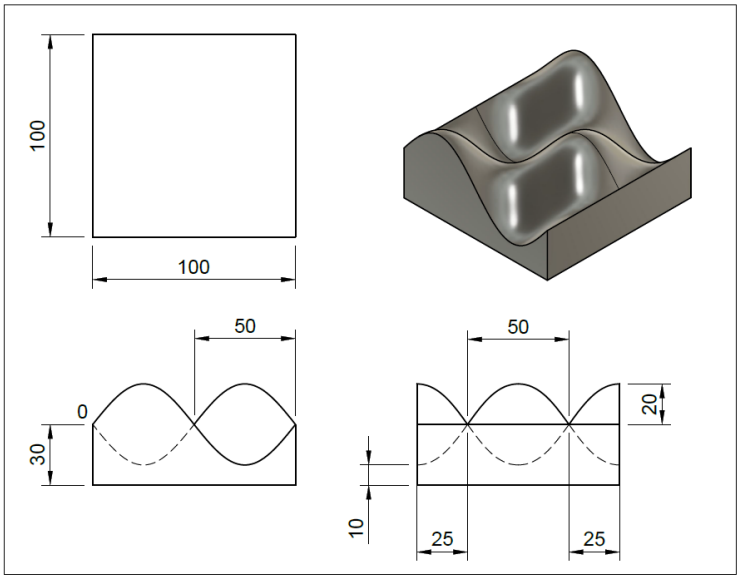

Figure C.1: Large Egg Box Surface Reference Object.

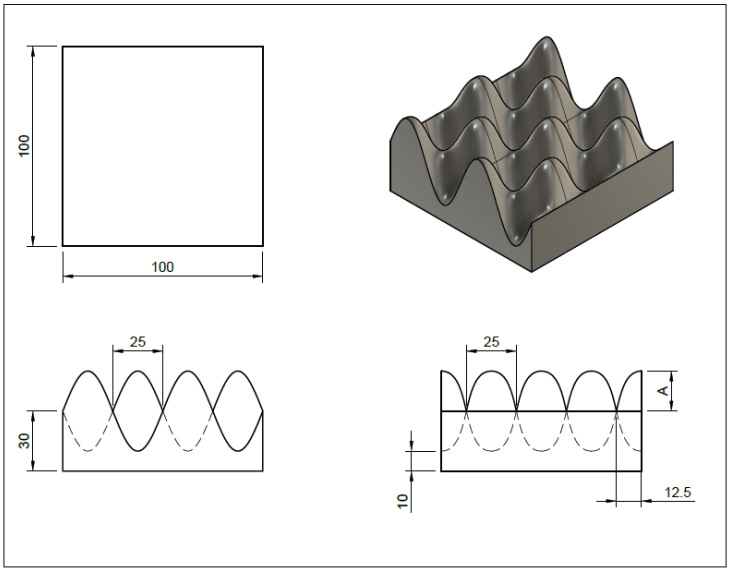

Figure C.2: Small Egg Box Surface Reference Object, Printed with amplitude A = 15mm.

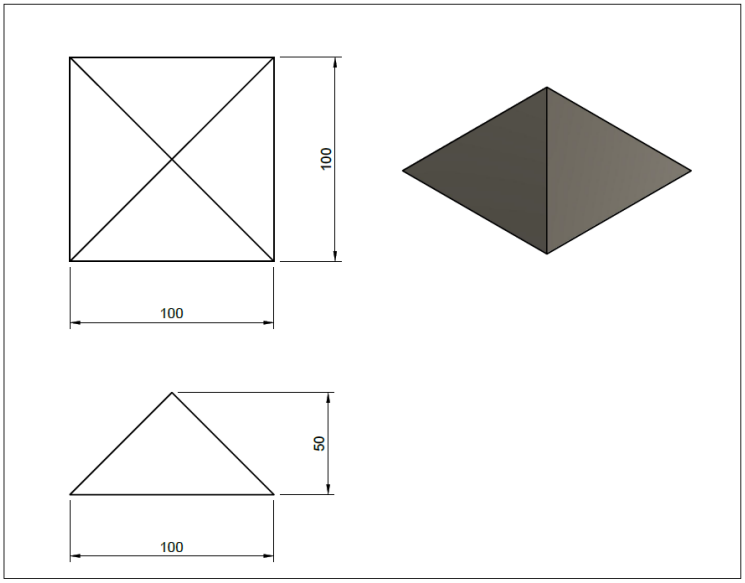

Figure C.3: Large Pyramid Surface Reference Object.

Appendix D

Characterisation Results

D.1 Planar Characterisation Results

Table D.1: Results for Planar Characterisation Against Distance. Per-Trial Values Over n = 385-552 Frames.

Distance along optical axis (mm)	Median depth error (mm) per-trial median [Q1, Q3]	Depth RMSE (mm) per-trial median [Q1, Q3]	Plane-fit RMS (mm) per-trial median [Q1, Q3]	Valid pixels (%) per-trial median [Q1, Q3]
604	4.70 [4.70, 4.70]	4.90 [4.84, 4.96]	1.22 [1.21, 1.25]	99.6 [99.6, 99.7]
628	4.60 [4.60, 4.80]	4.86 [4.81, 4.90]	1.29 [1.27, 1.31]	99.5 [99.5, 99.6]
667	6.40 [6.40, 6.60]	6.64 [6.58, 6.70]	1.56 [1.53, 1.59]	99.5 [99.4, 99.5]
723	5.60 [5.30, 5.60]	6.13 [6.04, 6.23]	2.64 [2.57, 2.70]	97.7 [97.5, 98.0]
760	6.10 [5.80, 6.10]	6.66 [6.54, 6.76]	2.81 [2.73, 2.88]	97.7 [97.5, 97.9]
821	7.80 [7.50, 7.80]	8.35 [8.21, 8.47]	3.27 [3.18, 3.35]	97.4 [97.1, 97.6]
880	7.30 [6.90, 7.30]	8.37 [8.21, 8.54]	4.05 [3.95, 4.17]	96.9 [96.6, 97.2]
903	8.10 [7.70, 8.10]	9.41 [9.20, 9.62]	4.58 [4.45, 4.70]	96.6 [96.2, 96.9]

Table D.2: Results for Planar Characterisation Against Scan Angle. Per-trial values over n=527-574 frames.

AbsMean IMU angle	Axis	Median depth error (mm) per-trial median [Q1, Q3]	Depth RMSE (mm) per-trial median [Q1, Q3]	Plane-fit RMS (mm) per-trial median [Q1, Q3]	Valid pixels (%) per-trial median [Q1, Q3]
1.7	pitch	5.0 [5.0, 5.1]	5.4 [5.3, 5.5]	1.62 [1.57, 1.66]	99.4 [99.3, 99.5]
2.5	roll	5.8 [5.7, 5.9]	6.2 [6.1, 6.3]	1.69 [1.65, 1.74]	99.2 [99.1, 99.3]
2.8	pitch	6.2 [6.1, 6.3]	6.8 [6.7, 6.8]	1.60 [1.55, 1.65]	99.5 [99.4, 99.6]
4.4	pitch	5.9 [5.8, 6.0]	6.1 [6.0, 6.2]	1.59 [1.55, 1.64]	99.4 [99.3, 99.5]
5.0	roll	5.5 [5.4, 5.6]	5.9 [5.8, 6.0]	1.67 [1.62, 1.72]	99.2 [99.1, 99.4]
6.2	roll	6.5 [6.4, 6.6]	7.0 [6.9, 7.1]	1.69 [1.64, 1.73]	99.2 [99.0, 99.3]
7.0	pitch	6.1 [6.0, 6.2]	6.4 [6.3, 6.4]	1.63 [1.59, 1.68]	99.5 [99.3, 99.6]
8.8	roll	9.8 [9.7, 9.9]	10.9 [10.8, 11.0]	1.72 [1.67, 1.77]	99.2 [99.1, 99.4]
9.3	pitch	6.7 [6.6, 6.8]	7.0 [7.0, 7.1]	1.63 [1.59, 1.68]	99.4 [99.3, 99.5]
10.8	pitch	6.6 [6.5, 6.7]	7.0 [6.9, 7.0]	1.68 [1.63, 1.73]	99.4 [99.3, 99.5]
11.0	roll	8.2 [8.1, 8.3]	8.7 [8.6, 8.8]	1.73 [1.69, 1.77]	99.3 [99.1, 99.4]
13.1	pitch	6.6 [6.5, 6.7]	7.0 [6.9, 7.1]	1.74 [1.69, 1.79]	99.3 [99.2, 99.4]
13.5	roll	7.4 [7.3, 7.5]	7.8 [7.7, 7.9]	1.79 [1.74, 1.83]	99.2 [99.1, 99.3]
15.8	roll	7.9 [7.8, 8.0]	8.4 [8.3, 8.5]	1.79 [1.75, 1.84]	99.2 [99.1, 99.3]
15.8	pitch	7.9 [7.8, 8.0]	8.5 [8.4, 8.5]	1.79 [1.74, 1.84]	99.3 [99.2, 99.5]
17.8	roll	8.9 [8.8, 9.0]	9.4 [9.3, 9.5]	1.87 [1.81, 1.92]	99.1 [98.9, 99.3]
17.8	pitch	7.4 [7.3, 7.5]	8.0 [7.9, 8.1]	1.91 [1.85, 1.97]	99.2 [99.0, 99.4]
18.8	roll	6.5 [6.4, 6.6]	7.1 [7.0, 7.2]	1.88 [1.83, 1.93]	99.1 [98.9, 99.2]
20.1	pitch	7.5 [7.4, 7.6]	8.1 [8.0, 8.2]	2.02 [1.96, 2.09]	99.0 [98.8, 99.2]
21.9	roll	10.9 [10.8, 11.0]	11.4 [11.3, 11.5]	1.96 [1.91, 2.02]	98.9 [98.8, 99.1]
22.0	pitch	7.8 [7.7, 7.9]	8.1 [8.0, 8.2]	2.01 [1.94, 2.08]	99.1 [98.9, 99.2]
23.8	roll	9.9 [9.7, 10.0]	10.3 [10.2, 10.4]	2.05 [2.00, 2.10]	98.7 [98.5, 98.9]
24.1	pitch	8.8 [8.7, 8.9]	9.2 [9.1, 9.4]	2.11 [2.03, 2.18]	98.9 [98.7, 99.1]
25.5	roll	11.1 [11.0, 11.2]	11.5 [11.4, 11.6]	2.11 [2.06, 2.18]	98.4 [98.2, 98.7]
26.5	pitch	5.0 [4.9, 5.1]	5.4 [5.3, 5.6]	2.21 [2.14, 2.29]	98.8 [98.6, 99.0]
27.9	pitch	10.8 [10.6, 10.9]	11.4 [11.3, 11.6]	2.29 [2.22, 2.37]	98.8 [98.5, 99.0]

Table D.3: Results for Planar Characterisation Against Light Intensity. Per-trial values over n = 592-599 Frames.

Illuminance (lx)	Median depth error (mm) per-trial median [Q1, Q3]	Depth RMSE (mm) per-trial median [Q1, Q3]	Plane-fit RMS (mm) per-trial median [Q1, Q3]	Valid pixels (%) per-trial median [Q1, Q3]
0	4.80 [4.80, 5.20]	6.22 [5.97, 6.41]	1.80 [1.75, 1.84]	99.9 [99.9, 100.0]
17	7.50 [7.00, 7.50]	8.01 [7.76, 8.20]	2.74 [2.64, 2.82]	99.4 [99.3, 99.5]
65	5.20 [5.20, 5.60]	6.81 [6.67, 6.97]	2.84 [2.75, 2.94]	99.3 [99.2, 99.5]
78	7.50 [7.00, 7.50]	8.29 [8.04, 8.50]	3.44 [3.33, 3.56]	98.8 [98.5, 98.9]
91	7.00 [7.00, 7.50]	8.41 [8.17, 8.65]	4.14 [4.01, 4.29]	97.0 [96.7, 97.4]
151	7.50 [7.50, 7.90]	8.83 [8.60, 9.06]	4.43 [4.29, 4.59]	96.3 [95.9, 96.8]

D.2 Camera Configuration Results

Table D.4: Each value is over n = 15 scans. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Depth unit (3), Preset (5), Preset actual (5).

Resolution	Object	Accuracy RMSE (mm) median [Q1, Q3]	Signed bias (mm) median [Q1, Q3]	Completeness (full reference) at 2 mm (%) median [Q1, Q3]	F-score at 2 mm (%) median [Q1, Q3]
640 × 360	Big pyramid	2.14 [1.99, 2.19]	−0.013 [−0.019, −0.008]	50.7 [47.4, 52.7]	56.1 [52.9, 58.5]
	Big Egg Box	3.20 [2.86, 3.28]	−0.466 [−0.547, −0.396]	28.0 [25.5, 31.9]	33.9 [32.2, 37.6]
848 × 480	Big pyramid	1.70 [1.61, 1.77]	−0.001 [−0.007, 0.001]	65.7 [62.2, 68.5]	69.9 [67.2, 72.7]
	Big Egg Box	2.39 [2.34, 2.52]	−0.222 [−0.253, −0.115]	47.5 [45.0, 50.9]	52.3 [51.8, 55.4]
1280 × 720	Big pyramid	1.47 [1.34, 1.57]	−0.017 [−0.027, −0.002]	77.5 [69.6, 79.1]	79.9 [76.2, 81.0]
	Big Egg Box	2.07 [2.00, 2.21]	−0.100 [−0.118, −0.068]	57.1 [54.9, 61.4]	62.1 [59.8, 64.2]

Table D.5: Each value is over n = 15 scans. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Resolution (3), Preset (5), Preset actual (5).

Depth unit (mm)	Object	Accuracy RMSE (mm) median [Q1, Q3]	Signed bias (mm) median [Q1, Q3]	Completeness (full reference) at 2 mm (%) median [Q1, Q3]	F-score at 2 mm (%) median [Q1, Q3]
0.1	Big pyramid	1.62 [1.52, 2.03]	−0.016 [−0.021, −0.001]	64.6 [51.5, 70.3]	71.2 [57.3, 76.2]
	Big Egg Box	2.35 [2.26, 2.85]	−0.154 [−0.399, −0.113]	45.8 [31.6, 54.9]	52.9 [37.5, 58.7]
0.5	Big pyramid	1.70 [1.55, 1.85]	−0.006 [−0.012, 0.003]	62.2 [56.8, 75.2]	69.9 [63.2, 77.9]
	Big Egg Box	2.39 [2.15, 2.86]	−0.230 [−0.340, −0.101]	47.9 [31.6, 54.4]	52.6 [37.3, 61.2]
1	Big pyramid	1.74 [1.56, 2.01]	−0.008 [−0.023, −0.006]	62.9 [52.9, 73.4]	68.8 [58.8, 75.5]
	Big Egg Box	2.47 [2.23, 2.74]	−0.269 [−0.372, −0.104]	47.5 [35.8, 56.6]	51.8 [41.5, 60.8]

Table D.6: Each value is over n = 8–10 scans. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Resolution (3), Depth unit (3), Preset (5).

Preset actual	Object	Accuracy RMSE (mm) median [Q1, Q3]	Signed bias (mm) median [Q1, Q3]	Completeness (full reference) at 2 mm (%) median [Q1, Q3]	F-score at 2 mm (%) median [Q1, Q3]
Custom	Big pyramid	1.70 [1.56, 1.83]	−0.011 [−0.020, −0.001]	64.6 [56.4, 70.9]	69.9 [62.7, 75.2]
	Big Egg Box	2.43 [2.27, 2.71]	−0.232 [−0.383, −0.135]	48.0 [38.3, 53.5]	52.7 [43.8, 58.7]
Default	Big pyramid	1.88 [1.69, 2.14]	0.001 [−0.018, 0.005]	58.7 [48.2, 74.6]	63.2 [52.9, 75.1]
	Big Egg Box	2.51 [2.29, 2.99]	−0.353 [−0.388, −0.122]	46.7 [29.6, 56.3]	51.4 [35.2, 59.1]
High Accuracy	Big pyramid	1.74 [1.33, 2.14]	−0.017 [−0.019, −0.014]	59.9 [48.6, 62.2]	67.8 [54.8, 71.2]
	Big Egg Box	2.36 [2.00, 3.29]	−0.098 [−0.466, −0.060]	44.0 [25.1, 51.9]	51.8 [32.5, 60.1]
High Density	Big pyramid	1.60 [1.55, 1.86]	−0.007 [−0.014, −0.004]	72.3 [57.2, 78.8]	75.9 [63.6, 79.9]
	Big Egg Box	2.28 [2.05, 2.64]	−0.109 [−0.248, −0.104]	55.0 [36.2, 62.0]	59.2 [41.9, 64.4]
Medium density	Big pyramid	1.65 [1.48, 1.97]	−0.007 [−0.013, −0.001]	67.1 [53.6, 77.5]	71.2 [59.5, 80.2]
	Big Egg Box	2.59 [2.21, 3.20]	−0.236 [−0.493, −0.135]	47.1 [28.1, 55.0]	51.8 [33.9, 58.4]

D.3 Filter Configuration Results

Table D.7: Each value is over n = 4–5 scans. Fixed: Tilt = 0; Spatial filter on = Off; HDR merge on = Off; Status = ok. Median taken over: Rotation (4).

Object	Height (mm)	Configuration	Accuracy RMSE (mm) median [Q1, Q3]	Signed bias (mm) median [Q1, Q3]	Completeness (full reference at 2 mm (%) median [Q1, Q3
Big Egg Box	641	Baseline (no filters)	1.71 [1.67, 1.80]	−0.060 [−0.073, −0.053]	70.6 [68.0, 70.6
		Best overall	1.71 [1.70, 1.76]	−0.086 [−0.108, −0.086]	68.9 [68.9, 69.9
		Best for this shape & height	1.68 [1.66, 1.70]	−0.073 [−0.078, −0.072]	70.3 [69.8, 70.6
	701	Baseline (no filters)	1.99 [1.96, 1.99]	−0.047 [−0.060, −0.012]	63.5 [61.3, 64.0
		Best overall	1.88 [1.88, 1.89]	−0.055 [−0.064, −0.044]	62.6 [61.9, 62.8
		Best for this shape & height	1.89 [1.88, 1.91]	−0.041 [−0.089, −0.028]	11.2 [11.1, 11.2
	763	Baseline (no filters)	2.14 [2.12, 2.25]	−0.194 [−0.226, −0.174]	58.9 [57.0, 62.0
		Best overall	2.09 [2.08, 2.10]	−0.171 [−0.236, −0.154]	60.9 [60.2, 61.0
		Best for this shape & height	2.06 [2.03, 2.09]	−0.237 [−0.237, −0.226]	52.9 [52.6, 52.9
	810	Baseline (no filters)	2.15 [2.15, 2.35]	−0.103 [−0.150, −0.044]	54.6 [54.5, 56.7
		Best overall	2.11 [2.10, 2.14]	−0.085 [−0.090, −0.082]	56.1 [55.9, 56.9
		Best for this shape & height	2.09 [2.08, 2.12]	−0.110 [−0.115, −0.104]	46.7 [45.6, 47.0
	863	Baseline (no filters)	2.64 [2.55, 2.76]	−0.121 [−0.160, −0.016]	49.8 [48.7, 50.7
		Best overall	2.40 [2.37, 2.46]	−0.202 [−0.202, −0.201]	51.6 [50.1, 52.9
		Best for this shape & height	2.37 [2.27, 2.42]	−0.141 [−0.232, −0.136]	54.1 [49.1, 54.7
	933	Baseline (no filters)	2.95 [2.84, 3.20]	−0.345 [−0.385, −0.341]	43.3 [37.5, 45.2
		Best overall	2.80 [2.78, 2.86]	−0.311 [−0.362, −0.303]	43.6 [40.1, 44.6
		Best for this shape & height	2.89 [2.73, 2.96]	−0.221 [−0.322, −0.088]	47.6 [42.6, 52.5
Small Egg Box	641	Baseline (no filters)	2.04 [2.01, 2.05]	−0.109 [−0.113, −0.105]	54.2 [53.9, 55.7
		Best overall	2.04 [2.01, 2.06]	−0.112 [−0.113, −0.110]	54.7 [54.6, 55.3
		Best for this shape & height	2.00 [1.97, 2.01]	−0.107 [−0.110, −0.103]	55.3 [53.9, 55.3
	701	Baseline (no filters)	2.47 [2.43, 2.49]	−0.106 [−0.109, −0.053]	44.9 [42.9, 45.4
		Best overall	2.43 [2.41, 2.44]	−0.092 [−0.102, −0.088]	45.2 [45.2, 45.7
		Best for this shape & height	2.39 [2.39, 2.41]	−0.085 [−0.097, −0.076]	46.1 [45.4, 46.2
	763	Baseline (no filters)	2.67 [2.67, 2.68]	−0.194 [−0.199, −0.190]	36.5 [36.3, 37.7
		Best overall	2.63 [2.61, 2.63]	−0.182 [−0.183, −0.182]	36.5 [36.4, 36.7
		Best for this shape & height	2.59 [2.59, 2.60]	−0.180 [−0.195, −0.177]	37.3 [36.8, 37.5
	810	Baseline (no filters)	2.76 [2.76, 2.83]	−0.216 [−0.248, −0.216]	35.0 [34.3, 35.2
		Best overall	2.77 [2.77, 2.78]	−0.199 [−0.203, −0.195]	34.3 [33.9, 34.3
		Best for this shape & height	2.75 [2.74, 2.75]	−0.194 [−0.199, −0.175]	34.9 [34.1, 36.0
	863	Baseline (no filters)	3.07 [3.05, 3.20]	−0.208 [−0.213, −0.144]	32.3 [32.1, 32.4
		Best overall	2.99 [2.99, 3.01]	−0.163 [−0.186, −0.158]	32.7 [32.3, 33.0
		Best for this shape & height	2.98 [2.97, 2.99]	−0.184 [−0.191, −0.176]	25.3 [25.0, 25.5
	933	Baseline (no filters)	3.17 [3.14, 3.24]	−0.263 [−0.271, −0.259]	27.3 [26.5, 28.6
		Best overall	3.03 [3.02, 3.04]	−0.263 [−0.271, −0.260]	28.5 [28.0, 28.6
		Best for this shape & height	3.03 [3.03, 3.03]	−0.258 [−0.261, −0.257]	28.5 [28.3, 28.5

D.4 Fusion Combination Results

Table D.8: Each value is over n = 1 point clouds. Fixed: Status = ok.

Object	Height (mm)	View family	Accuracy RMSE (mm) median [Q1, Q3]	Signed bias (mm) median [Q1, Q3]	Completeness (full reference) at 2 mm (%) median [Q1, Q3]	F-score at 2 mm (%) median [Q1, Q3]
Big Egg Box	641	all	1.15	−0.054	88.8	90.9
		single	1.86	−0.096	65.3	68.7
	701	all	1.34	−0.061	81.0	83.6
		single	1.94	−0.041	64.7	67.8
	763	all	1.57	−0.087	74.0	76.9
		single	2.33	−0.205	59.2	63.0
	810	all	1.76	−0.108	66.6	70.1
		single	2.11	−0.057	60.4	64.4
	863	all	1.89	−0.098	62.0	65.3
		single	2.24	−0.219	53.1	57.5
	933	all	2.31	−0.168	51.9	55.0
		single	3.20	−0.362	31.8	36.5
Small Egg Box	641	all	1.93	−0.082	54.4	60.8
		single	2.09	−0.081	53.8	59.8
	701	all	2.15	−0.083	48.9	55.4
		single	2.56	−0.127	40.4	46.8
	763	all	2.43	−0.072	43.5	50.3
		single	2.71	−0.192	34.9	42.0
	810	all	2.63	−0.113	39.1	46.0
		single	2.73	−0.189	34.4	41.9
	863	all	2.72	−0.080	38.6	45.4
		single	3.12	−0.153	31.5	39.1
	933	all	2.94	−0.105	35.1	41.7
		single	3.09	−0.253	28.4	35.4

Appendix E

Characterisation Plots

E.1 Filter Heatmaps

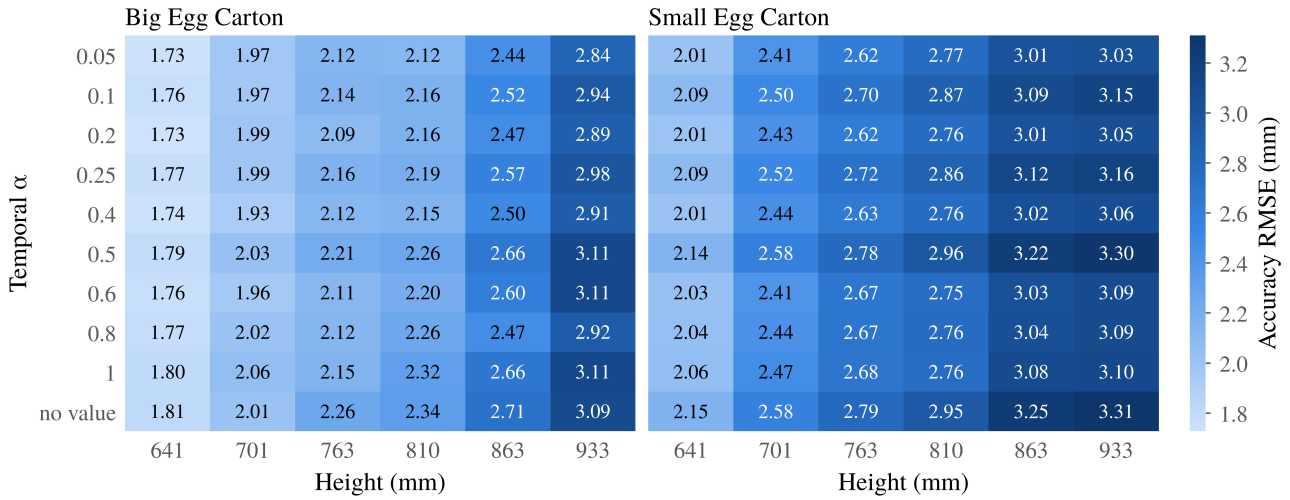

Figure E.1: Median with IRQ % CI of the mean of n = 5–295 scans per cell. Fixed: Tilt = 0. Averaged over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial filter on (2), Spatial magnitude (6), Spatial α (8) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded.

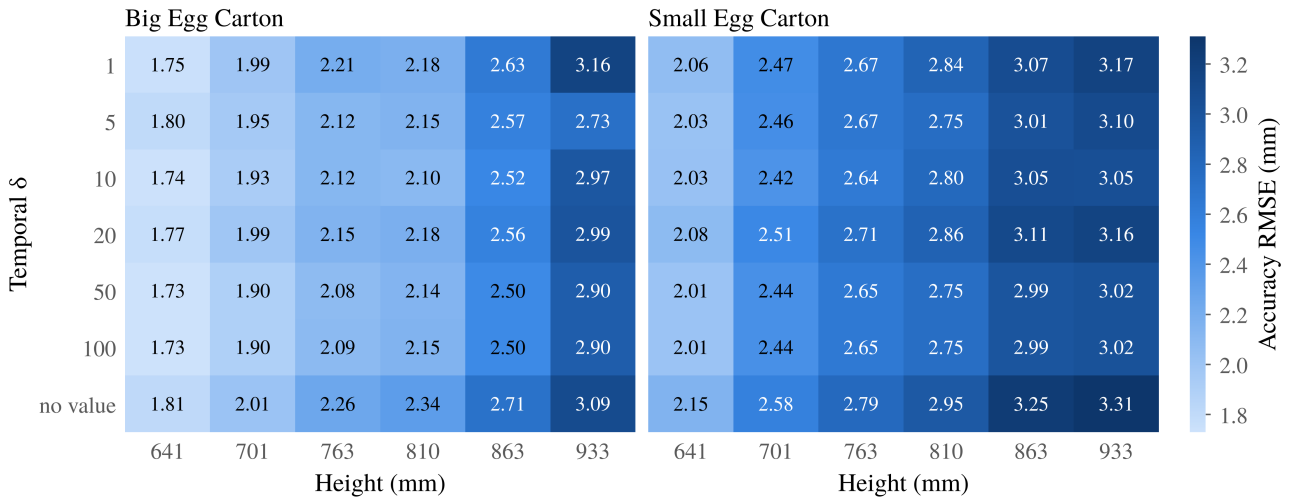

Figure E.2: Median with IRQ % CI of the mean of n = 5–815 scans per cell. Fixed: Tilt = 0. Averaged over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial filter on (2), Spatial magnitude (6), Spatial α (8) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded.

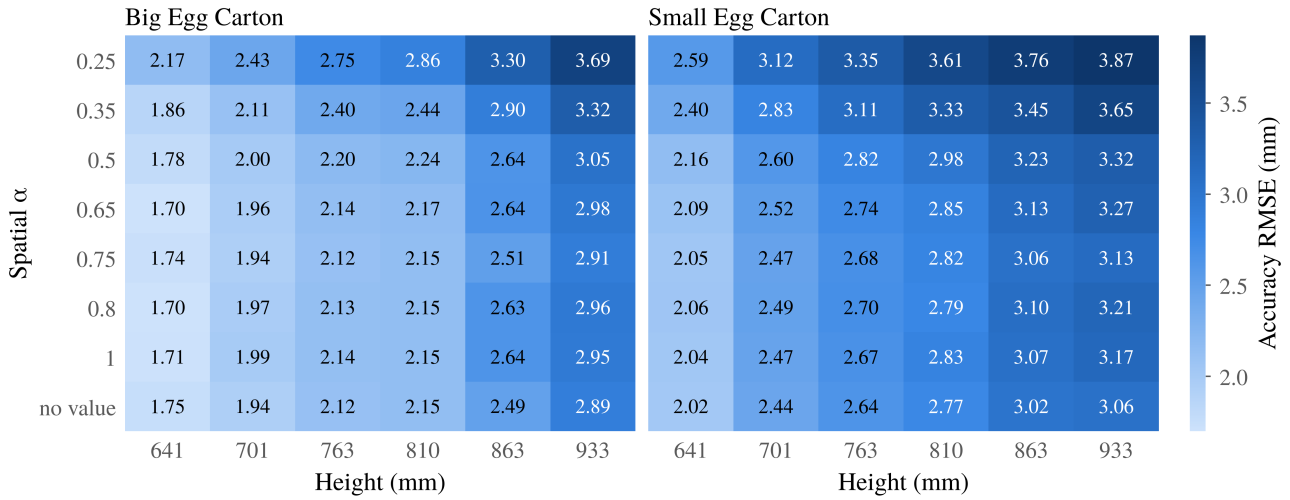

Figure E.3: Median with IQR of n = 5–315 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial magnitude (6), Spatial δ (9), Spatial hole fill (7) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded.

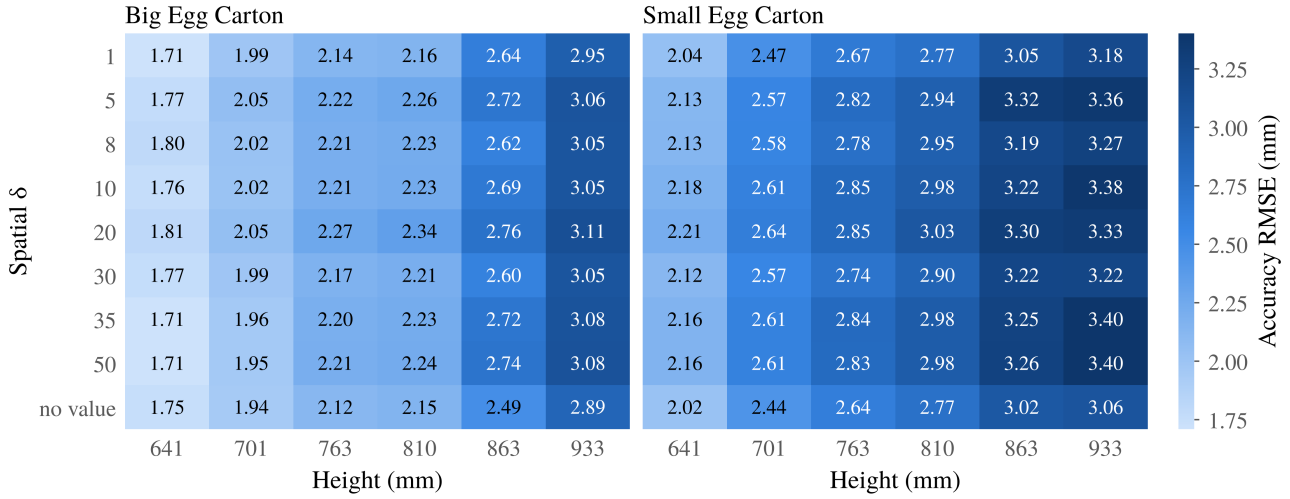

Figure E.4: Median with IQR of n = 5–385 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial magnitude (6), Spatial α (8), Spatial hole fill (7) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded.

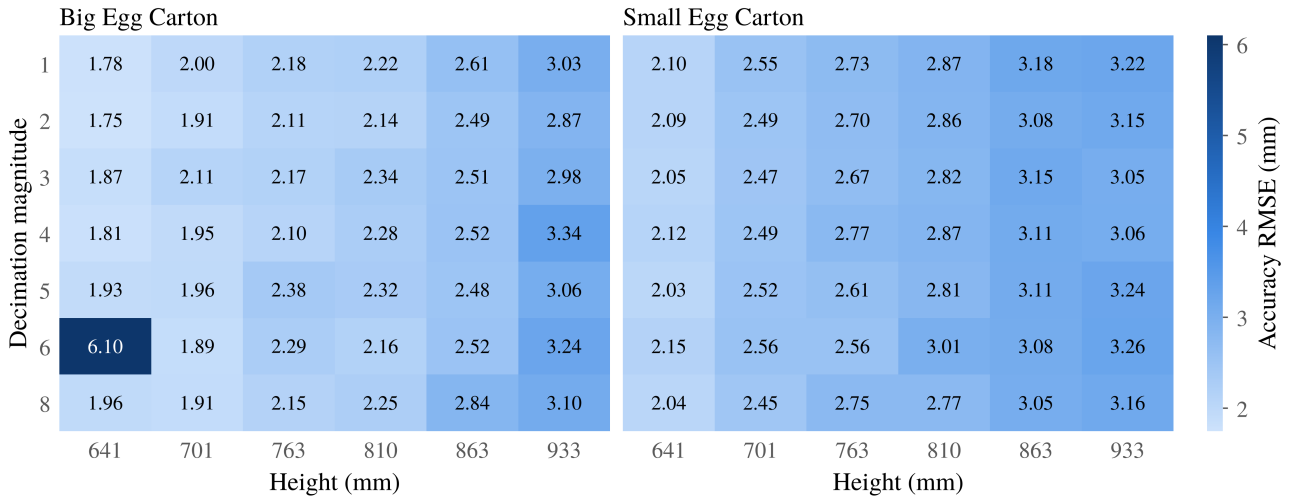

Figure E.5: Median with IQR of n = 5–950 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Spatial filter on (2), Spatial magnitude (6), Spatial α (8), Spatial δ (9), Spatial hole fill (7) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded.

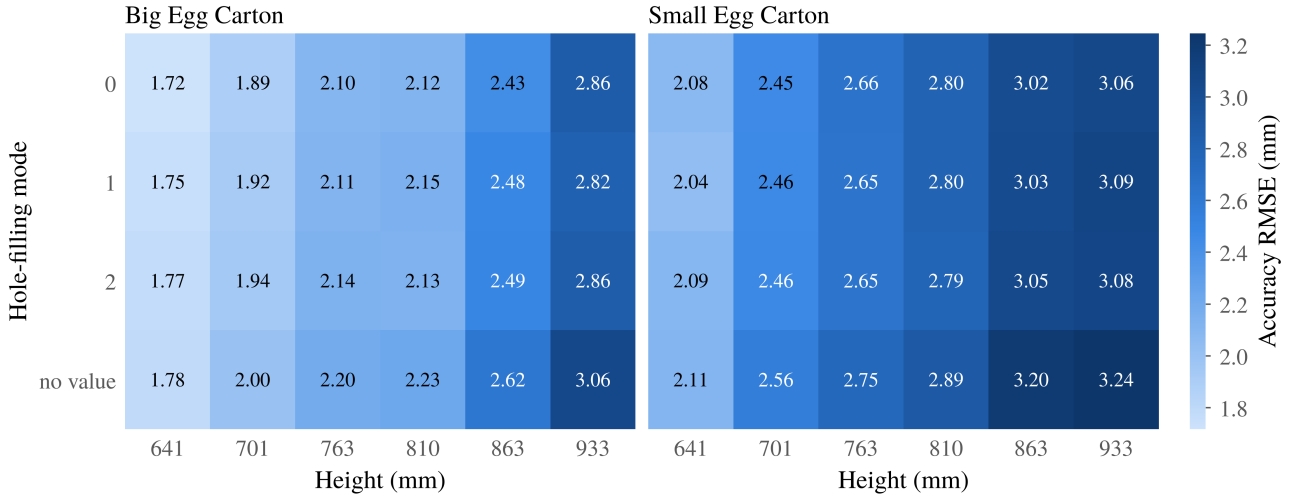

Figure E.6: Median with IQR of n = 55–915 scans per cell. Fixed: Tilt = 0. Median taken over: Rotation (4), Decimation on (2), Decimation magnitude (7), Spatial filter on (2), Spatial magnitude (6), Spatial α (8) and 8 more. Spanning 257 filter configurations. 2,460 rows without a value excluded.

E.2 Combination Heatmaps

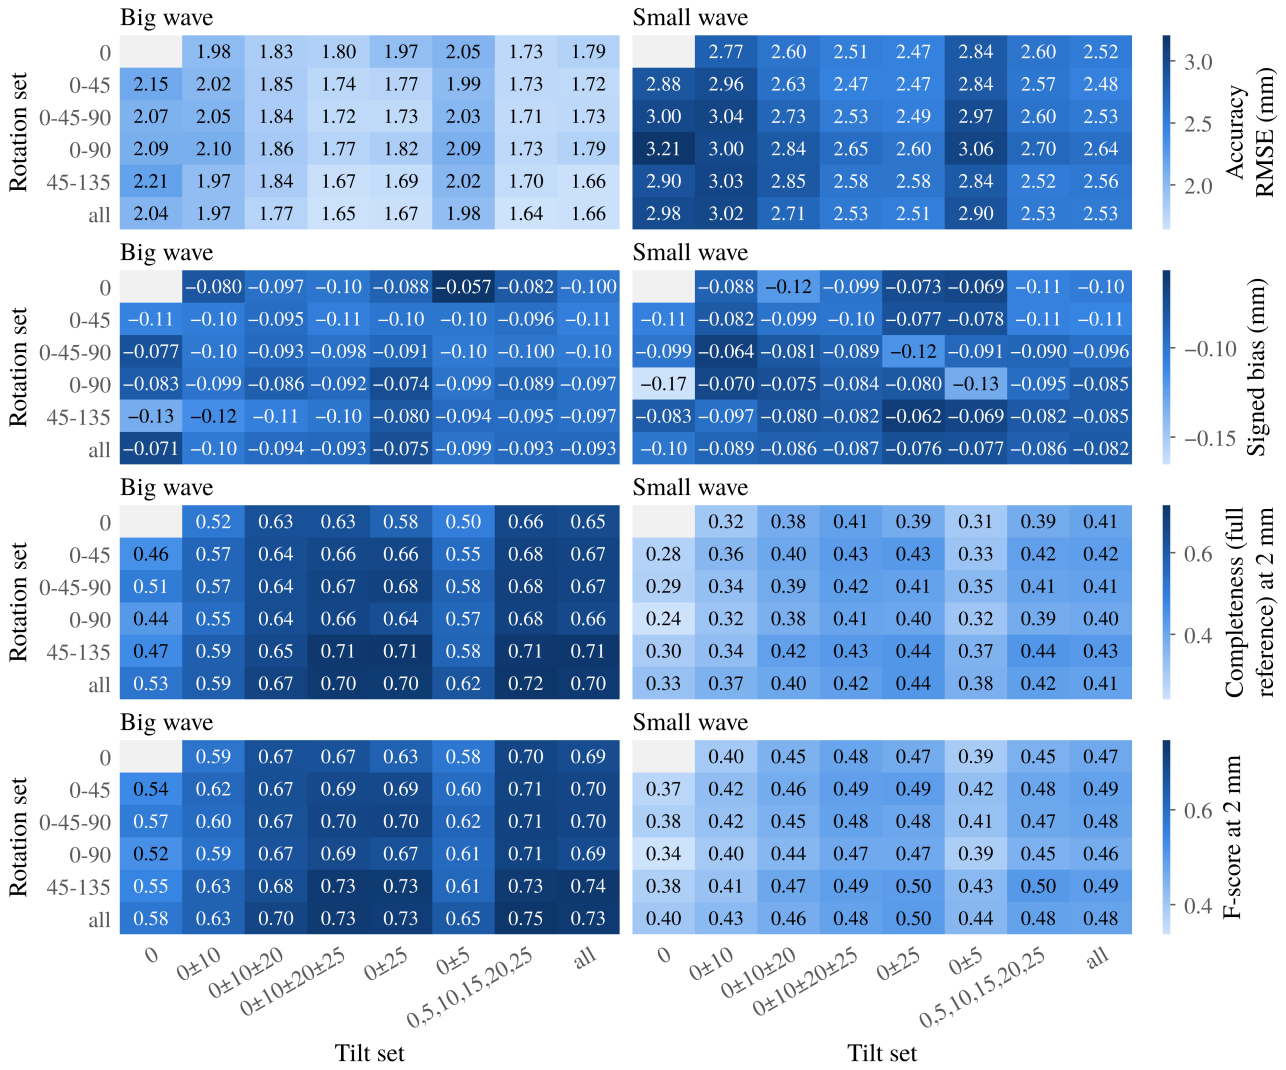

Figure E.7: Grey cells: no data for that combination. Median with IQR of n = 6 point clouds per cell. Fixed: Kind = fused; Fusion cell = 0.0005; Scan voxel = 0.0005; Status = ok. Median taken over: Height (6).

E.3 Parameter Heatmap

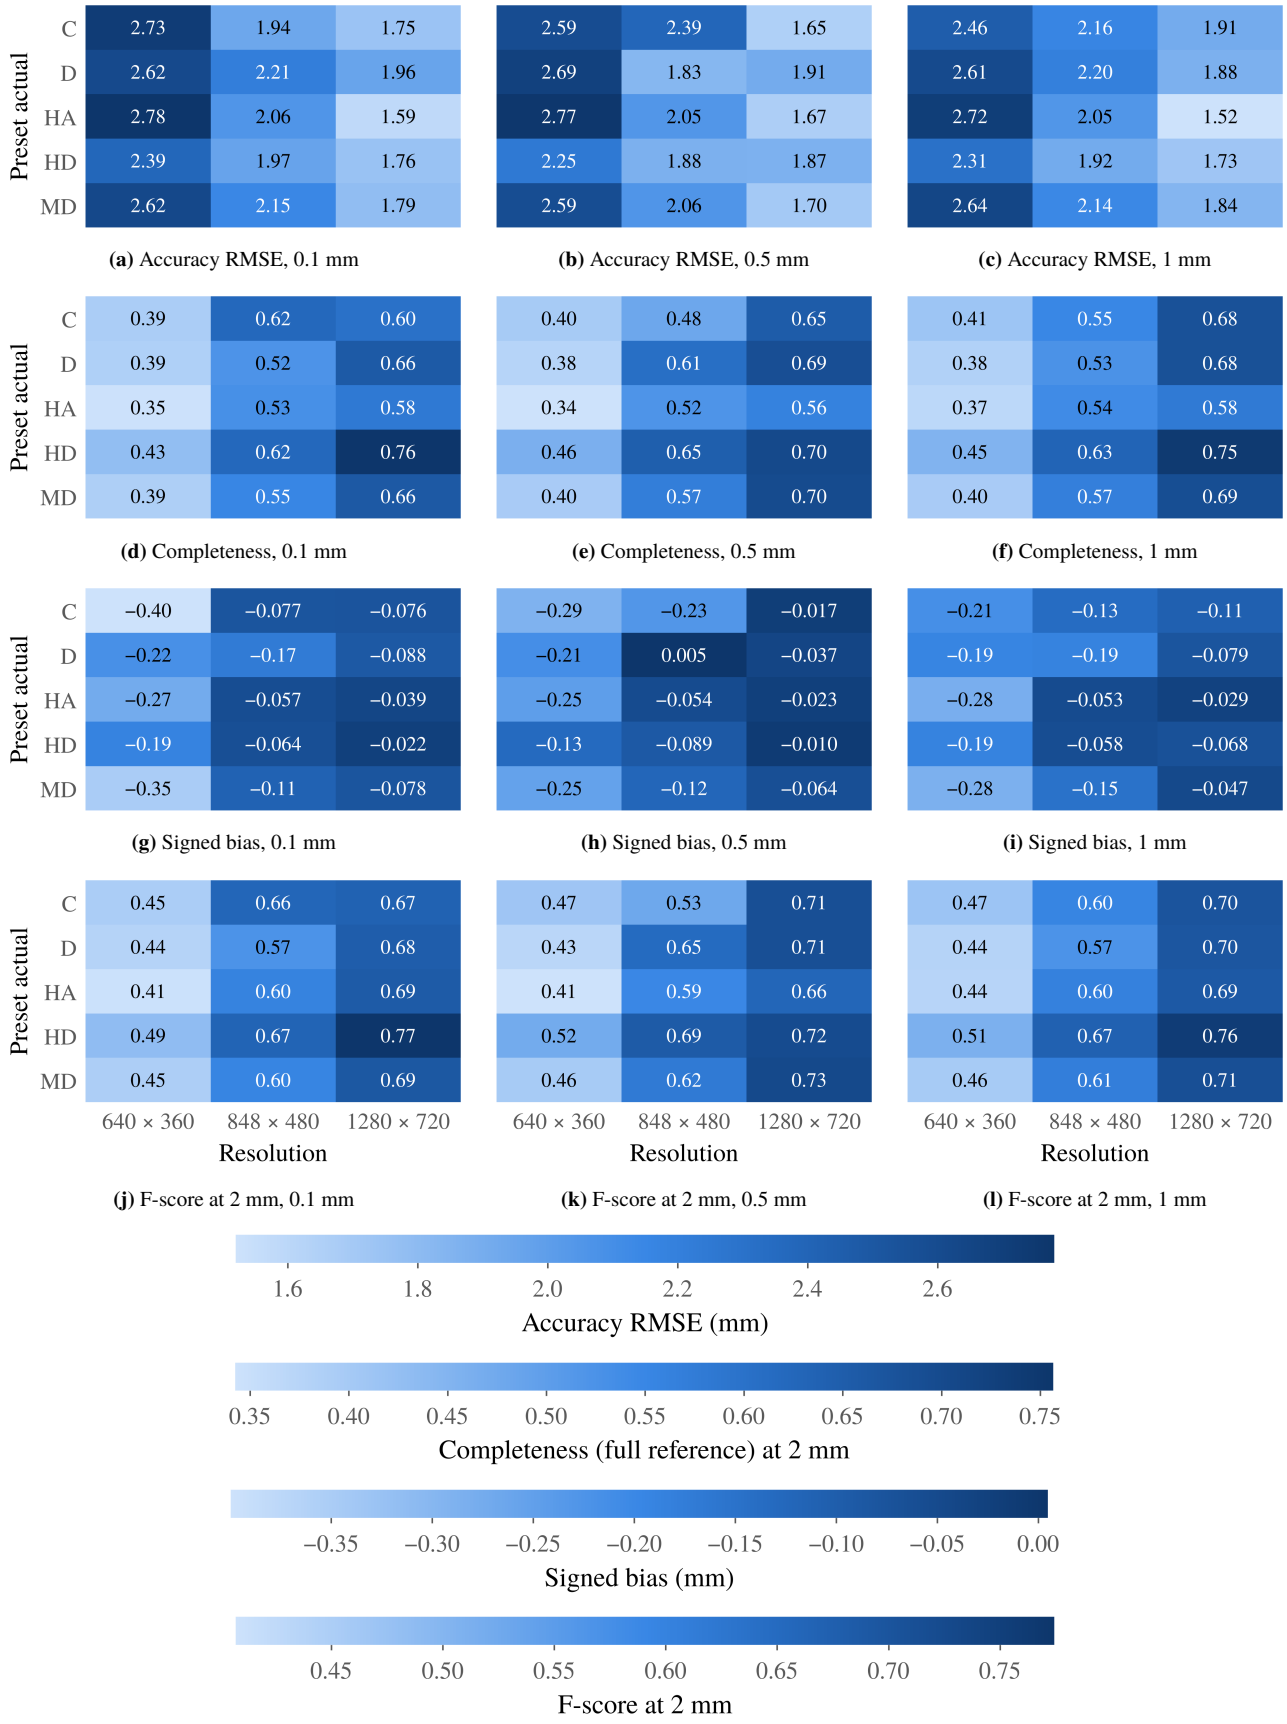

(a) Accuracy RMSE, 0.1 mm **(b)** Accuracy RMSE, 0.5 mm **(c)** Accuracy RMSE, 1 mm

(d) Completeness, 0.1 mm **(e)** Completeness, 0.5 mm **(f)** Completeness, 1 mm

(g) Signed bias, 0.1 mm **(h)** Signed bias, 0.5 mm **(i)** Signed bias, 1 mm

(j) F-score at 2 mm, 0.1 mm **(k)** F-score at 2 mm, 0.5 mm **(l)** F-score at 2 mm, 1 mm

Figure E.8: Median with IQR of n = 1–3 scans per cell. Fixed: Laser power = 150; Emitter = Laser; Fps = 30; Status = ok. Median taken over: Object (2), Preset (5).

E.4 Accuracy Vs Tilt

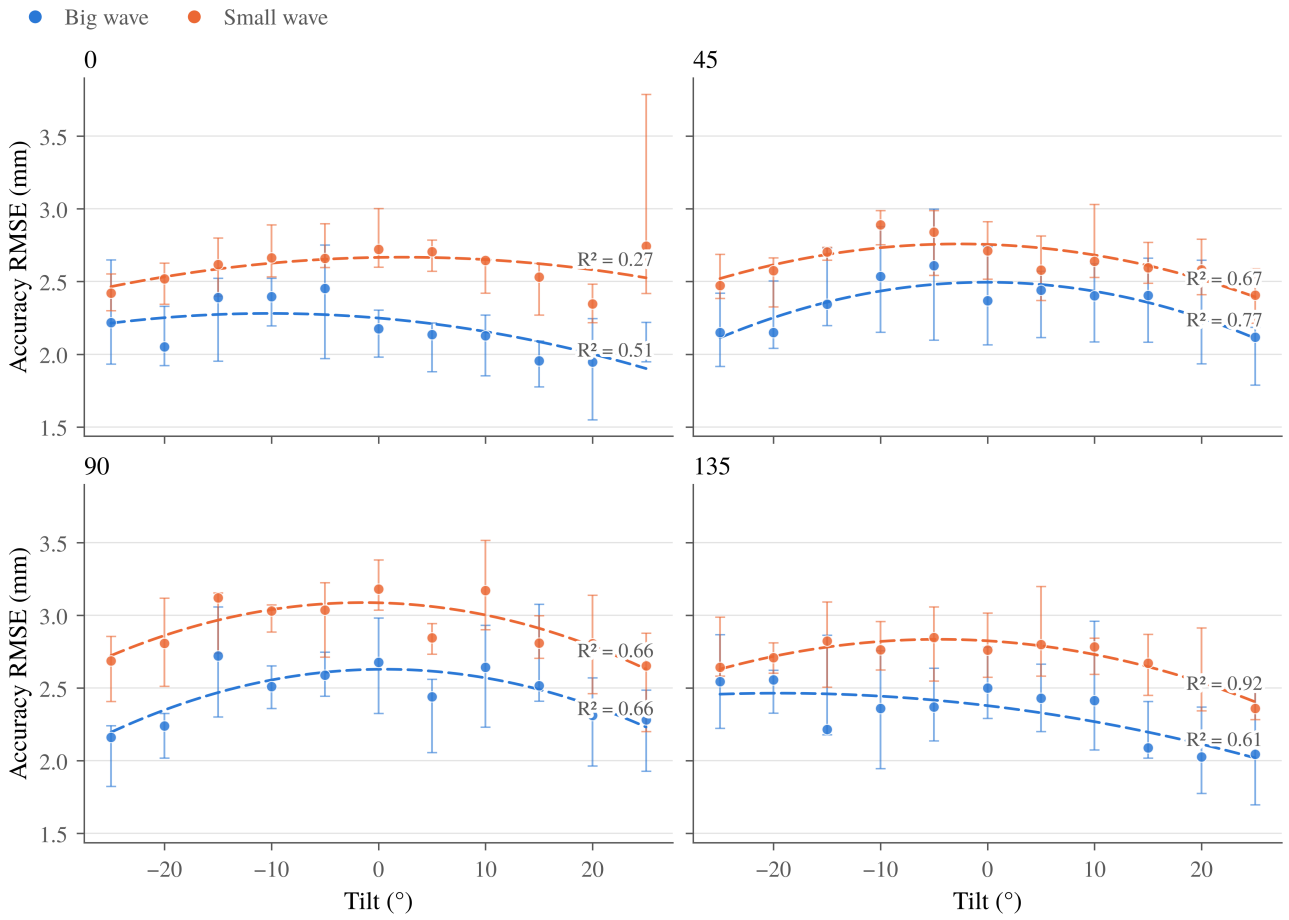

Figure E.9: Dashed lines: quadratic least-squares fit per series and panel (R^2 0.27 to 0.92). Median with IQR of n = 5–6 point clouds per point. Fixed: Kind = single; View family = single; Number of rotations = 1; Number of tilts = 1; Number of scans = 1. Median taken over: Height (6), Status (3). 3 rows without a value excluded.

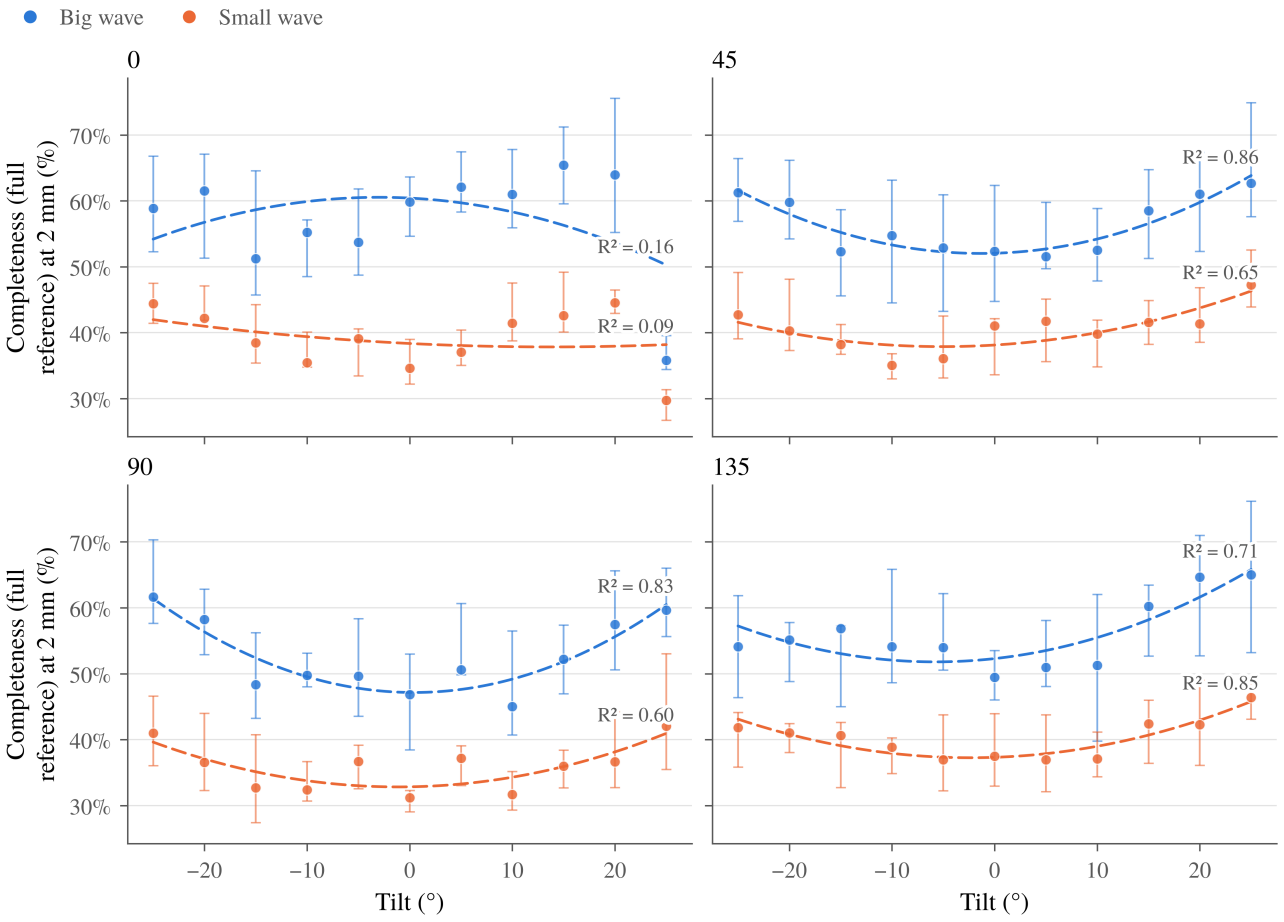

Figure E.10: Dashed lines: quadratic least-squares fit per series and panel (R^2 0.09 to 0.86). Median with IQR of n = 5–6 point clouds per point. Fixed: Kind = single; View family = single; Number of rotations = 1; Number of tilts = 1; Number of scans = 1. Median taken over: Height (6), Status (3). 3 rows without a value excluded.

E.5 Single Vs Fused Error Histograms

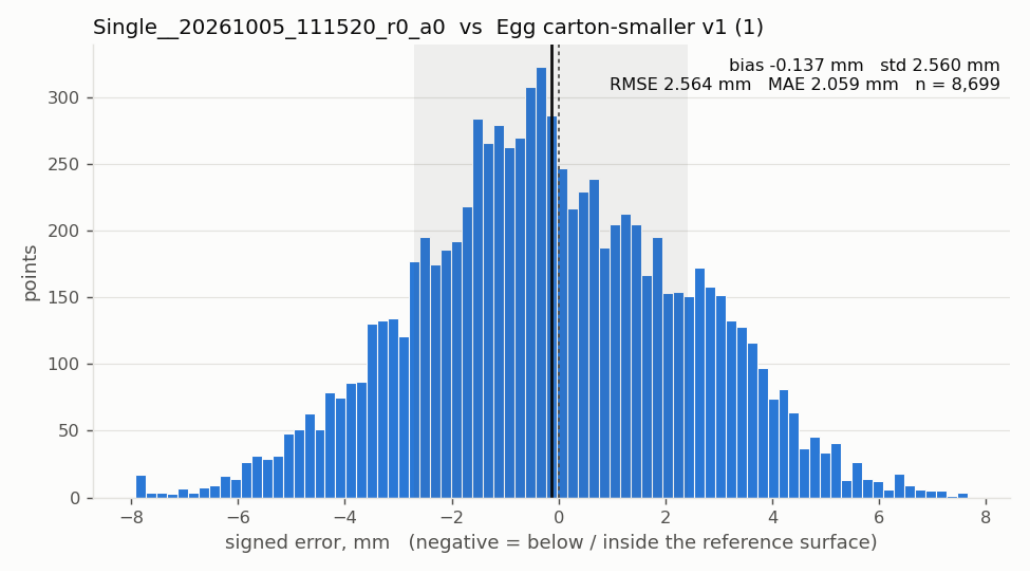

Figure E.11: Histogram Showing Spread of Signed Error

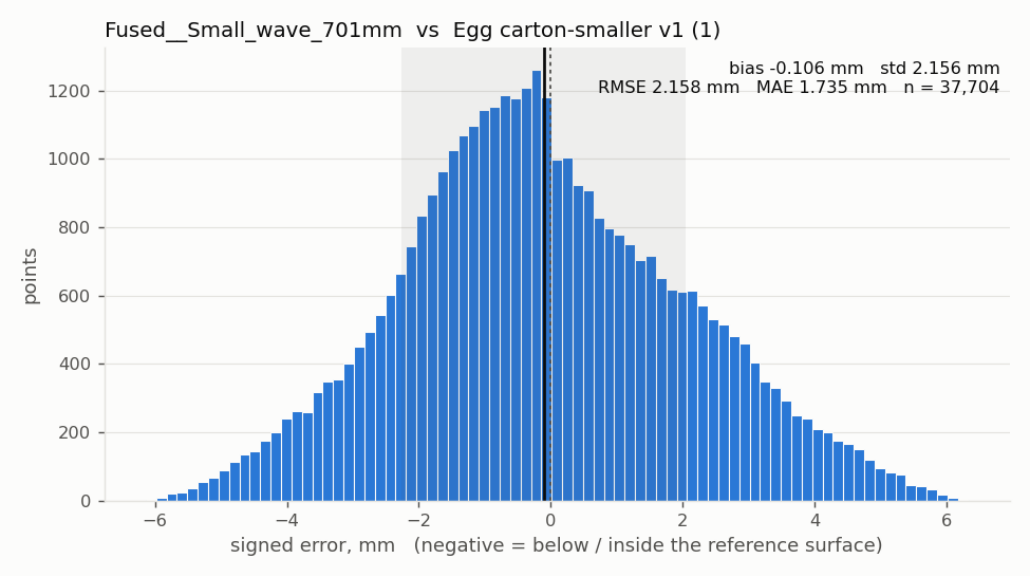

Figure E.12: Histogram Showing Spread of Signed Error

Appendix F

Validation Results

F.1 Reference Shapes Rugosity

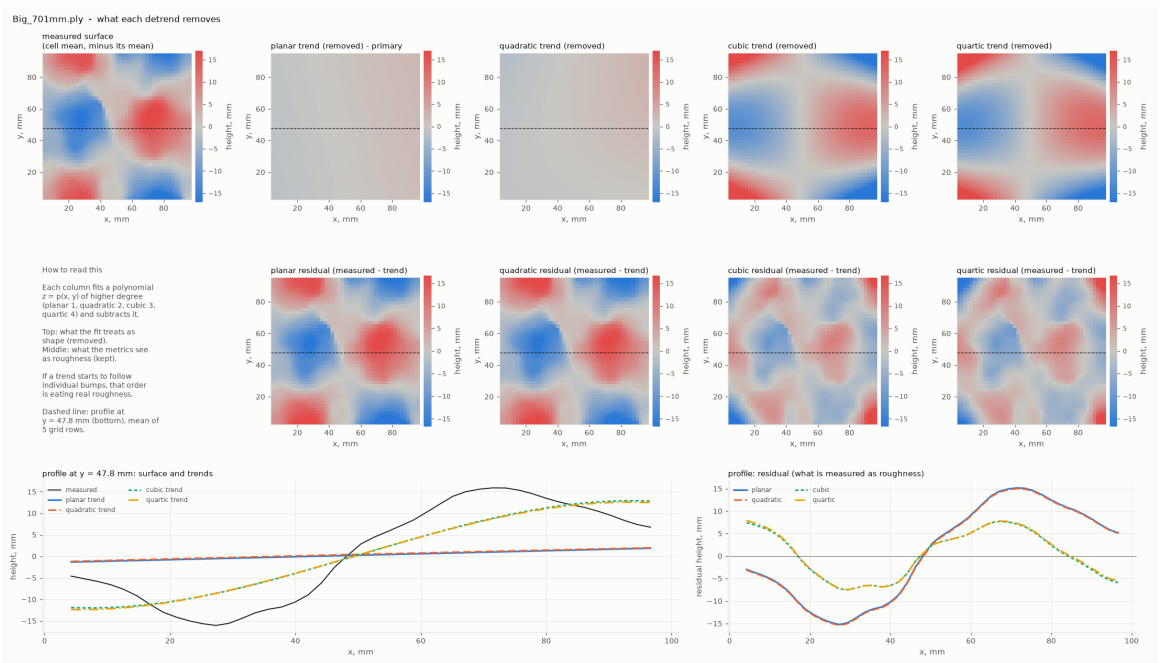

Figure F.1: Big Wave Detrend Explained Single

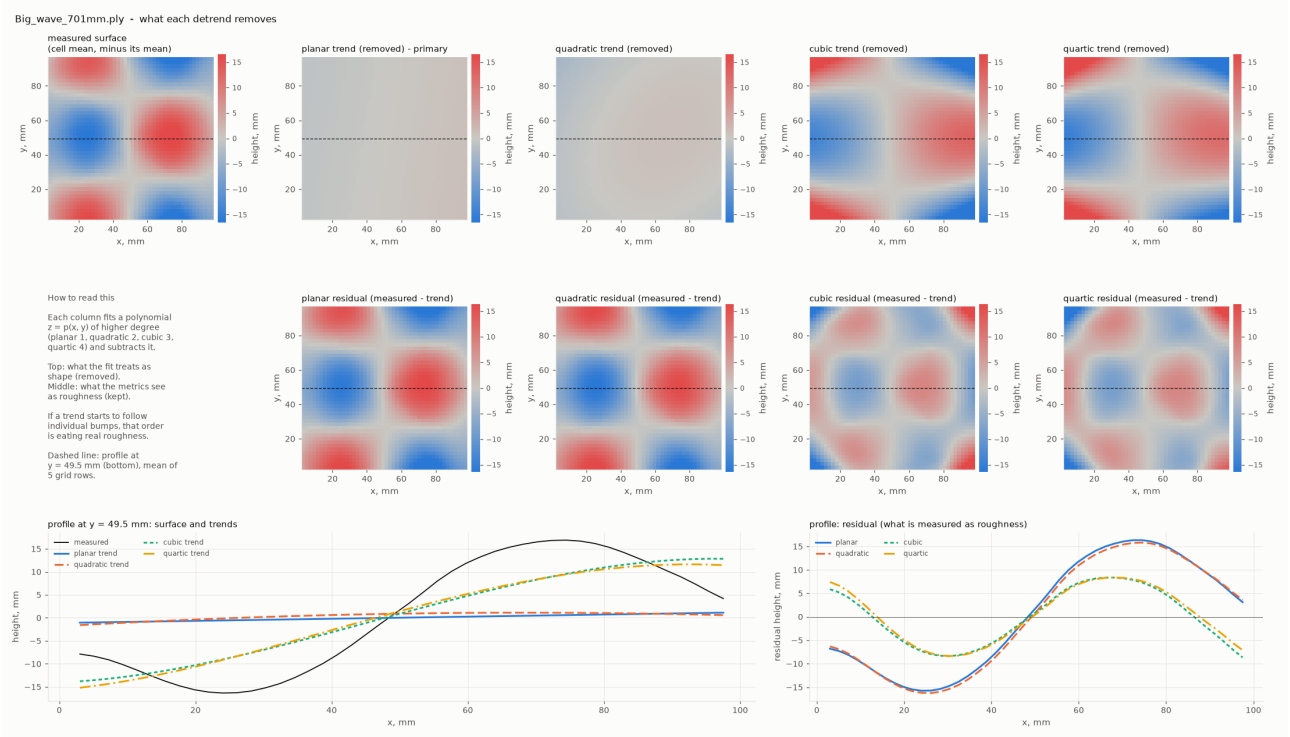

Figure F.2: Big Wave Detrend Explained Fused

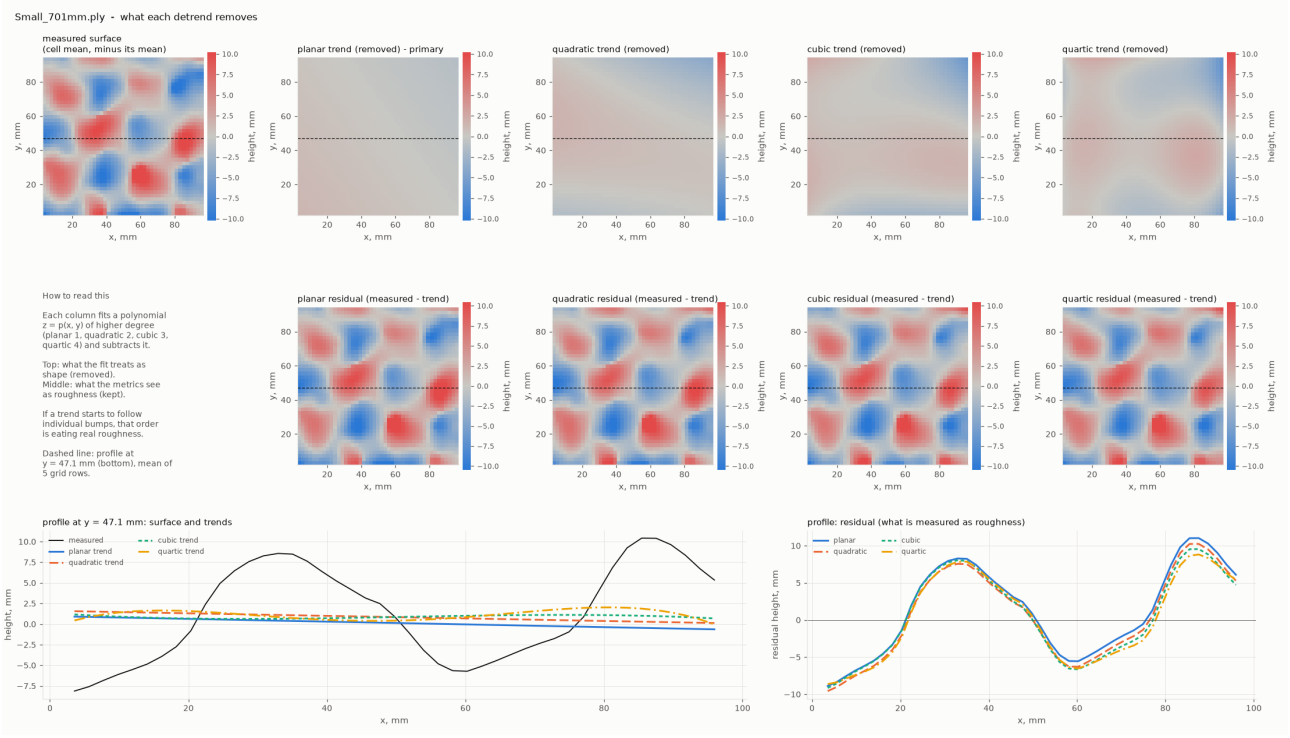

Figure F.3: Small Wave Detrend Explained Single

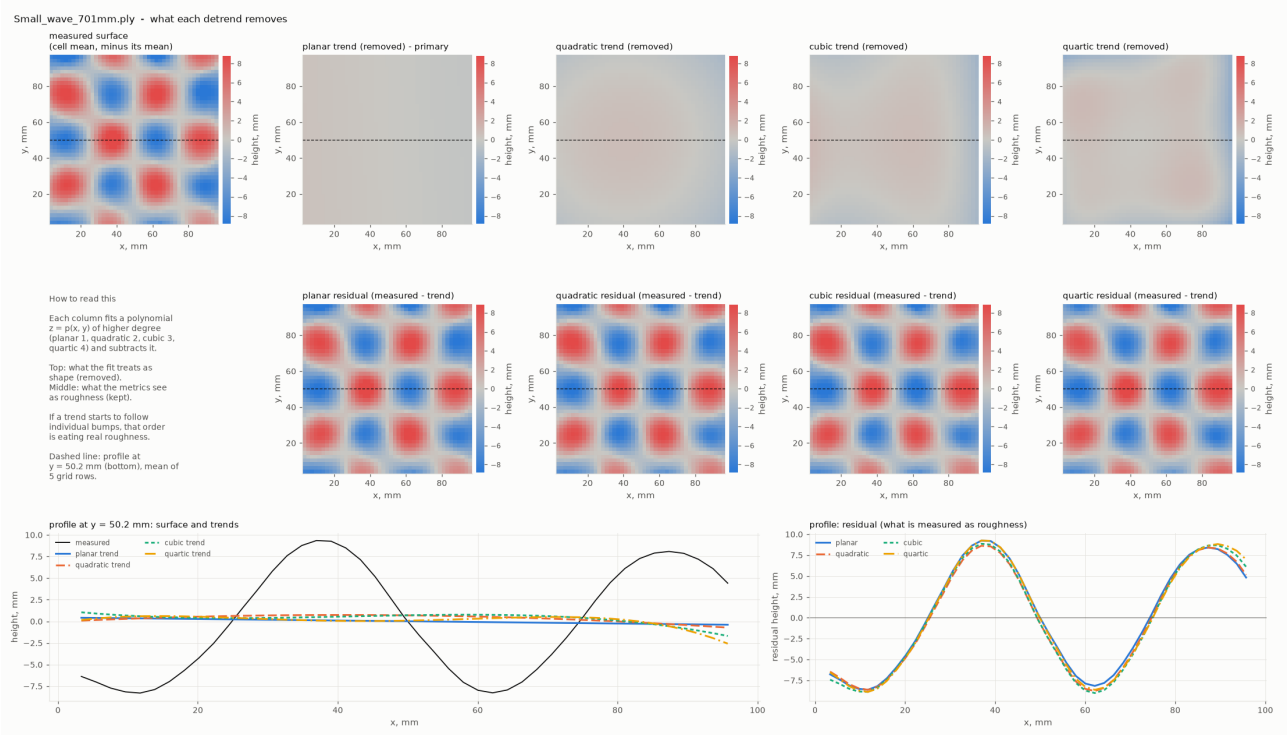

Figure F.4: Small Wave Detrend Explained Fused

Appendix B: Ethics approval

Appendix C: Topic description

The original topic description used during the project selection process is included below. This topic has since been further refined, and the exact problem that was investigated is described in Section 1.2.

Appendix D: ECSA Graduate Attributes

Project mapping to ECSA Graduate Attributes

The table below provides a structured reflection on where each of the required Graduate Attributes (GAs) as defined by Engineering Council of South Africa (ECSA) associated with the final year project is met. This includes reference to the specific Chapter/Section (where applicable).

GA	Graduate Attribute Description	Student reflection: How this outcome was achieved during the project
GA1	*Problem solving*	
GA4	*Investigations, experiments and data analysis*	
GA5	*Use of engineering tools*	
GA6	*Professional and technical communication*	
GA8	*Individual and collaborative teamwork*	
GA9	*Independent learning ability*	

Bibliography

[1] S. Schenone, I. Bartl, and S. F. Thrush, "Sediment microtopography predicts localised benthic ecosystem functioning and the effect of species interactions," *Marine Ecology Progress Series*, vol. 749, pp. 47–55, Nov. 2024.

[2] C. Nolet, A. Poortinga, P. Roosjen, H. Bartholomeus, and G. Ruessink, "Measuring and Modeling the Effect of Surface Moisture on the Spectral Reflectance of Coastal Beach Sand," *PLOS ONE*, vol. 9, no. 11, p. e112151, Nov. 2014.

[3] E. Kristensen, G. Penha-Lopes, M. Delefosse, T. Valdemarsen, C. O. Quintana, and G. T. Banta, "What is bioturbation? The need for a precise definition for fauna in aquatic sciences," *Marine Ecology Progress Series*, vol. 446, pp. 285–302, Feb. 2012.

[4] D. Pillay and G. Branch, "Bioengineering effects of burrowing thalassinidean shrimps on marine soft-bottom ecosystems," in *Oceanography and Marine Biology: An Annual Review*, Jun. 2011, vol. 49, pp. 137–192.

[5] O. Venter, D. Pillay, and K. Prayag, "Water filtration by burrowing sandprawns provides novel insights on endobenthic engineering and solutions for eutrophication," *Scientific Reports*, vol. 10, no. 1, p. 1913, Feb. 2020.

[6] S. Schenone, J. E. Hewitt, J. Hillman, R. Gladstone-Gallagher, J. Gammal, C. Pilditch, A. M. Lohrer, E. Ferretti, M. Azhar, P. Delmas, and S. F. Thrush, "Seafloor sediment microtopography as a surrogate for biodiversity and ecosystem functioning," *Ecological Applications*, vol. 35, no. 1, p. e3069, 2025.

[7] "ISO 25178-2:2021," https://www.iso.org/standard/74591.html.

[8] M. Azhar, J. R. Hillman, T. Gee, S. Schenone, W. van der Mark, S. F. Thrush, and P. Delmas, "An RGB-D framework for capturing soft-sediment microtopography," *Methods in Ecology and Evolution*, vol. 13, no. 8, pp. 1730–1745, 2022.

[9] S. Kamal, S. Y. Lee, and J. Warnken, "Investigating three-dimensional mesoscale habitat complexity and its ecological implications using low-cost RGB-D sensor technology," *Methods in Ecology and Evolution*, vol. 5, no. 9, pp. 845–853, 2014.

[10] G. Halmetschlager-Funek, M. Suchi, M. Kampel, and M. Vincze, "An Empirical Evaluation of Ten Depth Cameras: Bias, Precision, Lateral Noise, Different Lighting Conditions and Materials, and Multiple Sensor Setups in Indoor Environments," *IEEE Robotics & Automation Magazine*, vol. 26, no. 1, pp. 67–77, Mar. 2019.

[11] "Intel RealSense Product Family D400 Series," 2023.

[12] L. Keselman, J. I. Woodfill, A. Grunnet-Jepsen, and A. Bhowmik, "Intel(R) RealSense(TM) Stereoscopic Depth Cameras," in *2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, Jul. 2017, pp. 1267–1276.

[13] A. Grunner-Jepsen, J. N. Sweetser, P. Winer, A. Takagi, and J. Woodfill, "Projectors for D400 Series Depth Cameras."

[14] M. Carfagni, R. Furferi, L. Governi, C. Santarelli, M. Servi, F. Uccheddu, and Y. Volpe, "Metrological and Critical Characterization of the Intel D415 Stereo Depth Camera," *Sensors*, vol. 19, no. 3, p. 489, Jan. 2019.

[15] L. Rustler, V. Volprecht, and M. Hoffmann, "Empirical Comparison of Four Stereoscopic Depth Sensing Cameras for Robotics Applications," *IEEE Access*, vol. 13, pp. 67 564–67 577, 2025.

[16] P. Nel and GM. Branch, "Assessment of the abundance and distribution of burrowing sandprawns and mudprawns (Callichirus and Upogebia species) in Langebaan Lagoon, South Africa," *African Journal of Marine Science*, vol. 35, no. 2, pp. 195–208, Aug. 2013.

[17] G. M. Branch and A. Pringle, "The impact of the sand prawn *Callianassa kraussi* Stebbing on sediment turnover and on bacteria, meiofauna, and benthic microflora," *Journal of Experimental Marine Biology and Ecology*, vol. 107, no. 3, pp. 219–235, Apr. 1987.

[18] D. Pillay, G. M. Branch, and A. T. Forbes, "Experimental evidence for the effects of the thalassinidean sandprawn Callianassa kraussi on macrobenthic communities," *Marine Biology*, vol. 152, no. 3, pp. 611–618, Sep. 2007.

[19] S. F. Thrush, J. E. Hewitt, G. A. Funnell, V. J. Cummings, J. Ellis, D. Schultz, D. Talley, and A. Norkko, "Fishing disturbance and marine biodiversity: Role of habitat structure in simple soft-sediment systems," *Marine Ecology Progress Series*, vol. 221, pp. 255–264, Oct. 2001.

[20] S. Schenone and S. F. Thrush, "Unraveling ecosystem functioning in intertidal soft sediments: The role of density-driven interactions," *Scientific Reports*, vol. 10, no. 1, p. 11909, Jul. 2020.

[21] A. Friedman, O. Pizarro, S. B. Williams, and M. Johnson-Roberson, "Multi-Scale Measures of Rugosity, Slope and Aspect from Benthic Stereo Image Reconstructions," *PLOS ONE*, vol. 7, no. 12, p. e50440, Dec. 2012.

[22] C. D. Storlazzi, P. Dartnell, G. A. Hatcher, and A. E. Gibbs, "End of the chain? Rugosity and fine-scale bathymetry from existing underwater digital imagery using structure-from-motion (SfM) technology," *Coral Reefs*, vol. 35, no. 3, pp. 889–894, Sep. 2016.

[23] R. Leach, *Characterisation of Areal Surface Texture*. Springer Science & Business Media, Apr. 2013.

[24] T. D. B. Jacobs, T. Junge, and L. Pastewka, "Quantitative characterization of surface topography using spectral analysis," *Surface Topography: Metrology and Properties*, vol. 5, no. 1, p. 013001, Jan. 2017.

[25] K. D. Mankoff and T. A. Russo, "The Kinect: A low-cost, high-resolution, short-range 3D camera," *Earth Surface Processes and Landforms*, vol. 38, no. 9, pp. 926–936, 2013.

[26] A. Vit and G. Shani, "Comparing RGB-D Sensors for Close Range Outdoor Agricultural Phenotyping," *Sensors*, vol. 18, no. 12, p. 4413, Dec. 2018.

[27] A. Grunner-Jepsen, J. N. Sweetser, and J. Woodfill, "Tuning depth cameras for best performance."

[28] A. Grunnet-Jepsen and Tong, "Depth Post-Processing for Intel® RealSense™™ D400 Depth Cameras," https://www.intel.com/content/www/us/en/content-details/842031/depth-post-processing-for-intel-realsense-d400-depth-cameras.html.

[29] "Camera Depth Testing Methodology."

[30] M. Servi, E. Mussi, A. Profili, R. Furferi, Y. Volpe, L. Governi, and F. Buonamici, "Metrological Characterization and Comparison of D415, D455, L515 RealSense Devices in the Close Range," *Sensors*, vol. 21, no. 22, p. 7770, Nov. 2021.

[31] C. V. Nguyen, S. Izadi, and D. Lovell, "Modeling Kinect Sensor Noise for Improved 3D Reconstruction and Tracking," in *Visualization & Transmission 2012 Second International Conference on 3D Imaging, Modeling, Processing*, Oct. 2012, pp. 524–530.

[32] B. Curless and M. Levoy, "A volumetric method for building complex models from range images," in *Proceedings of the 23rd Annual Conference on Computer Graphics and Interactive Techniques*, ser. SIGGRAPH '96. New York, NY, USA: Association for Computing Machinery, Aug. 1996, pp. 303–312.

[33] R. A. Newcombe, S. Izadi, O. Hilliges, D. Molyneaux, D. Kim, A. J. Davison, P. Kohi, J. Shotton, S. Hodges, and A. Fitzgibbon, "KinectFusion: Real-time dense surface mapping and tracking," in *2011 10th IEEE International Symposium on Mixed and Augmented Reality*, Oct. 2011, pp. 127–136.

[34] N. Gelfand, L. Ikemoto, S. Rusinkiewicz, and M. Levoy, "Geometrically stable sampling for the ICP algorithm," in *Fourth International Conference on 3-D Digital Imaging and Modeling, 2003. 3DIM 2003. Proceedings.*, Oct. 2003, pp. 260–267.

[35] E. Olson, "AprilTag: A robust and flexible visual fiducial system," in *2011 IEEE International Conference on Robotics and Automation*, May 2011, pp. 3400–3407.

[36] J. Wang and E. Olson, "AprilTag 2: Efficient and robust fiducial detection," in *2016 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, Oct. 2016, pp. 4193–4198.

[37] R. Hartley, A. Zisserman *et al.*, *Multiple View Geometry in Computer Vision*. Cambridge university press, Cambridge, 2003, vol. 2.

[38] Z. Zhang, "A flexible new technique for camera calibration," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 22, no. 11, pp. 1330–1334, Nov. 2000.

[39] A. Grunnet-Jepsen, J. N. Sweetser, and J. Woodfill, "Best-Known-Methods for Tuning Intel® RealSense™ D400 Depth Cameras for Best Performance."

[40] K. S. Arun, T. S. Huang, and S. D. Blostein, "Least-Squares Fitting of Two 3-D Point Sets," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. PAMI-9, no. 5, pp. 698–700, Sep. 1987.

[41] S. Umeyama, "Least-squares estimation of transformation parameters between two point patterns," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 13, no. 4, pp. 376–380, Apr. 1991.

[42] G. Schweighofer and A. Pinz, "Robust Pose Estimation from a Planar Target," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 28, no. 12, pp. 2024–2030, Dec. 2006.

[43] R. B. Rusu, N. Blodow, and M. Beetz, "Fast Point Feature Histograms (FPFH) for 3D registration," in *2009 IEEE International Conference on Robotics and Automation*, May 2009, pp. 3212–3217.

[44] M. A. Fischler and R. C. Bolles, "Random sample consensus: A paradigm for model fitting with applications to image analysis and automated cartography," *Communications of the ACM*, vol. 24, no. 6, pp. 381–395, Jun. 1981.

[45] P. J. Besl and N. D. McKay, "Method for registration of 3-D shapes," in *Sensor Fusion IV: Control Paradigms and Data Structures*, vol. 1611. SPIE, Apr. 1992, p. 586.

[46] Y. Chen and G. Medioni, "Object modelling by registration of multiple range images," *Image and Vision Computing*, vol. 10, no. 3, pp. 145–155, Apr. 1992.

[47] K.-L. Low, "Linear Least-Squares Optimization for Point-to-Plane ICP Surface Registration."

[48] S. Rusinkiewicz and M. Levoy, "Efficient variants of the ICP algorithm," in *Proceedings Third International Conference on 3-D Digital Imaging and Modeling*, May 2001, pp. 145–152.

[49] P. J. Huber and E. M. Ronchetti, *Robust Statistics*. John Wiley & Sons, Sep. 2011.

[50] D. Levin, "The approximation power of moving least-squares," *Mathematics of computation*, vol. 67, no. 224, pp. 1517–1531, 1998.

[51] M. Alexa, J. Behr, D. Cohen-Or, S. Fleishman, D. Levin, and C. Silva, "Computing and rendering point set surfaces," *IEEE Transactions on Visualization and Computer Graphics*, vol. 9, no. 1, pp. 3–15, Jan. 2003.

[52] A. Knapitsch, J. Park, Q.-Y. Zhou, and V. Koltun, "Tanks and temples: Benchmarking large-scale scene reconstruction," *ACM Transactions on Graphics (TOG)*, vol. 36, no. 4, pp. 78:1–78:13, Jul. 2017.

[53] D. Huttenlocher, G. Klanderman, and W. Rucklidge, "Comparing images using the Hausdorff distance," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 15, no. 9, pp. 850–863, Sep. 1993.